


CS 31006: Computer Networks – Internet Routing

**Department of Computer
Science and Engineering**



**INDIAN INSTITUTE OF TECHNOLOGY
KHARAGPUR**



Rajat Subhra Chakraborty
rschakraborty@cse.iitkgp.ac.in

Sandip Chakraborty
sandipc@cse.iitkgp.ac.in

A Practice Problem

| IP | Netmask | Next hop |
|--------------|---------------|----------|
| 128.96.170.0 | 255.255.254.0 | Eth0 |
| 128.96.168.0 | 255.255.254.0 | Eth1 |
| 128.96.166.0 | 255.255.254.0 | R2 |
| 128.96.164.0 | 255.255.252.0 | R3 |
| 0.0.0.0 | 0.0.0.0 | R4 |

- What would be next hop for the following destination IPs?
 - 128.96.171.92
 - 128.96.167.151
 - 128.96.163.151

Problem courtesy: Computer Networks, Larry L Peterson and Bruce S Davie

A Practice Problem

| IP | Netmask | Next hop |
|--------------|---------------|----------|
| 128.96.170.0 | 255.255.254.0 | Eth0 |
| 128.96.168.0 | 255.255.254.0 | Eth1 |
| 128.96.166.0 | 255.255.254.0 | R2 |
| 128.96.164.0 | 255.255.252.0 | R3 |
| 0.0.0.0 | 0.0.0.0 | R4 |

- What would be next hop for the following destination IPs?

- 128.96.171.92

171 -- 10101011

254 -- 11111110

Logical AND -- 10101010 --
170

Destination is Eth0

Problem courtesy: Computer Networks, Larry L Peterson and Bruce S Davie

A Practice Problem

| IP | Netmask | Next hop |
|--------------|---------------|----------|
| 128.96.170.0 | 255.255.254.0 | Eth0 |
| 128.96.168.0 | 255.255.254.0 | Eth1 |
| 128.96.166.0 | 255.255.254.0 | R2 |
| 128.96.164.0 | 255.255.252.0 | R3 |
| 0.0.0.0 | 0.0.0.0 | R4 |

- What would be next hop for the following destination IPs?

- 128.96.167.151

167 -- 10100111

254 -- 11111110

Logical AND -- 10100110 -- 166

Problem courtesy: Computer Networks, Larry L Peterson and Bruce S Davie

A Practice Problem

| IP | Netmask | Next hop |
|--------------|---------------|----------|
| 128.96.170.0 | 255.255.254.0 | Eth0 |
| 128.96.168.0 | 255.255.254.0 | Eth1 |
| 128.96.166.0 | 255.255.254.0 | R2 |
| 128.96.164.0 | 255.255.252.0 | R3 |
| 0.0.0.0 | 0.0.0.0 | R4 |

- What would be next hop for the following destination IPs?

- 128.96.167.151

167 -- 10100111

252 -- 11111100

Logical AND -- 10100100 --
164

Longest Match is 254
Next hop Eth1


Problem courtesy: Computer Networks, 4th Edition, by Andrew S. Tanenbaum and Bruce S. Davie

Routing in the Internet

- Two different variants of routing protocols – (a) **Intra-AS routing protocols**, and (b) **Inter-AS routing protocols**
- **Intra-AS (Intra Domain) Routing Protocols:** Routing within an AS; Example- Routing Information Protocol (RIP), Open Shortest Path First (OSPF). Sometimes, they are called as **Interior Gateway Protocols (IGP)**
- **Inter-AS (Inter Domain) Routing Protocols:** Routing among the various ASes based on peering relationship; Example: Border Gateway Protocol (BGP). These are known as **Exterior Gateway Protocols (EGP)**



The Optimality Principle (Bellman 1957)

If Router J is on the optimal path from Router I to Router K, then the optimal path from J to K also falls among the same route.



Sink Tree

The set of optimal routes from **all sources** to a **given destination** form a tree **rooted at the destination**. Such a tree is called a **sink tree**.



**Can you tell, why
we are more
focused to a
destination
centric tree,
rather than a**

Sink Tree

- Sink tree for a destination may not be unique – there can be an alternate path between two routers with the same path length.
- We may use a more generic data structure – a **Directed Acyclic Graph (DAG)**, denoting all possible paths from all the source to a destination.
- Both a tree or a DAG is okay – our objective here is to ensure **loop-free routing**

Distance Vector Routing

- Operates by having each router maintain a table (i.e., a vector) giving the best known distance to each destination and which link to use to get there.
- A distributed version of the **Bellman-Ford Algorithm for Shortest Path.**
- Used in **Routing Information Protocol (RIP)** – the original ARPANET routing algorithm.


- **Bellman Ford Equation (Dynamic Programming)**

: Cost of least cost path from node u to node v

: Cost from node u to node v where N_u is the set of neighbors for node u

- We start from the sink (the destination), and recursively construct the shortest path from all the nodes in the subnet.

Bellman Ford in Centralized Network




Initialize

| S | A | B | C | D | E |
|---|---|----|----------|----------|----------|
| 0 | 8 | 20 | ∞ | ∞ | ∞ |
| - | S | S | - | - | - |

Iteration 1

| S | A | B | C | D | E |
|---|---|----|----|----|----|
| 0 | 8 | 20 | 14 | 21 | 22 |
| - | S | S | A | B | B |

Bellman Ford in Centralized Network




Iteration 1

| S | A | B | C | D | E |
|---|---|----|----|----|----|
| 0 | 8 | 20 | 14 | 21 | 22 |
| - | S | S | A | B | B |

Iteration 2

| S | A | B | C | D | E |
|---|---|----|----|----|----|
| 0 | 8 | 17 | 14 | 21 | 15 |
| - | S | C | A | B | C |

Bellman Ford in Centralized Network




Iteration 2

| S | A | B | C | D | E |
|---|---|----|----|----|----|
| 0 | 8 | 17 | 14 | 21 | 15 |
| - | S | C | A | B | C |

Iteration 3

| S | A | B | C | D | E |
|---|---|----|----|----|----|
| 0 | 8 | 17 | 14 | 17 | 15 |
| - | S | C | A | E | C |

Bellman Ford in Centralized Network




Iteration 3

| S | A | B | C | D | E |
|---|---|----|----|----|----|
| 0 | 8 | 17 | 14 | 17 | 15 |
| - | S | C | A | E | C |

Iteration 4


| S | A | B | C | D | E |
|---|---|----|----|----|----|
| 0 | 8 | 17 | 14 | 17 | 15 |
| - | S | C | A | E | C |

Distributed Bellman Ford To Distance Vector Routing



- How do we move from one iteration to the next iteration for Bellman Ford algorithm?
- We need the **currently computed distance information** from each node to others in the network - this is called the **distance vector**
- Used to find out the routes between any two nodes
- In RIP with CIDR, the routing table

Distributed Bellman Ford To Distance Vector Routing



| | S | A | | S |
|---|------|------|--|------|
| S | 00,- | 08,S | | 00,- |
| A | 08,A | 00,- | | 08,A |
| B | 20,B | 9,C | | 17,A |
| C | 23,B | 6,C | | 14,A |
| D | 21,B | 9,C | | 17,A |
| E | 15,A | 7,C | | 15,A |


Centralized Bellman Ford Versus Distance Vector Routing

- Note the differences between **centralized Bellman Ford algorithm** and its distributed version used in **distance vector routing**.
- We are destination oriented (path to sink) for both the variants
- **Centralized Bellman Ford:** We find out the shortest path from all other nodes to the sink
- **Distance Vector Routing:** We find out the shortest path from the current node to all other nodes that can work as the sink


Centralized Bellman Ford Versus Distance Vector Routing

- **Can you guess the reason for this? – Think of the convergence criteria**
- **Centralized Bellman Ford:** You know the number of iterations required. For n number of nodes, you need $n-1$ iterations.
 - **Requires synchronous updates of distance tables**
- **Distance Vector Routing:** You do not know the number of iterations required, number of nodes (n) changes dynamically, distance vector updates are asynchronous
 - **Update the local routing table asynchronously as and when you receive an update from a neighbor**

The Count to Infinity Problem



(a)



(b)

Link A-B comes

up

**Link A-B gets
down**

Solving Count to Infinity

- **Split Horizon:** Routing information is prevented from exiting the router on an interface through which the information was received.
 - A router needs to wait for a timeout to remove an entry from a routing table
- **Poisson Reverse:** All known networks are advertised in each routing update. However, those networks learned through a specific interface are advertised as **unreachable** in the routing announcements sent out to that interface.
 - Routes are immediately removed when they are marked as unreachable from all the routing updates from all the neighbors.

Limitations of Distance Vector Routing / RIP

- The resolution to the counting to infinity problem enforces a maximum cost for a network path (generally 15 in RIP). This limits the diameter of a AS to a maximum of 15 hops.
- High signaling overhead - Periodic broadcasting of the distance vector table can result in increased utilization of the network resources for signaling.
- The algorithm is relatively slow to converge; you require information from all the nodes in the AS.