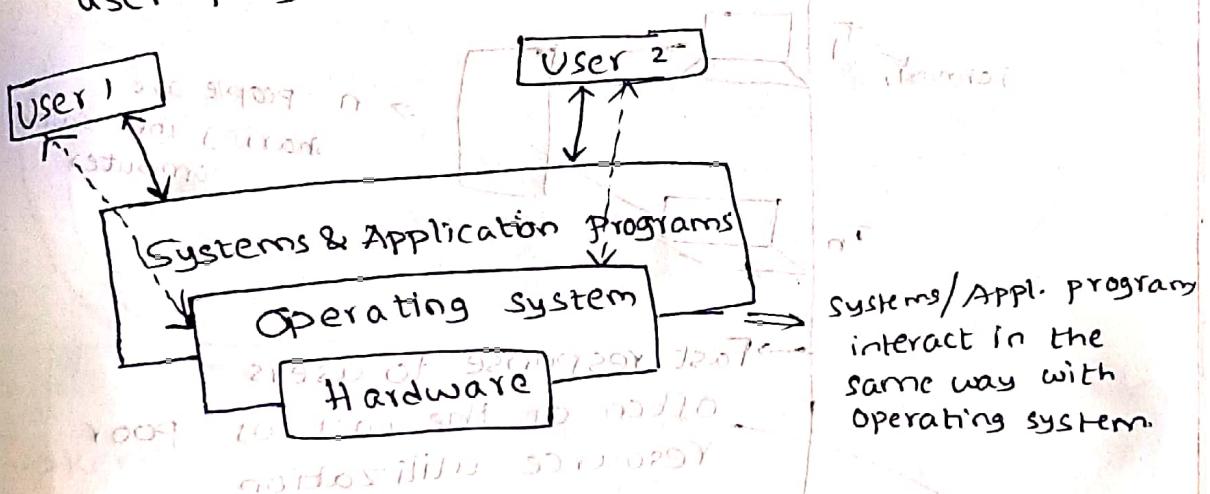


- Web link : <http://www.facweb.iltkgp.ac.in/~nisg/os>
- Books :
 - Operating System Concepts, 9th edition
Silberschatz, Galvin and Gagne

* What is an OS?

→ A program that acts as an interface between user program and computer hardware.



→ Systems programs → programs that assist the user in creating/running programs

Ex: Compiler, loader/linker, editor...

→ Application programs

Ex: Calculator, word processor...

→ Primary goals of an OS

→ Make a computer system easy to use.

→ Ensures efficient utilization of the resources

→ CPU, Time, Memory, files, I/O devices

* Evolution :-

→ Mini Computers and Mainframes

(1970-80)

Ex: IBM 1130

[Baudot code → in Punched Cards]

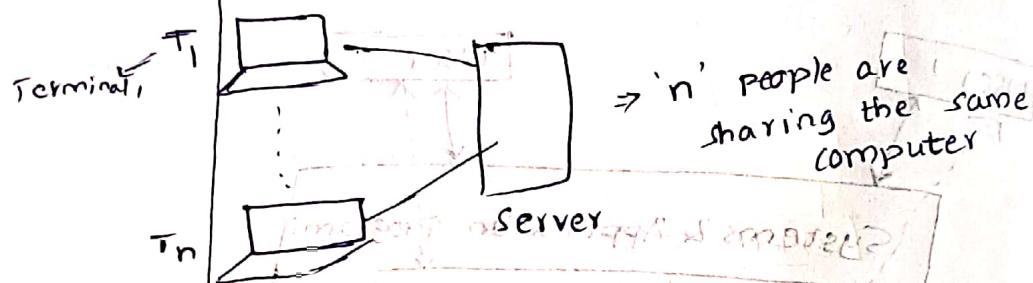
→ Users used to share a computer

→ Best utilization of resources

→ User experience has compromised

→ Workstations and Terminals

(1980- Present)



→ Fast response to users often at the cost of poor Resource utilization

→ ratifying global reliability of program

Ex: P_1 Read
10 min. → compute
↓ Print

→ while P_1 is in compute,
if P_2 is called, then P_1
has to saved & stopped

temporarily ⇒ more # registers
are used.

→ Desktops^(PCs) and Laptops (1985 - present)

→ Flexible user interface, easy to use

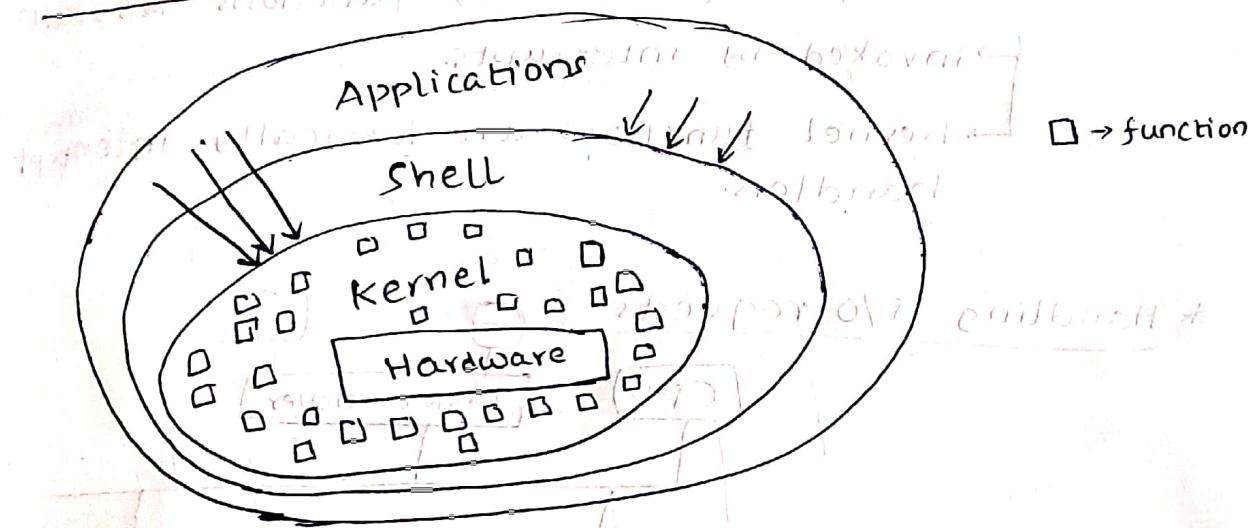
→ Single-user system (fully personal)

→ Mobile Systems

→ Constrained battery life

→ User experience

*A Typical structure of OS (w.r.t UNIX)

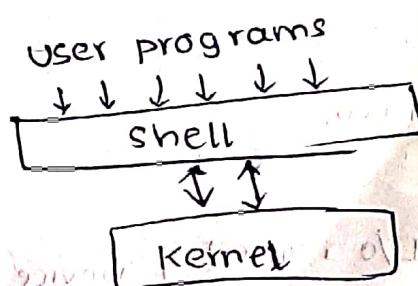


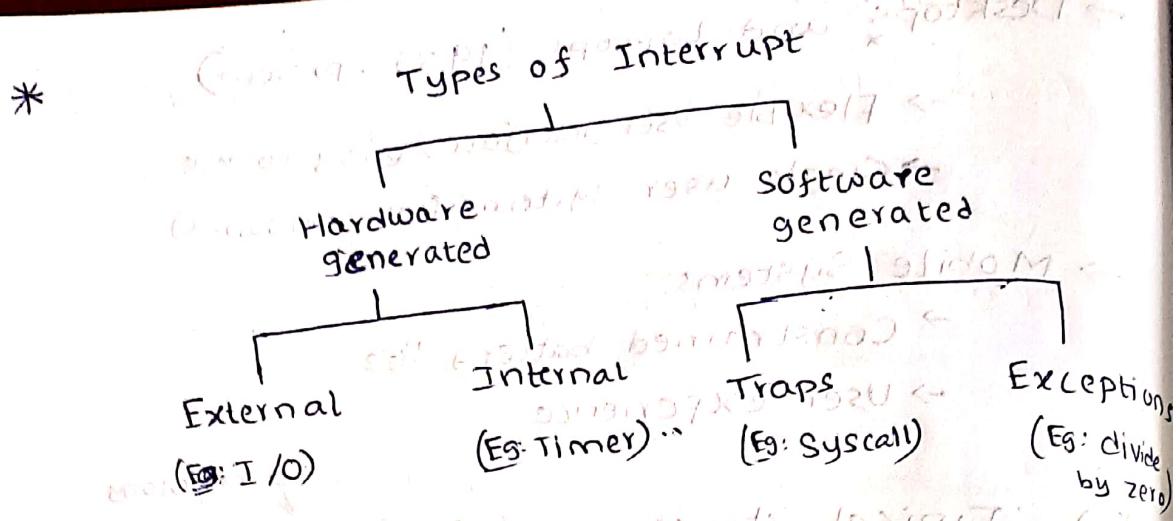
→ Kernel → A set of functions that provide the essential functionalities of the OS, and the associated data structures.

→ Kernel runs at all times → INCORRECT Statement

→ Shell → A program which allows user programs

to access the functions provided by the kernel.

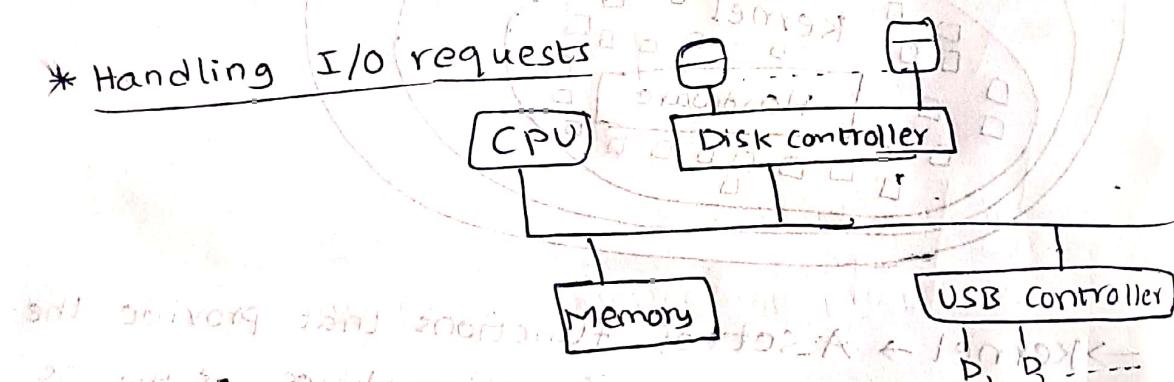




* How are the OS (i.e kernel) functions invoked?

- invoked by interrupts.
- Kernel functions are basically interrupt handlers.

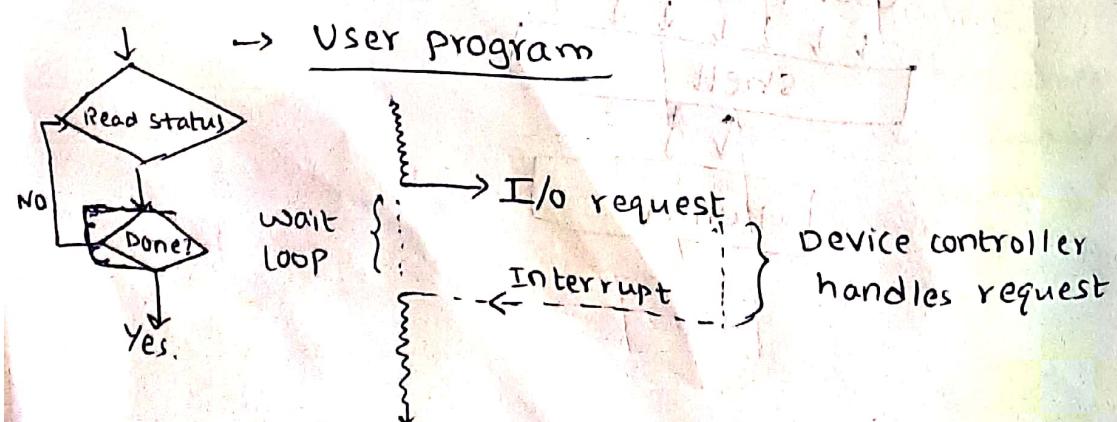
* Handling I/O requests



→ If there is no Disk controller, CPU has to spend so much time in disk access.

→ Disk controller → has a dedicated processor to carryout disk access and store the read data in a buffer.

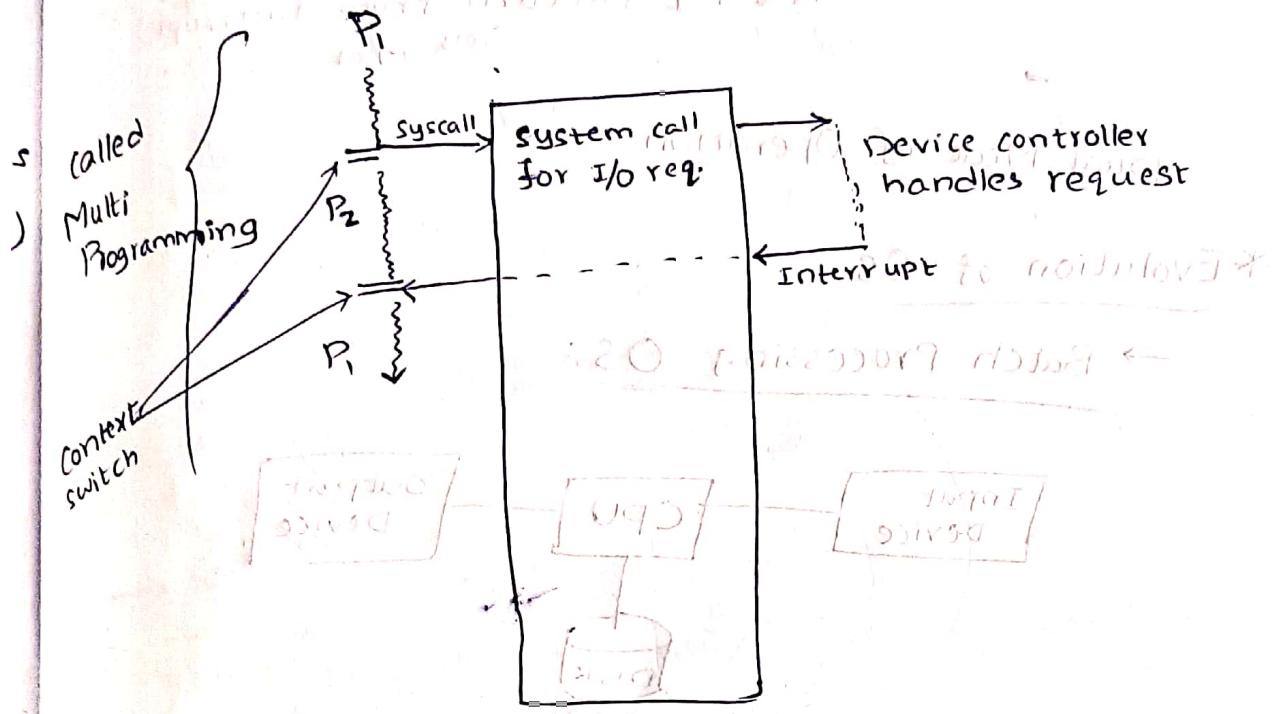
* Alternate :-



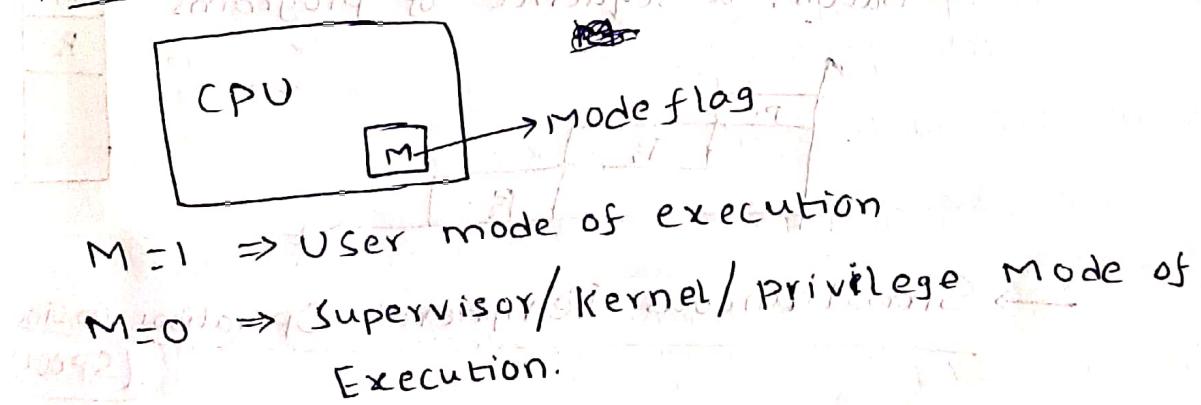
→ The above scheme is called ~~key~~ ~~process~~
"BUSY WAITING"

* Alternative 2 :-

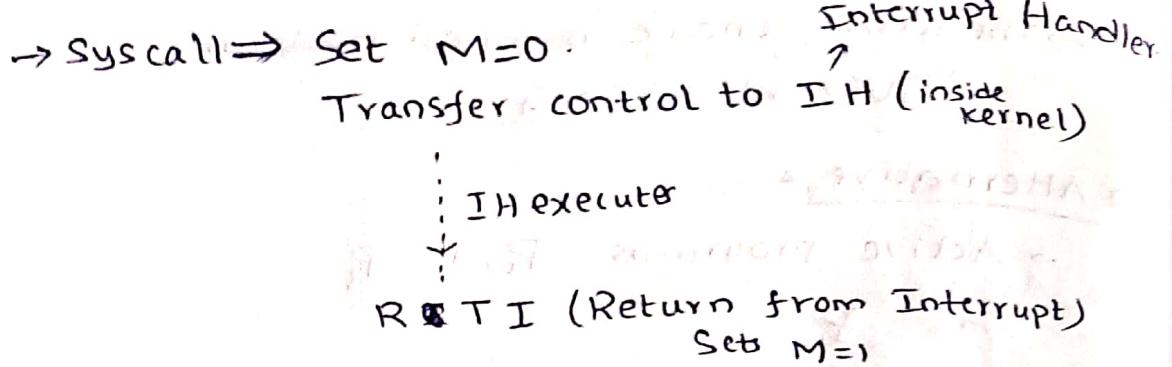
→ Active programs $P_1, P_2, P_3 \dots$



* Syscall



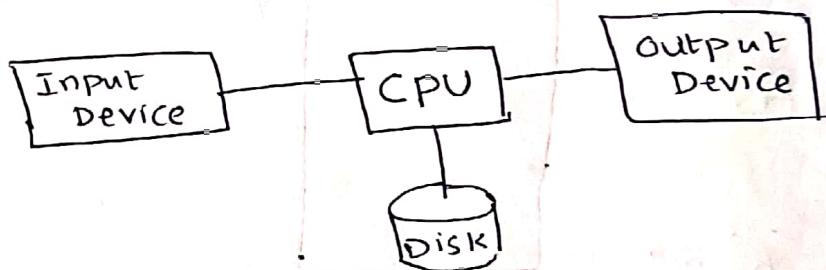
* → CPU can execute a privileged instr. only when $M=0$
↳ else it will generate an exception.



Dual-Mode of Operation.

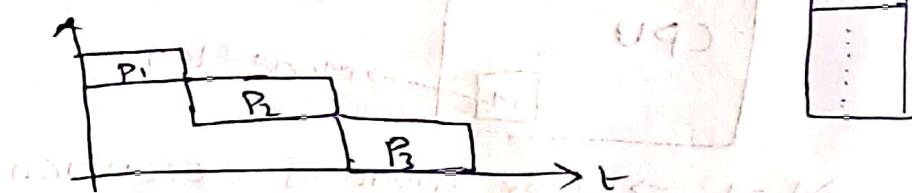
* Evolution of OS :-

→ Batch Processing OS :-

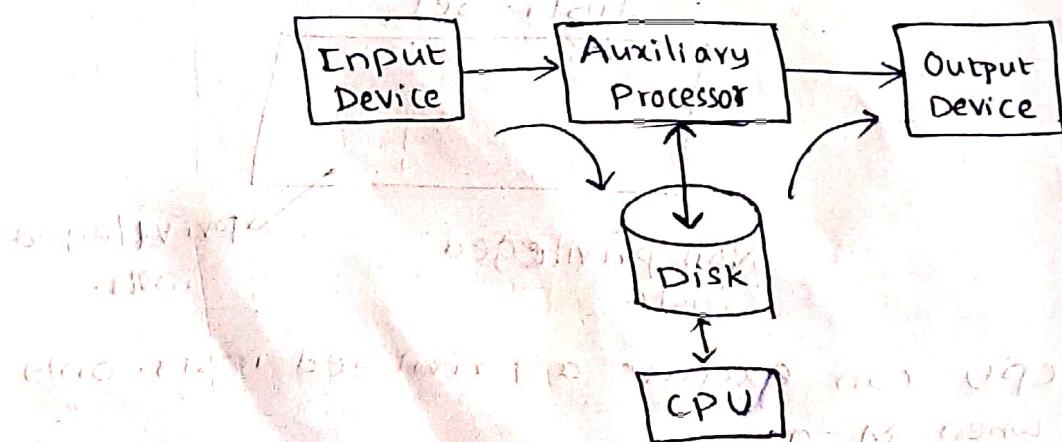


→ programs execute one at a time.

→ Batch: a sequence of programs

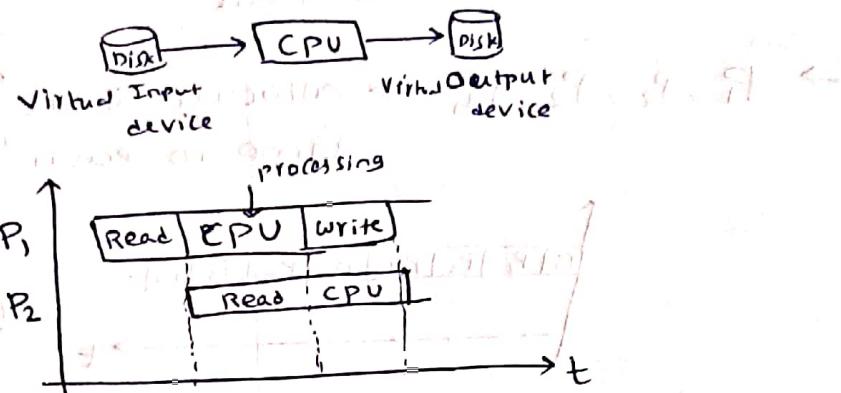


→ Simultaneous Peripheral Operations On Line (SPOOL)



- Auxiliary Processor
 - Just reads data from input device & ~~sends~~ sends to disk
 - OR reads data from disk and sends to output device

→ Concept of Virtual I/O devices



→ Multiprogrammed OS:-

- Multiple programs loaded in memory at the same time.
 - You can switch execution from one program to another.
 - Parallelism in CPU & I/O operation.
 - Concept of CPU-bound jobs & IO-bound jobs
- programs which do more computation than I/O.

Eg: Matrix multiply ($n \times n$)

Objective
Maximum resource utilization
by a good mix of CPU & IO-bound jobs

I/O: $O(n^2)$ \Rightarrow CPU-bound job
computation: $O(n^3)$

2. Payroll \rightarrow IO-bound job

→ Time sharing/Multi-tasking OS:-

- users sit on terminals & interact in real time.

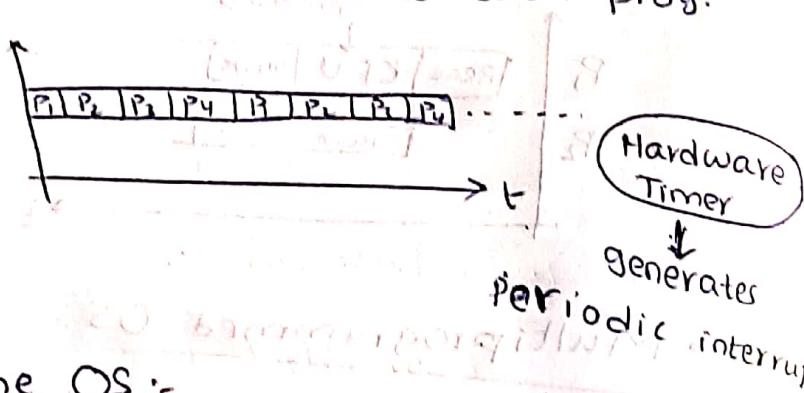
→ Objective → Maximum User satisfaction
minimization of response time

→ Very frequent context switching between programs

→ Let's say there are 50 users

if we follow Multiprogramming, one job is being done & remaining 49 users have to wait.

say, $P_1, P_2, P_3, P_4 \rightarrow$ allocate short bursts of time to each prog. $5-10\text{ ms}$



→ Real-time OS:

→ process - user requests based on deadlines

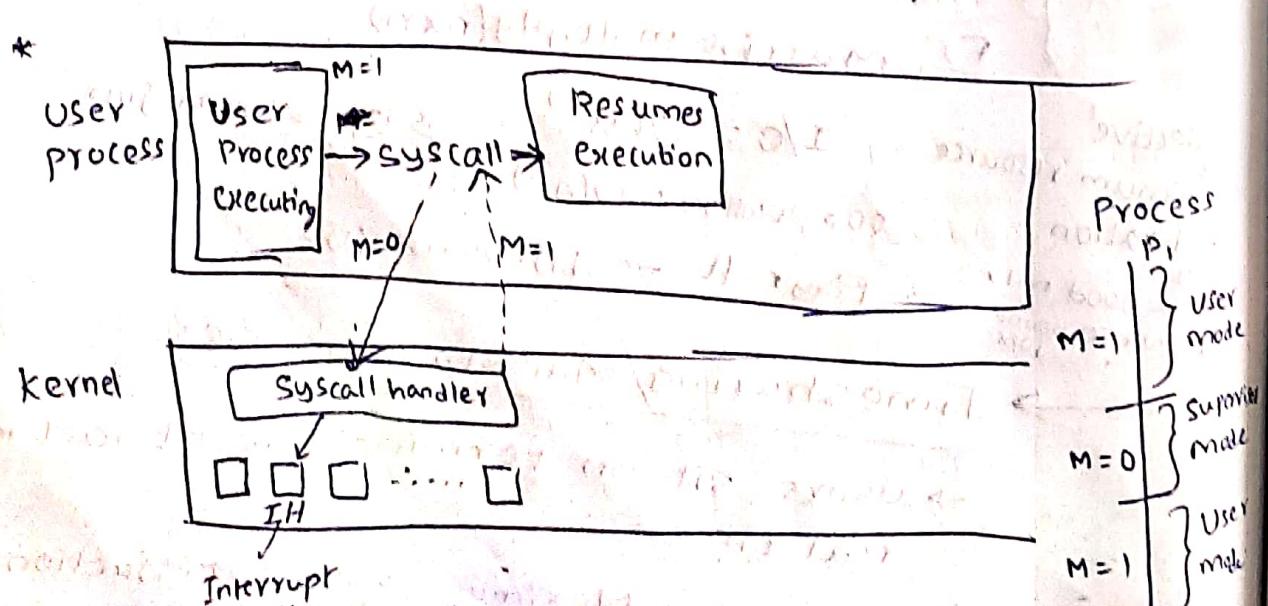
some work deadlines come as interrupts

additional interrupt

→ Distributed OS:

→ part of computation happens in one place and other part in other place.

* process → a program in execution



* System Calls

- It is a machine instr.
- different for different ISA.

Motorola
68000

MIPS
Syscall

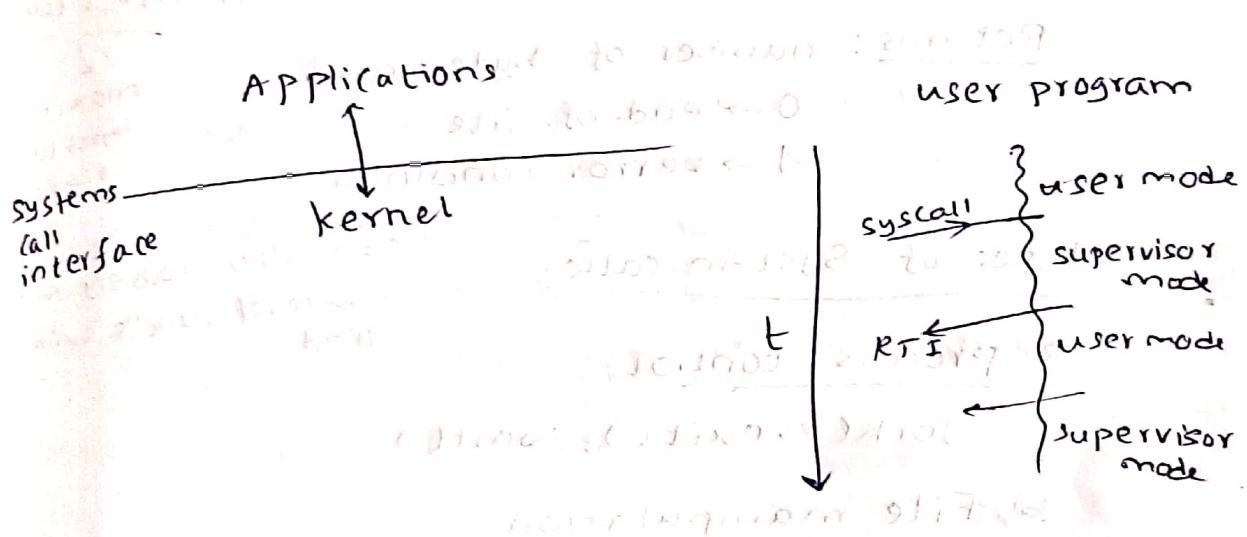
Intel
int n
n → parameter

ARM
SWI

$0 \leq n \leq 31 \rightarrow$ predefined

$n > 31 \rightarrow$ user defined

- system call instr. is executed
- software interrupt
- Mode = supervisor
- invoke interrupt handler
- Return from interrupt



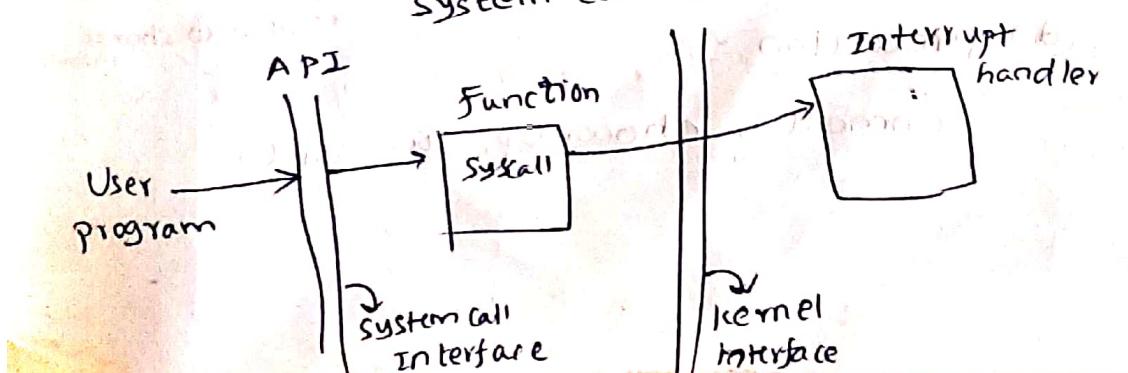
* System Call Interface:

- provided by the OS

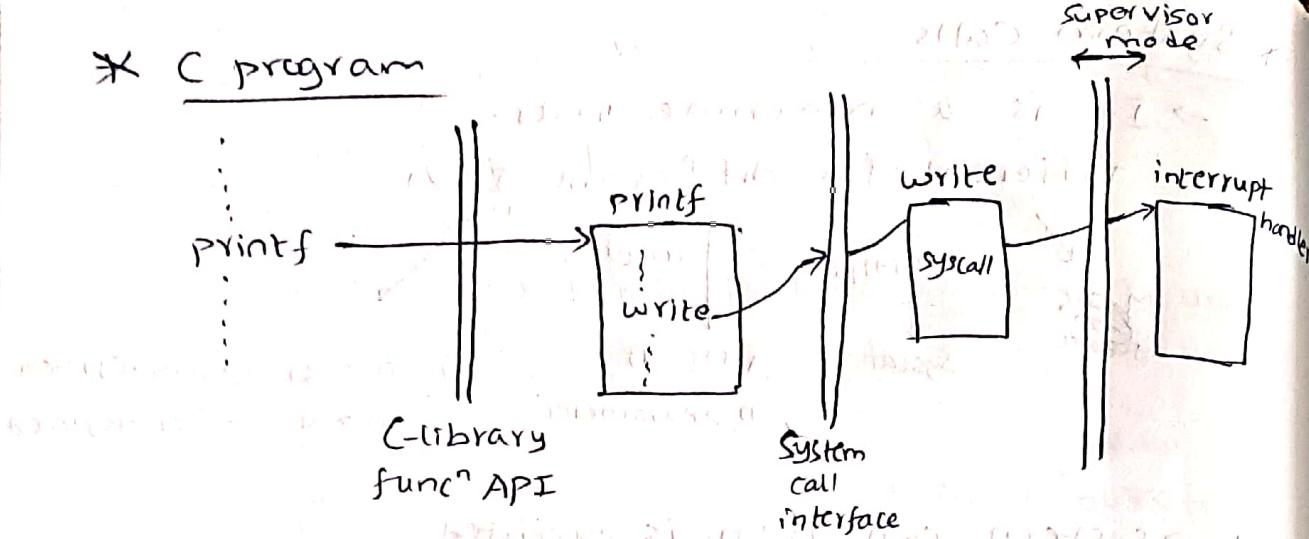
→ A set of low-level func's that can

be called from a user program.

→ These low-level func's execute some system call.



* C program



* Example API:

```
#include <unistd.h>
ssize_t read(int fd, void *buf, size_t count);
```

Returns: number of bytes read

0 → end-of-file

-1 → error condition

max. no. of bytes to read

→ Types of System calls:-

can also loosely mean low-level func's such as read.

(a) process control:

`fork()`, `exit()`, `wait()`.

(b) File manipulation:

`open()`, `close()`, `read()`, `write()`

(c) Inter-process communication: related to semaphores

`Pipe()`, `semget()`, `semop()`, `semctl()`,

`shmget()`, `shmat()`, `shmdt()`, `shmctl()`,

`kill()`, `signal()`.

(d) protection:

related to shared memory

`chmod()`, `chown()`, `umask()`

* Problems

1. Related to SPOOLING

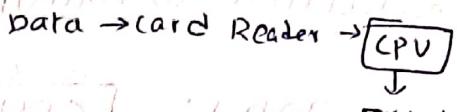
→ Payroll system (1000 employees)

→ Reading a card - 100 ms

→ Printing a line - 100 ms

→ Read/write from disk - 10 ms

→ processing time - 5 ms/employee



CPU

Printer

(a) No SPOOLING

$$\text{elapsed time} = 1000(100 + 5 + 100) \text{ ms}$$

only one line is to
be printed

$$= 205 \text{ s.}$$

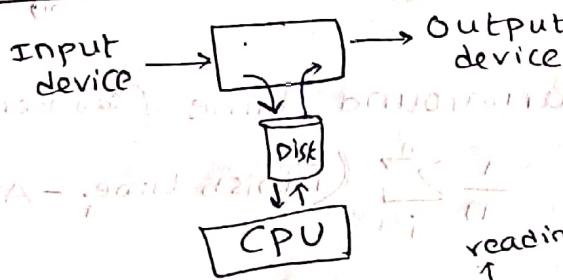
$$\text{CPU idle time} = 205 - \frac{1000 \times 5 \text{ ms}}{1000} = 200 \text{ s.}$$

card reader reads one card every second

$$(100 + 5 + 100) \text{ ms} \approx 5 \text{ cards/second}$$

$$\text{CPU utilisation} = \frac{5}{205} \times 100 = 2.5\%$$

(b) SPOOLING



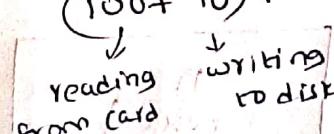
$$\text{Elapsed time} = 1000(10 + 5 + 10) \text{ ms}$$

writing to disk

$$\text{CPU idle time} = 25 - 1000 \times 5 = 20 \text{ s.}$$

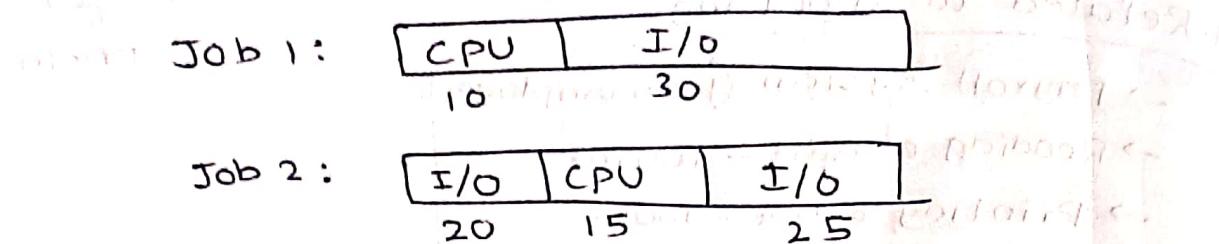
card reader reads one card every second

$$(100 + 10) \text{ ms} \approx 9 \text{ cards/second}$$

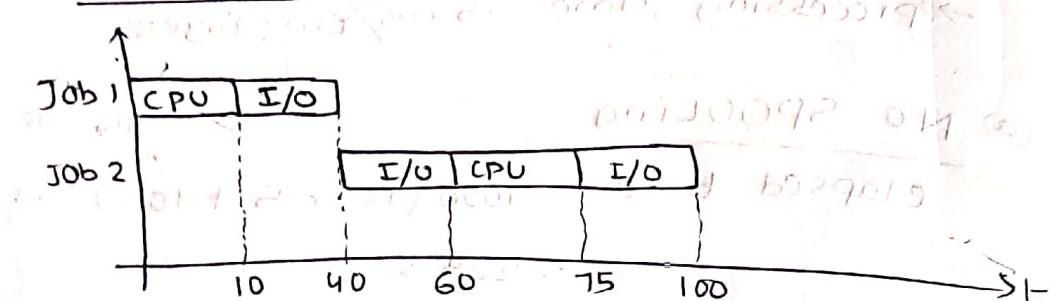


$$\text{CPU utilisation} = \frac{5}{25} \times 100 = 20\%$$

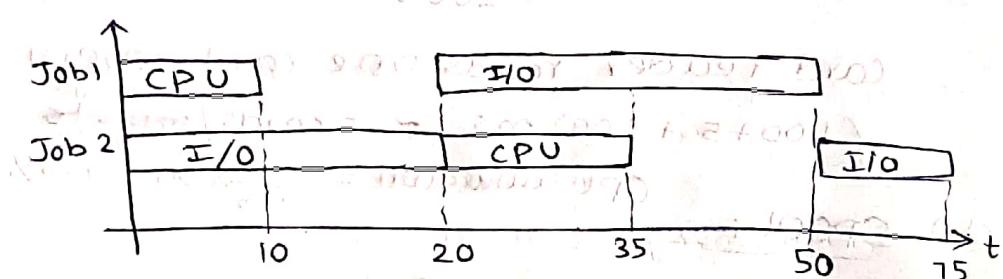
2. Related to multiprogramming:



(a) Batch execution:



(b) Multiprogramming:



→ Average turnaround time (say there are 'n' jobs)

$$= \frac{1}{n} \sum_{i=1}^n (\text{Finish time}_i - \text{Arrival time}_i)$$

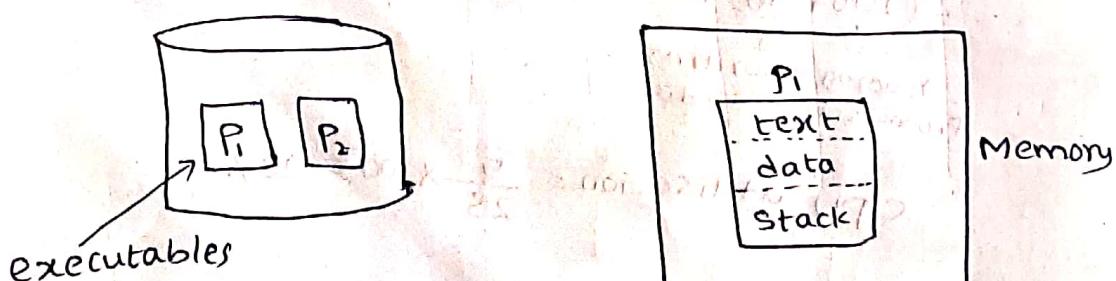
→ Average throughput

Throughput = No. of jobs completed per unit time

* process :-

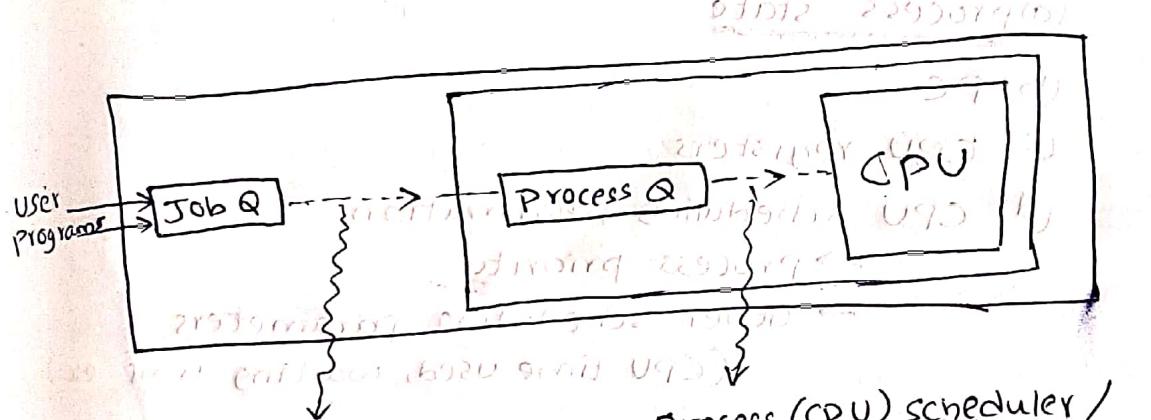
→ A program in execution

(terminologies of UNIX)



- A process consists of several regions
 - text → contains executable code
 - data → contains global data (static data)
 - stack → contains dynamic data (e.g. recursion)
 - heap → dynamic data controlled by the program.
(e.g. malloc).
- shared → Region shared by more than one process.

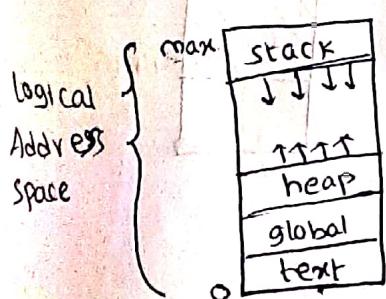
→ Batch-processing and multiprogramming
 ↳ user programs were called jobs



When the job gets loaded into memory, it becomes a process.

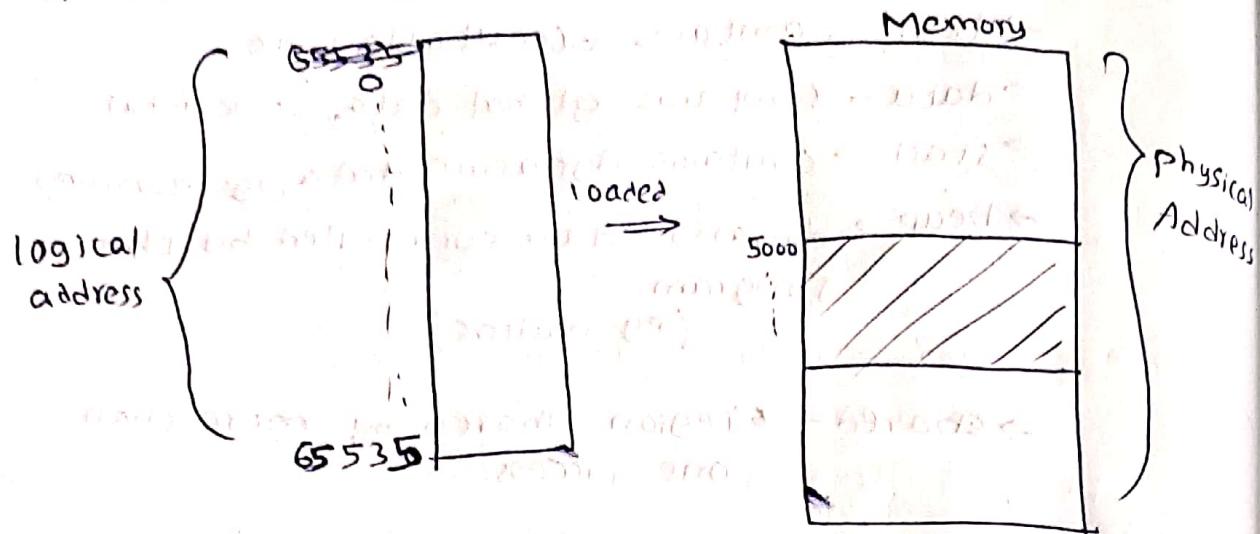
→ Structure of a process:

→ Each process is represented in the kernel using a data structure called Process Control Block (PCB).



Contains relevant information about a process.

* → Say Address - 16 bits



* TYPICAL PCB entries:-

(a) process state

(b) PC

(c) CPU registers

(d) CPU scheduling information

→ process priority

→ other scheduling parameters

(CPU time used, waiting time etc)

(e) Memory management information

→ Base and limit registers

→ Page table.

(f) Accounting information

→ Amount of CPU time used

→ Total waiting time

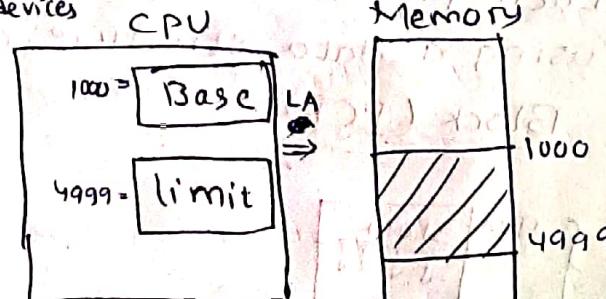
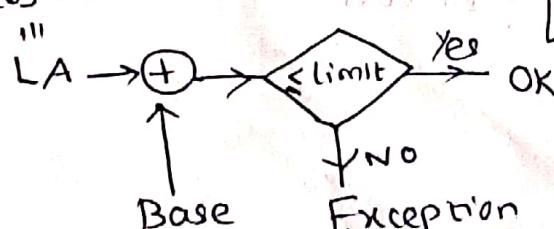
→ process id

(g) I/O status info

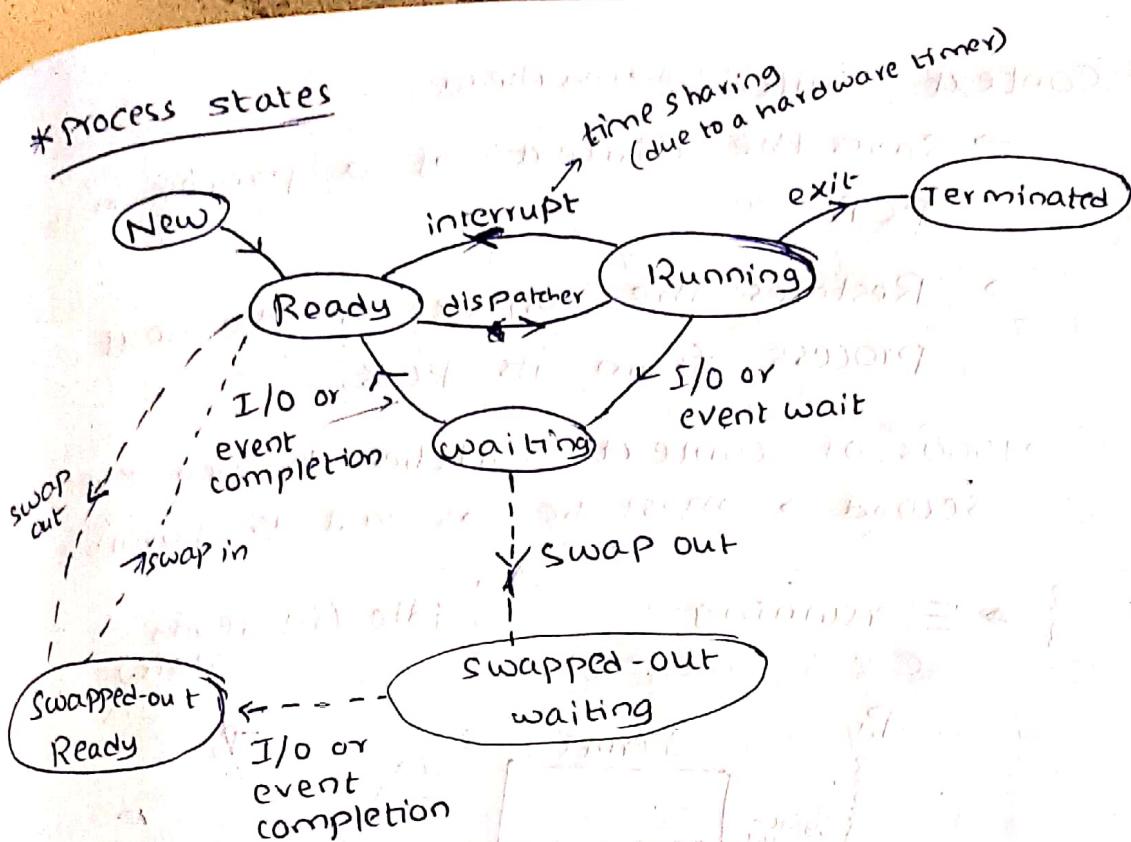
→ list of I/O devices being used

→ list of open files

logical address



* Process States



→ New: A process is being created

→ Ready: The process is ready to execute
(waiting for CPU).

→ Running: The process is running on the CPU.

→ waiting} Blocked: The process is waiting for some event to occur.
(I/O operation, signal from other process etc.).

→ Terminated: Finished execution and has left the system.

* Memory is limited

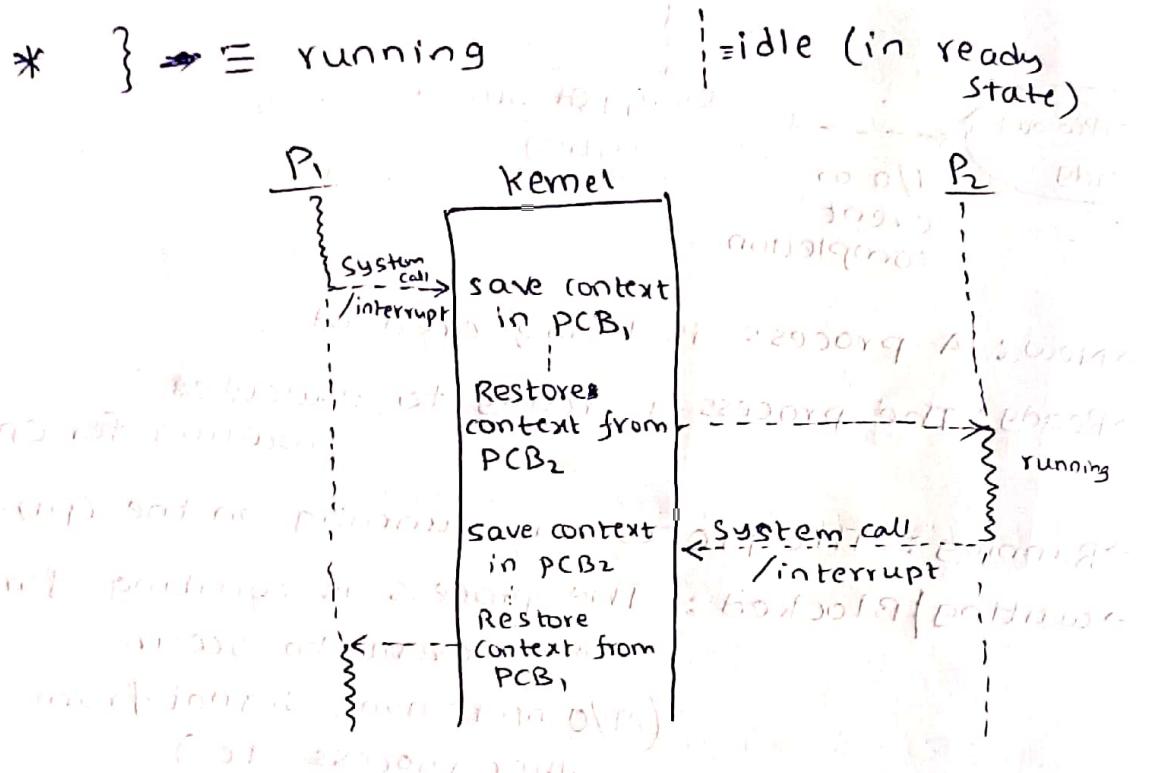
⇒ No. of processes that can be accommodated is limited.



Swap area

access to this is faster than the rest of the disk.

- Context switch → overhead
- Saves the context of a process in its PCB
- Restores the context of the next process from its PCB.
- 1000's of context switches happen every second ⇒ must be as fast as possible.



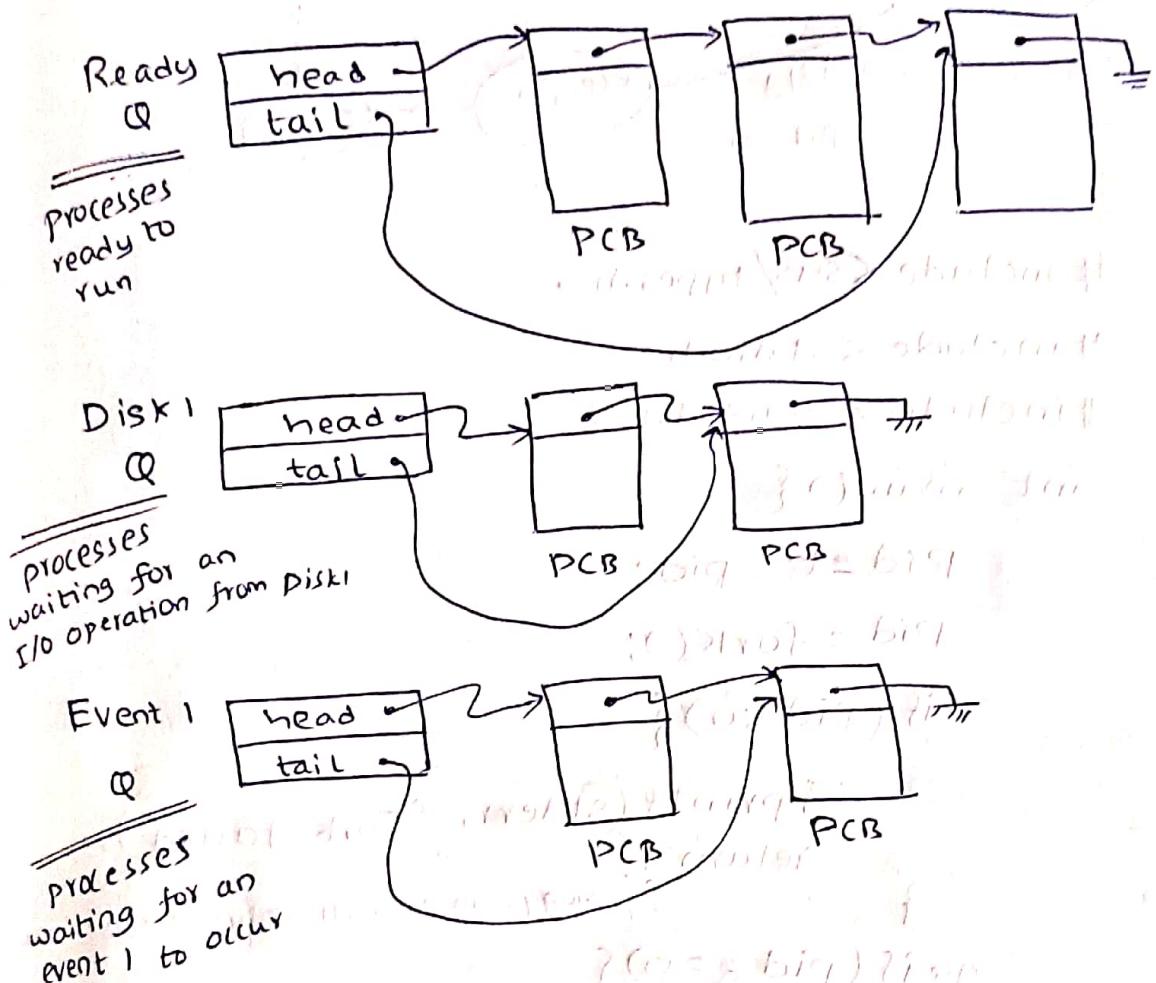
P_1 & P_2 are concurrent processes

* Process scheduling

→ Ready state → queue of processes
when the running process is interrupted,
there must be some algorithm to select a
process from the queue to make it running

* Process scheduler or dispatcher selects
a process from the ready queue for
execution whenever there is a context
switch.

- * One of the many scheduling queues maintained by kernel.
- linked list of PCBs



* Operations on Processes:

→ Process creation:-

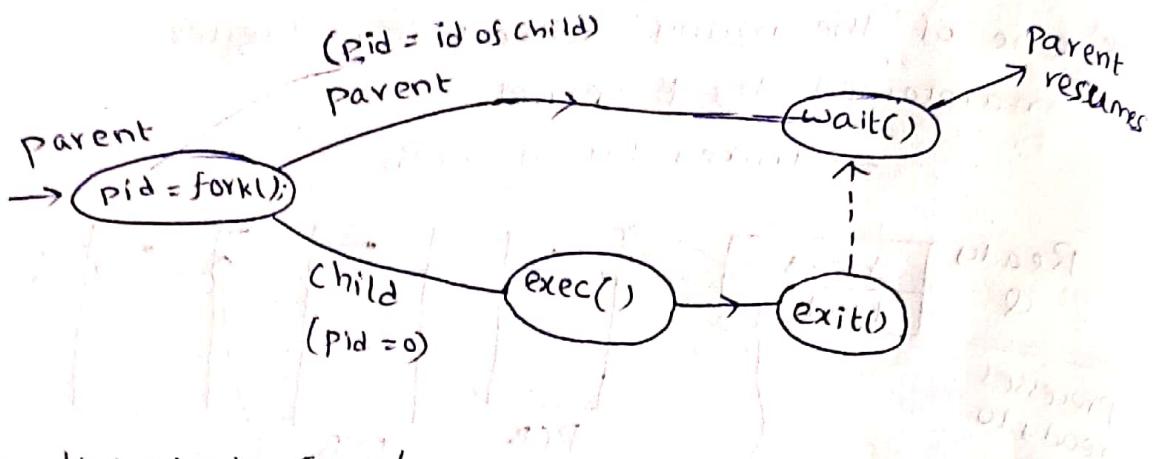
→ ps' command lists all the processes active.

→ every process has a unique id (pid).

→ A single process that gets created in the beginning → called "init" with pid=0

i.e when booted

→ all other processes are created using fork.



* #include <sys/types.h>

#include <stdio.h>

#include <unistd.h>

int main()

{ pid_t pid;

pid = fork();

if (pid < 0) {

 fprintf(stderr, "Fork failed");

 return 1;

} *non-zero return value ⇒ some error.*

else if (pid == 0) {

 execvp("/bin/ls", {"ls", NULL});

}

else {

 wait(NULL);

 printf("child complete");

}

return 0;

}

Process Termination:

→ The exit() and wait() system calls can be used.

child process

```
exit(2);
```

Parent process

```
pid_t pid;  
int status;
```

```
pid = wait(&status);
```

id. of child who has terminated
[∴ parent may have forked many processes, it can find which process terminated]

Value passed as argument to exit() in child.
i.e., Here it is 2.

Normal process termination

→ child executes exit()

→ parent executes wait()

If both these happen, the OS considers that it happened successfully & all entries are deleted from process tables.

→ A child process terminates but parent has not called wait().

⇒ The child becomes zombie

process
i.e., some entries will remain though it's not active

→ A child process terminates & parent also terminates but without calling wait()

⇒ The child process becomes orphan process

* All orphans are assigned to "init" as parent (i.e. init is parent) & init calls wait() periodically.

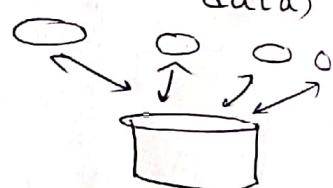
→ ∃ multiple process in various states of execution
↳ concurrent process

independent

A process that doesn't affect or be affected by other process running in system (can't share any data with other processes).

co-operative

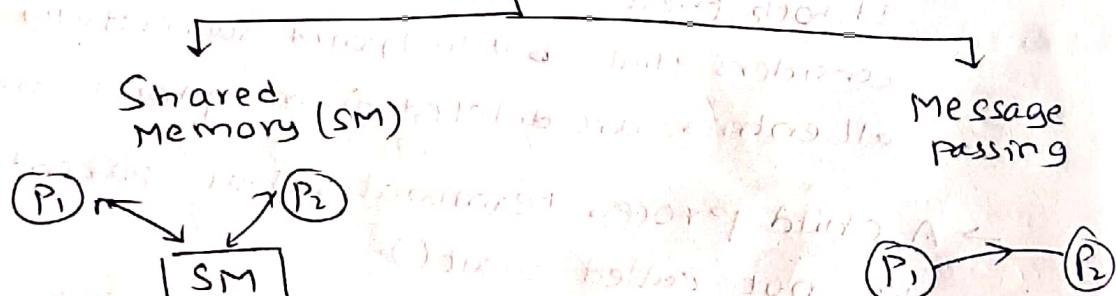
Process that are not dependent.
(share some data)



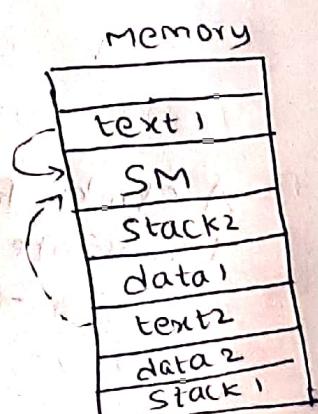
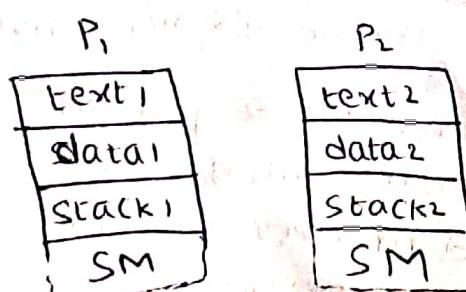
* Co-operating Processes :-

→ Some mechanism using which the processes can exchange data

Inter Process Communication (IPC)



(a) Shared Memory (SM)



System calls

- `shmget()` → creates a SM segment in Memory
- `shmat()` → attaches the SM segment to the address space of a process.
- `shmdt()` → detaches the SM segment from a process.
- `shmctl()` → removes the SM segment from memory

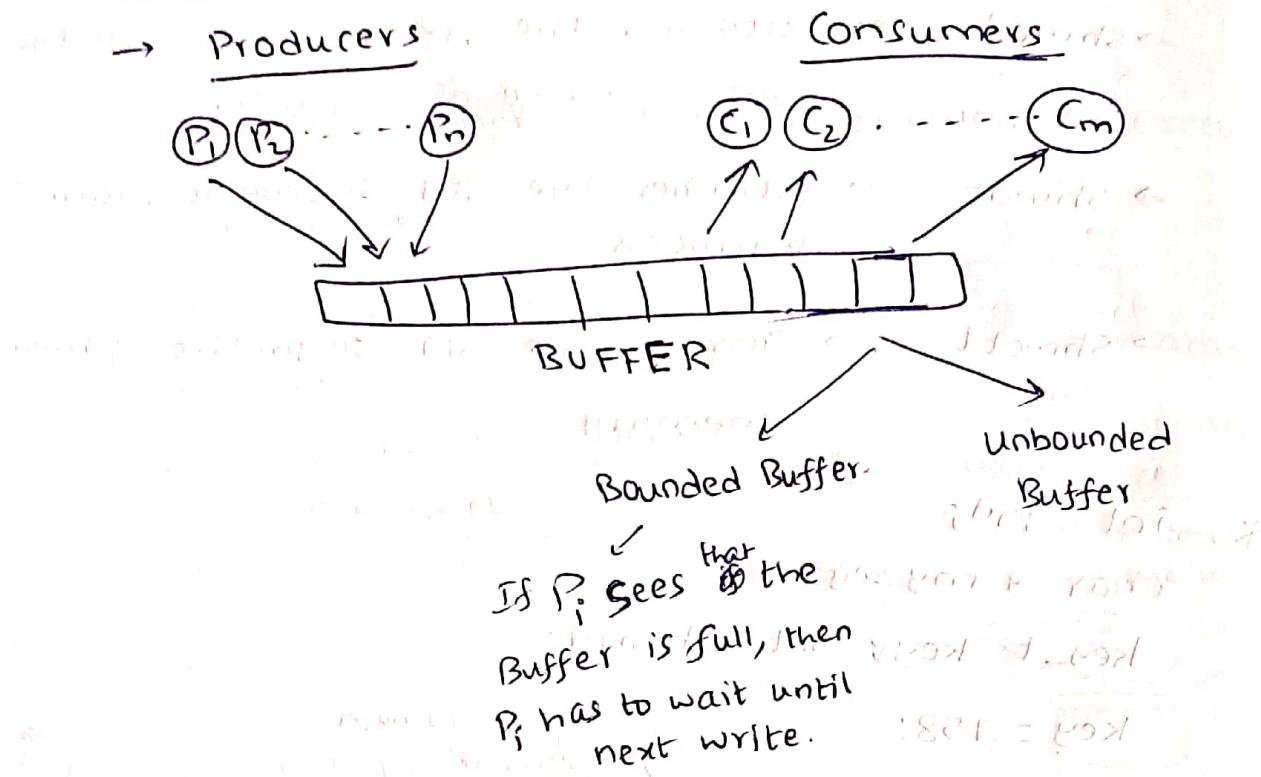
```
* int Pid;  
char * myseg; // user input  
key_t key; //int *shmid;  
key = 758; // user input  
shmid = shmget(key, 250, IPC_CREAT | 0666);  
pid = fork();  
if (pid == 0){  
    myseg = shmat(shmid, NULL, 0);  
    // user input  
    Shmdt(shmid);  
}  
else{  
    myseg = shmat(shmid, NULL, 0);  
    // user input  
    Shmdt(shmid);  
}  
Shmctl(shmid, IPC_RMID, NULL);
```

Diagram illustrating access permissions (octal) for the shared memory segment:

- 666 (110 110 110) is shown as three groups of three bits each.
- The first group (110) represents permissions for the owner (user).
- The second group (110) represents permissions for the group.
- The third group (110) represents permissions for others.
- Each bit corresponds to a permission: r (read), w (write), and x (execute). For example, the first bit (1) means r (read) for the owner.
- A note states: "every user has a group id also".
- A bracket labeled "Create a new SM segment with key if it does not exist" points to the `shmget` call.
- A bracket labeled "access permissions (octal)" points to the octal value 666.

* Example:-

Producer - Consumer Problem



Pseudo code

```
#define BUFF_SIZE 50
typedef struct {
    int value;
} item;
item buffer[BUFF_SIZE];
int in=0;
int out=0;
```

out = pointer to the first data item in the buffer
in = pointer to next free location in buffer
item = located in shared memory

producer

```
while(true){
```

 • < produce an item in next_P >

```
        while((int+1)%BUFF_SIZE == out);
```

```
            buffer[in] = next_P;
```

```
            in = (int+1)%BUFF_SIZE;
```

```
}
```

consumer

```

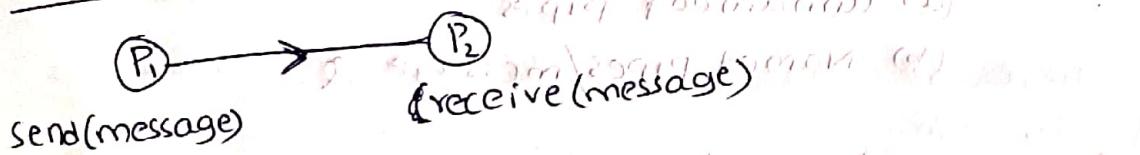
item = next_c;
while(true) {
    while(in == out)
        next_c = buffer[out];
    out = (out + 1) % BUFF_SIZE;
    <consume item in next_c>
}

```

Max data items that can be put in Buffer

$$= \text{BUFF_SIZE} - 1$$

(b) Message-Passing:-



Example:-

2 processes trying to update a variable is

SM.

SM [a=100]

P₁

a = a + 1

P₂

a = a - 1

in row level Lang	P ₁ P ₂
1. LOAD R ₁ , a 2. ADD R ₁ , R ₁ , 3. STORE R ₁ , a	4. LOAD R ₂ , a 5. SUB R ₂ , R ₃ , 6. STORE R ₂ , a

let's say st:1 was executing & there's a context switch

The sequence is like say

1 4 5 G 2 3 ⇒ Final value of a = 101

for 4 5 1 2 3 G 6 ⇒ " " = 99

The above is called RACE condition.

Also happens when ~~if~~ is given to S-R latch.

Example: Producer-Consumer Problem

<u>Producer</u>	<u>Consumer</u>
message next-p;	message next-c;
while(true){	while(true){
<produce item in nextP>	receive(next-c);
send(<u>next_P</u>);	<consume item in next-c>
}	}

→ Two approaches:

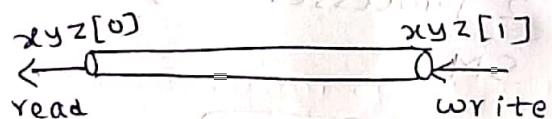
(a) unnamed pipes

(b) Named Pipes/message Q.

* Unnamed pipe:-

→ Pipe() system call.

```
int xyz[2];
pipe(xyz);
```



```
pid = fork();
```

```
if (pid == 0) {
```

```
    write(xyz[1], ...)
```

```
}
```

```
else {
```

```
    read(xyz[0], ...)
```

```
}
```

→ Parent & child can communicate
since they have access to both
ends of the pipe.

→ unnamed pipe can be used by processes that are generated by a common Parent.

* Named Pipe:

→ They use a common temporary file.

char *myfifo = "/tmp/myfifo"; files in /tmp/ are accessible to every user

mkfifo(myfifo, 0666);

more
parent

```
fd = open(myfifo, O_WRONLY);
write(fd, ...);
close(fd);
```

child

```
fd1 = open(myfifo, O_RDONLY);
read(fd1, ...);
close(fd1);
```

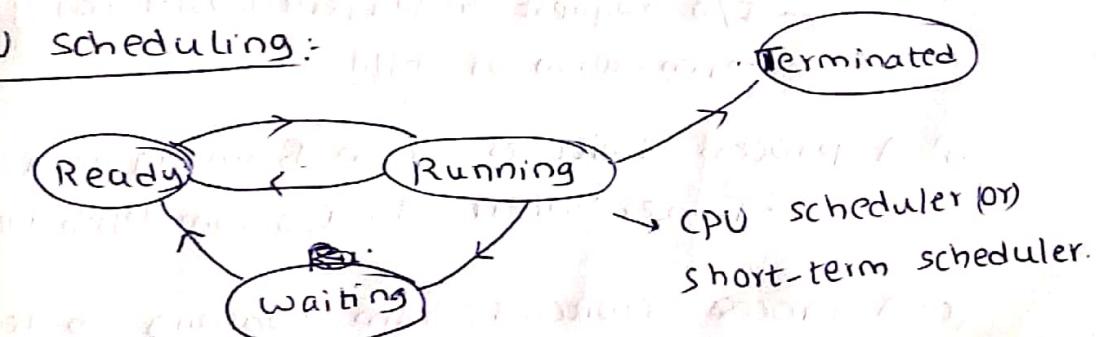
→ If even one of the processes calls mkfifo, then all the other processes can use it as named pipe.

(c) Process Comm" across machines

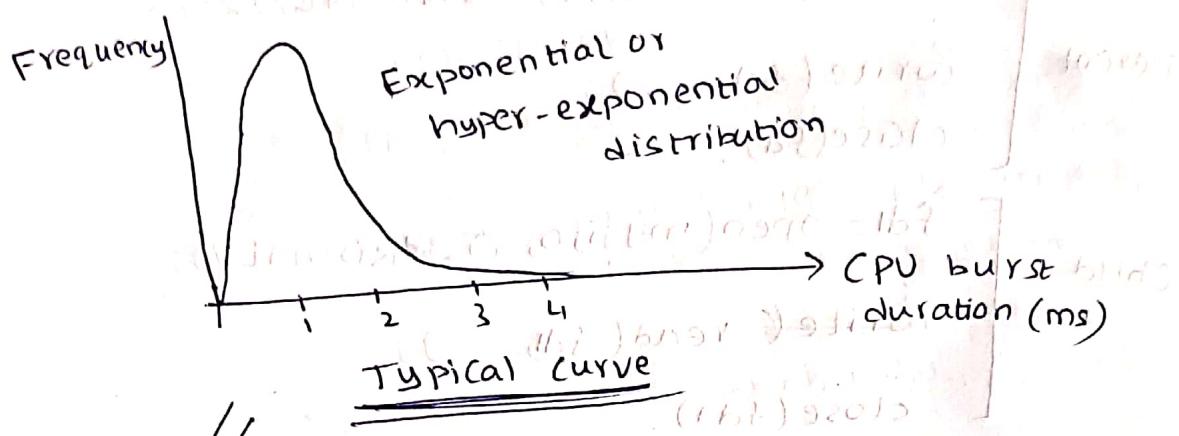
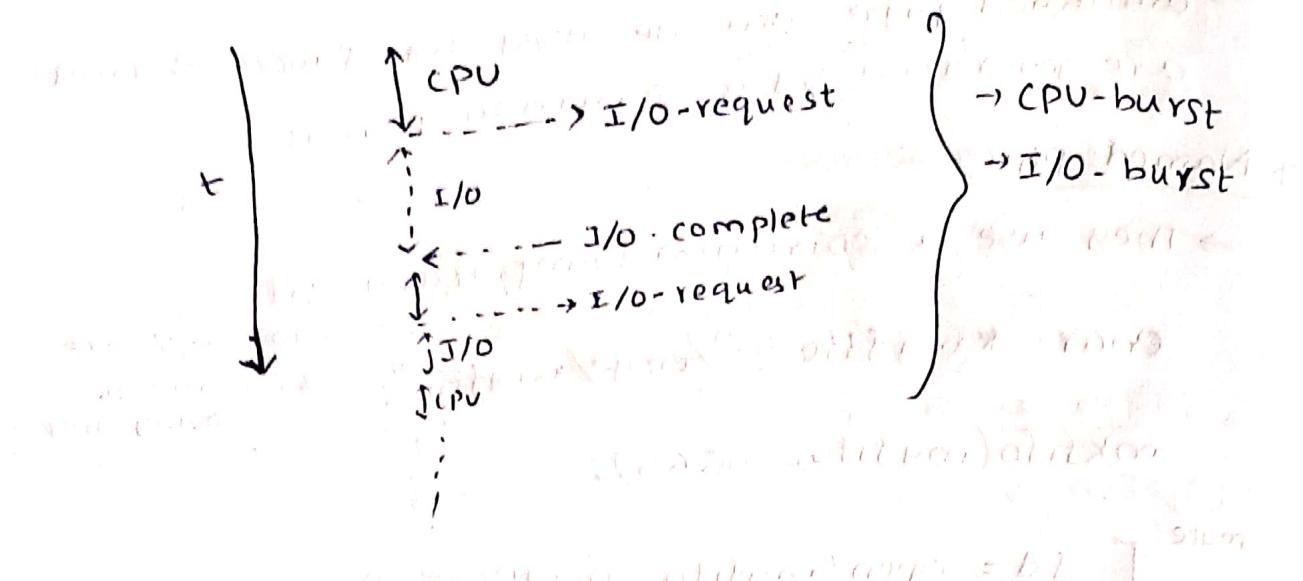
→ Remote procedure call

→ Socket communication

* CPU Scheduling:



* Typical process execution:



There are large no. of CPU bursts with very small duration & small no. of CPU bursts with large duration.

* Scheduling decisions

(a) A process switches from running to waiting

- I/O request (a) Parent invokes wait() for termination of child.

(b) A process switches from Running to Ready.

- Timer interrupt, I/O-completion interrupt

(c) A process switches from waiting to ready.

- I/O completion interrupt, child has terminated.

then, parent switches to ready.

(d) A process terminates

- exit() system call.

→ (a) and (d)

→ Non-Preemptive scheduling

↳ Allowing a process to continue for as long as it wants.

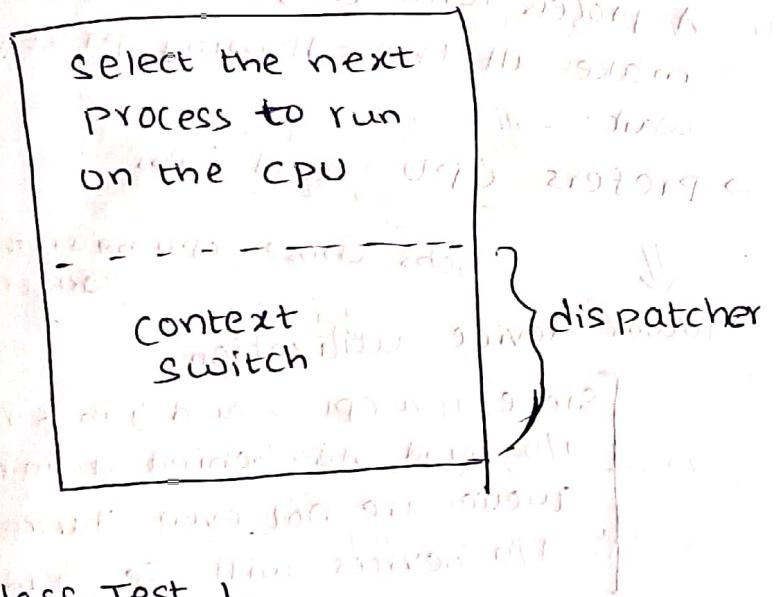
Here, the processes are coming out of CPU by their own.

→ (b) & (c)

→ Preemptive scheduling

↳ forcibly taking the CPU away from a process.

* CPU Scheduler



* Class Test 1

7th February 7-8 PM

Venue:- CSE Dept

Syllabus:- Up to CPU scheduling

* Scheduling Criteria

→ (a) CPU utilization

→ (b) Throughput - No. of processes finishing per unit time.

→ (c) Turnaround time - (Finishing time - Arrival time)

→ (d) Waiting time - Time taken in the ready and waiting states.

→ (e) Response time - after how much time a process receives a response.

* CPU Scheduling Algorithms

→ (a) First-Come First Serve (FCFS)

→ It is a non-preemptive algorithm.

→ Ready list is maintained as a FIFO queue.

* Drawback:

→ Convoy effect

↓
Traffic has to wait when convoy passes.

↓
A process with large CPU burst will make all the following processes wait.

→ prefers CPU-bound jobs

↓
Jobs whose CPU-bursts are longer.

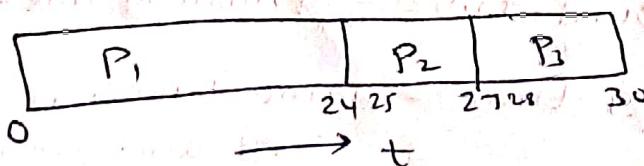
↓
lower device utilization.

Since, if a CPU-bound job is running I/O-bound jobs behind it in the queue are not even started, then I/O-devices will lie idle.

Ex:-

Processes	P ₁	P ₂	P ₃
Arrival Time	0	0	0
CPU Burst	24ms	3ms	3ms

(a) P₁-P₂-P₃ sequence



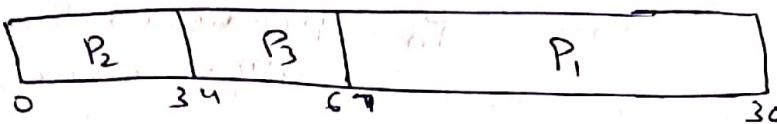
$$\text{Avg. turnaround time} = \frac{(24-0) + (27-0) + (30-0)}{3}$$

$$= 27 \text{ ms}$$

$$\text{Avg. waiting time} = \frac{0 + (27-3) + (30-3)}{3}$$

$$= 17 \text{ ms}$$

(b) $P_2 \rightarrow P_3 \rightarrow P_1$



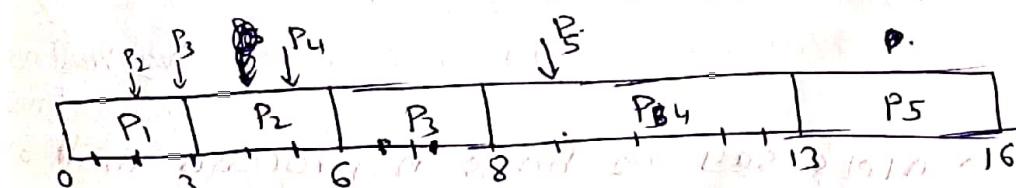
$$\text{Avg. turnaround time} = \frac{(30-0) + (6-0) + (3-0)}{3}$$

$$= 13 \text{ ms}$$

$$\text{AWT} = \frac{(30-24) + (3-3) + (6-3)}{3} = 2 \text{ ms}$$

Ex2:

Process	P_1	P_2	P_3	P_4	P_5
Arrival time (in ms)	0	2	3	5	8
CPU Burst	3	3	2	5	3



$$\text{AWT} = \frac{0 + 1 + 3 + 3 + 4}{5} = 2.2 \text{ ms}$$

(b) Shortest Job First/Next (SJF/SJN)

→ The process with the smallest next CPU burst is allocated the CPU.
↳ non-preemptive.

→ Predicting the next CPU burst time.
→ exponential averaging

t_n - length of n^{th} CPU burst

T_n - predicted length of n^{th} CPU burst

based on history
then, $T_{n+1} = \alpha t_n + (1-\alpha) T_n$,
where $0 \leq \alpha \leq 1$

$$\Rightarrow T_{n+1} = \alpha t_n + (1-\alpha)t_{n-1} + (1-\alpha)\alpha t_{n-2}$$

$$\dots + (1-\alpha)^{n-1} t_1 + (1-\alpha)^n T_0$$

initial value

$$\alpha=0 \Rightarrow T_{n+1} = T_n = T_0$$

→ recent history doesn't have any effect.

$$\alpha=1 \Rightarrow T_{n+1} = t_n$$

→ only recent CPU burst matters

→ The SJF algorithm minimizes the AWT

Avg-Waiting Time.

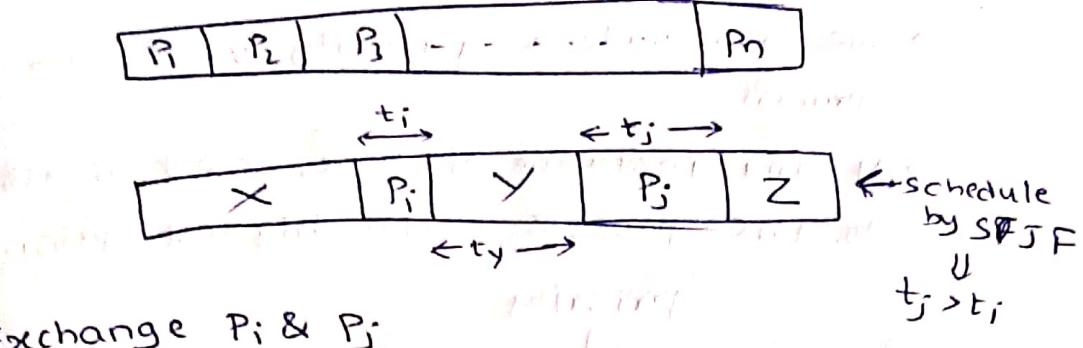
→ Let's say we have "n" processes that all arrive at $t=0$.

CPU burst times are $t_1 \leq t_2 \leq \dots \leq t_n$.
 P_1, P_2, \dots, P_n

Proof that in this case, SJF minimizes AWT.

Proof by contradiction

→ SJF Schedule



Exchange P_i & P_j:

→ X: no change in waiting time

→ Z: no change in waiting time

→ Y: All the processes will have their waiting times increased.

P_j: Waiting time decreases by t_y + t_i

P_j: Waiting time increases by t_y.

(c) Shortest Remaining Time First (SRT):

→ Preemptive version of SJN

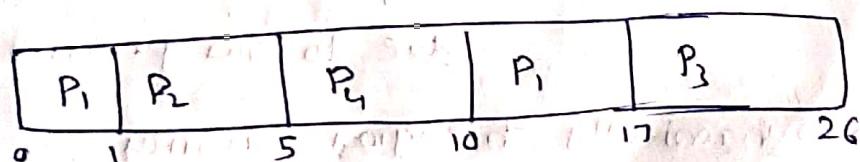
↳ A process with a shorter CPU burst can preempt a running process.

→ We decide when

↳ A new process arrives

Ex:

Process	P ₁	P ₂	P ₃	P ₄
Arrival time	0	1	2	3
CPU Burst	8	4	9	5



$$AWT = \frac{9+0+15+2}{4} = 6.5$$

(d) Priority Scheduling

→ We associate a priority value with each process

→ At any instance of time, CPU is assigned to the process with highest priority

Priority

Statically assigned

Dynamically assigned

→ Static Priority

→ Many versions of Linux

→ Each process is assigned a priority value

-20 → highest priority

+19 → ~~lowest~~ lowest priority

Default priority = 10.

→ Commands

→ nice → decreases the priority of a process

→ renice → increases " "

(only with "root" access)

→ Problem

Starvation or indefinite waiting

A process is waiting indefinitely due to low priority

→ Dynamically changing priority

→ Try to overcome starvation problem.

As a process waits, its priority increases

→ Response Ratio

$$RR = \frac{\text{Time since arrival} + \text{CPU burst time}}{\text{CPU burst time}}$$

Jobs arriving first, but are being delayed have high RR.

→ RR can indicate priority

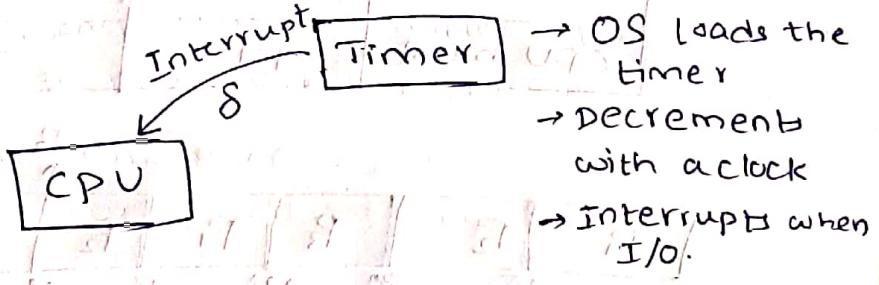
→ whenever there's a clash of priority, we can choose the one with highest RR

(e) Round-robin Scheduling

→ Designed primarily for interactive

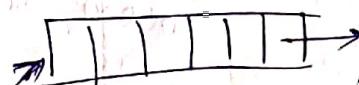
time-stationary systems

→ Let's say a small time quantum, δ is predefined ($\delta \approx 1-100 \text{ ms}$)

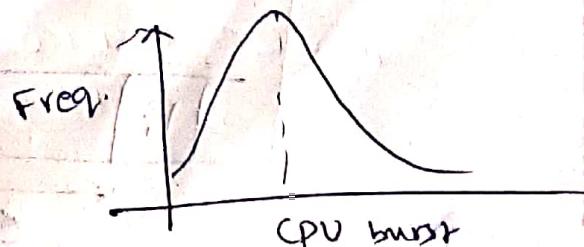


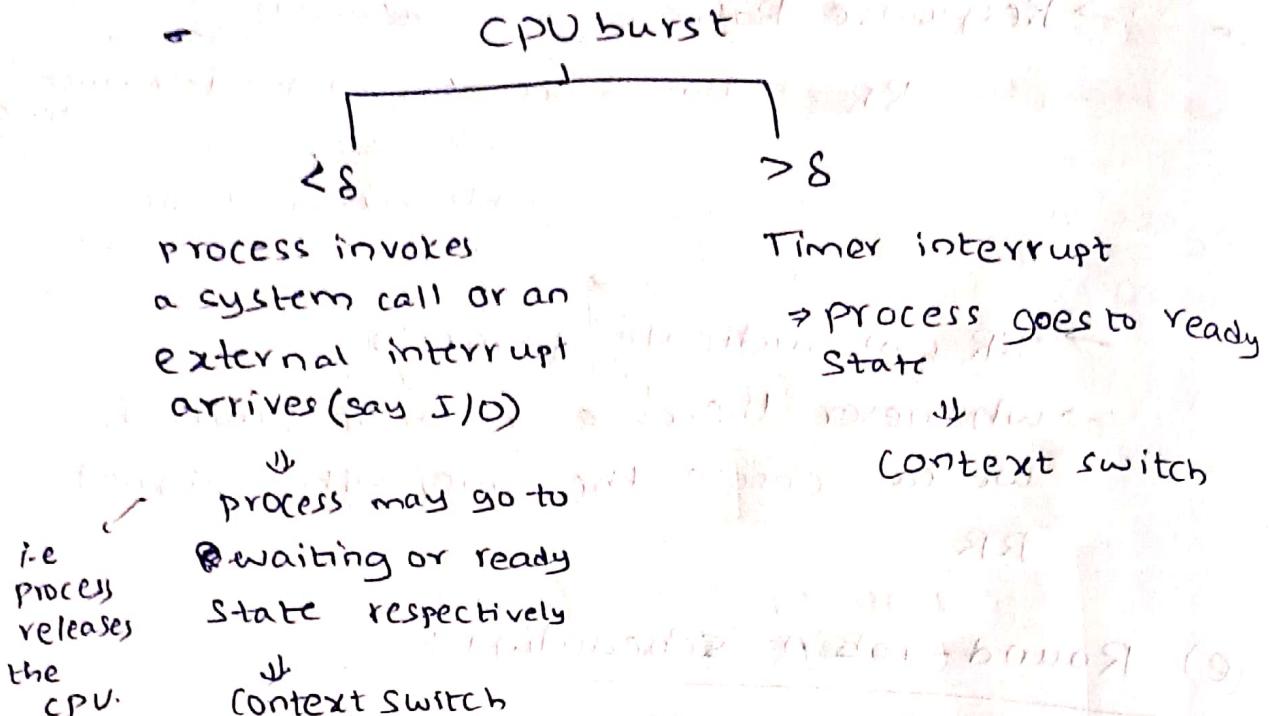
Ready list

→ stored as circular queue.



→ The instruction to set the timer is privileged instruction.

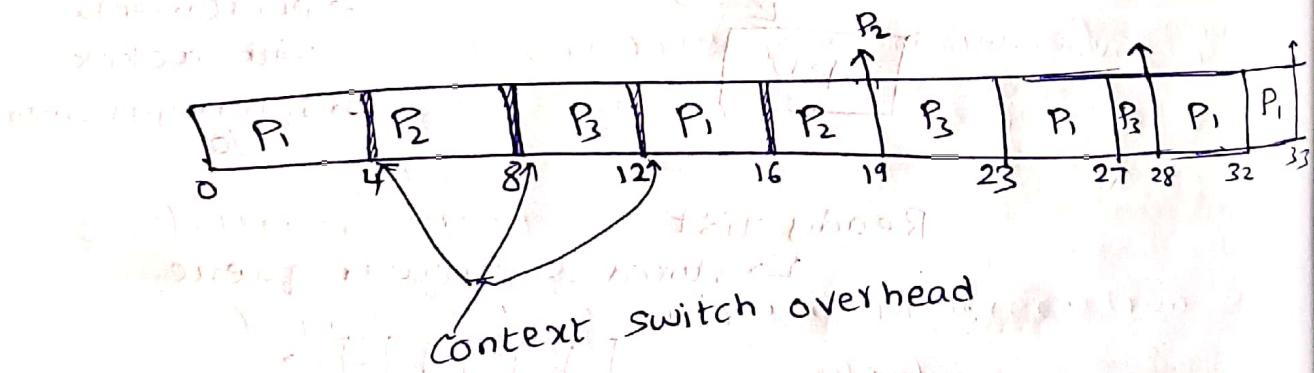




→ Example:

$$\delta = 4 \text{ ms}$$

Process	P ₁	P ₂	P ₃
Arrival Time	0	0	0
CPU burst	17	7	9

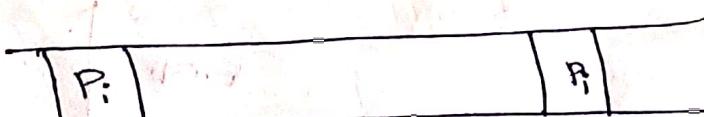


→ A Simple Analysis

→ n processes in READY queue.

→ δ is the time quantum.

→ σ is the context switch overhead.



$(n-1)(\delta + \sigma) \leftarrow$ maximum waiting time of a process before it gets back to CPU.

Very High $\delta \Rightarrow$ similar to FCFS.
 very small $\delta \Rightarrow$ large no. of context switches \Rightarrow overhead keeps adding up.

$\delta > \overline{G}$
 typically 10 ms
 typically 10 ms

Q.	Process	P ₁	P ₂	P ₃	P ₄
Arrival time		0	0	0	0
CPU Burst		6	3	1	7

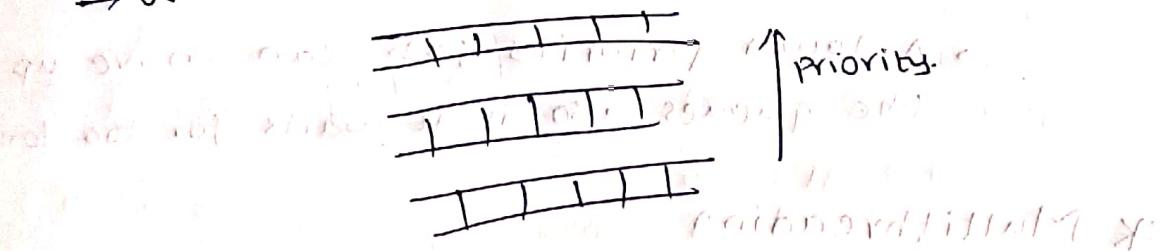
compute AWT for $\delta = 1, 2, 3, 4, 5, 6, 7$.

Ans. we get graph like



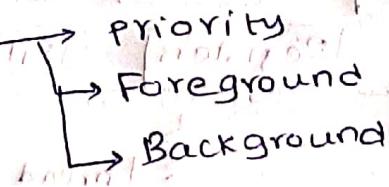
(f) Multi-level queue scheduling

→ We'll have multiple Ready queues



→ Multiple ready queues with different priorities and scheduling policies.

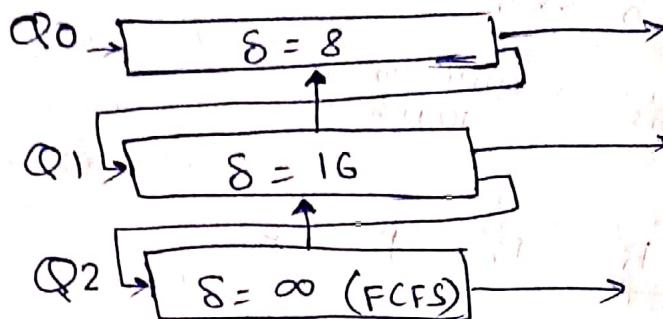
→ A process is assigned to one of the queues



→ processes do not move across queues.
 ↳ It can lead to starvation.

(9) Multilevel feedback queue scheduling

→ 3 queues



→ Q_0 has the highest priority & Q_2 has the lowest priority.

→ A new process is assigned to Q_0 .

→ Current running process is in Q_0

→ & ~~if~~ CPU burst < 8 ms

Then, the process remains in Q_0 .

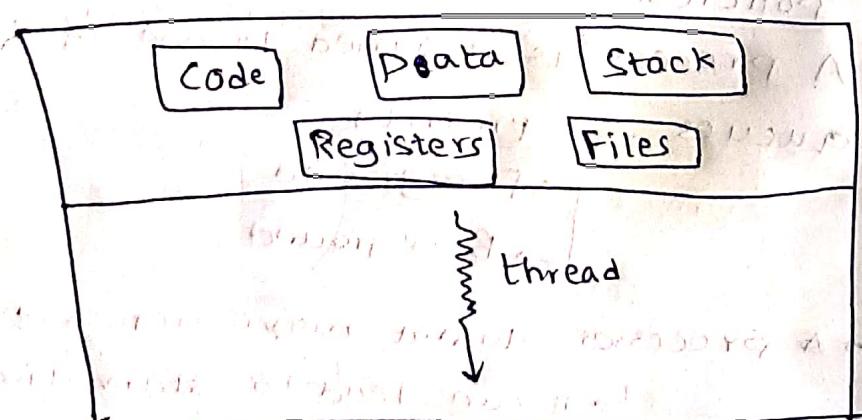
if CPU burst > 8 ms

Move the process to Q_1 :

→ A lower priority job can move up the queues also if it waits for too long.

* Multithreading

→ A process is a program in execution with a single thread of control.

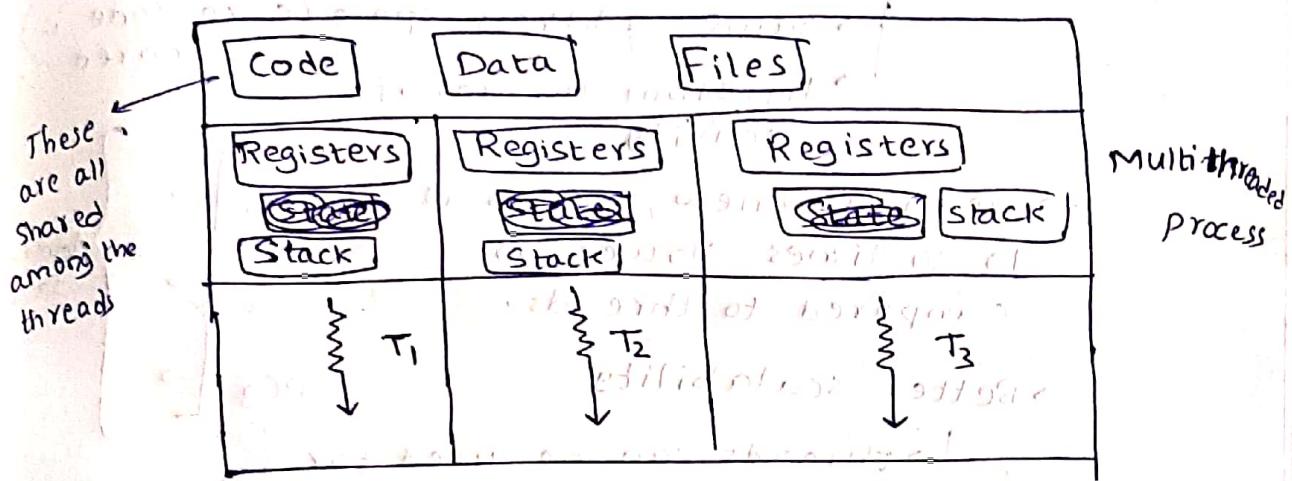


→ What is a thread?

→ Also referred to as light-weight process

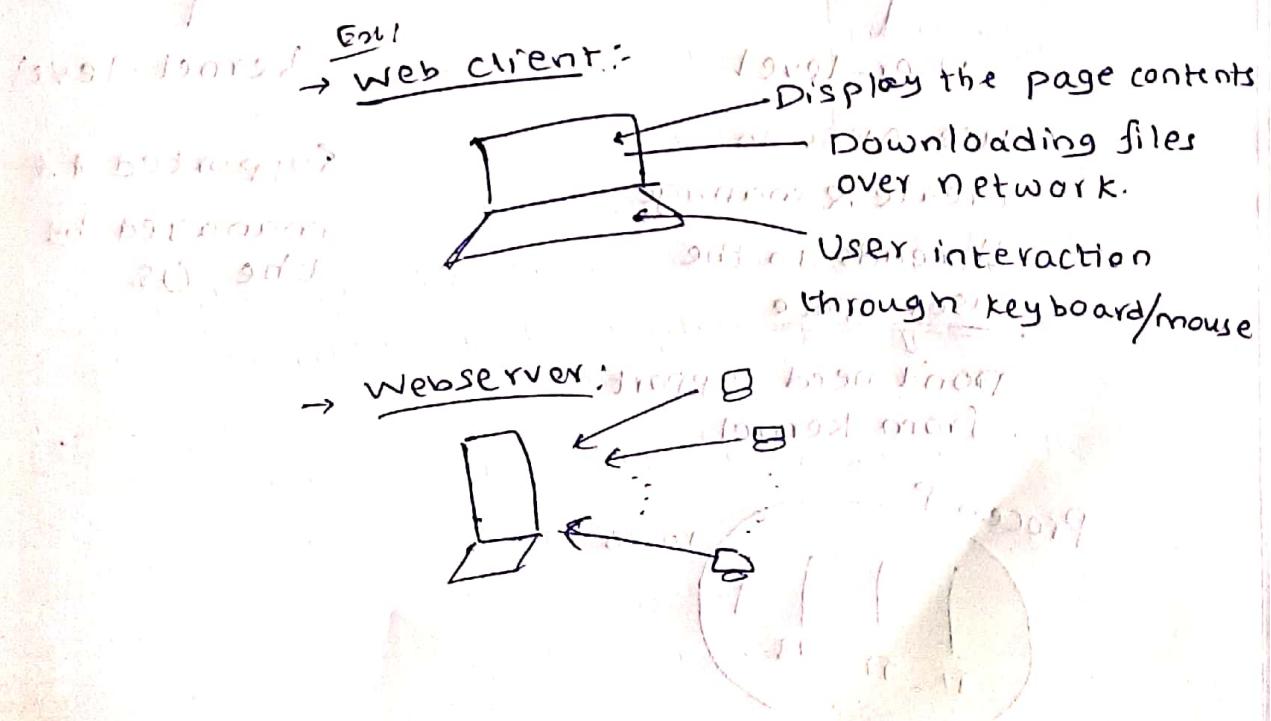
→ Contains a thread ID, Program Counter (PC), registers and stack.

→ Shares the code, data, files etc. with the peer threads.



→ Why do we require multiple threads?

→ Context switch overhead is significantly reduced. [since only Registers & stack need to be saved, remaining all are same].



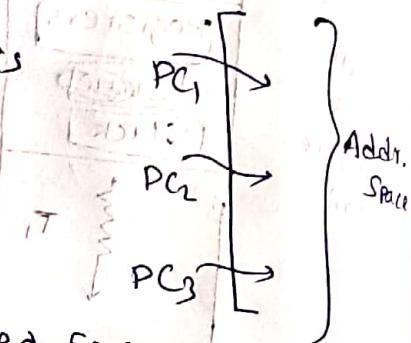
- Multithreading benefits
 - Context switch among threads is faster.
 - Better responsiveness in interactive applications.
 - Better resource utilization among peer threads

↳ same address space (since code is shared)
 ↳ different threads of activity

→ Creating a new process is 10-30 times slower as compared to threads.

→ Better scalability

↳ threads can be used for multiprocessor environments.



→ Multithreading models

Types of threads

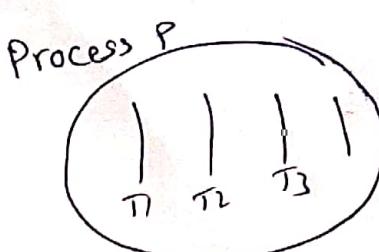
User-level threads

→ User-level threads use APIs to create & manage threads in the user space.

→ Don't need support from kernel.

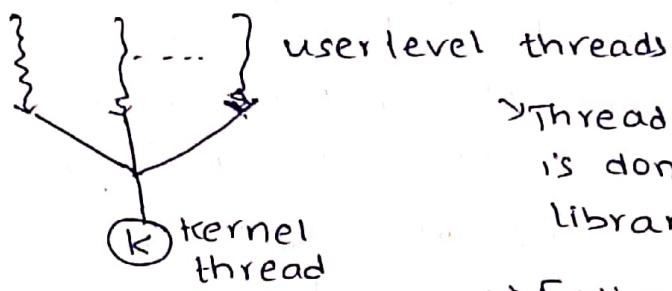
Kernel-level threads

→ Supported & managed by the OS.



If a thread makes a blocking system call to the kernel, whole process is blocked since the kernel doesn't know about threads as separate entities.

(a) Many-to-one Model:

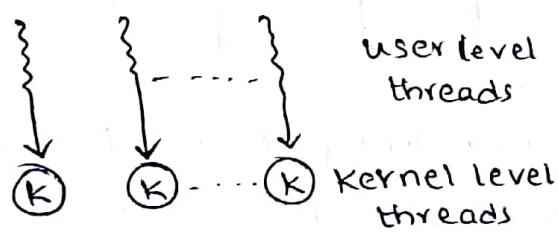


→ Thread management is done by a thread library in user space

→ Entire set of threads will be blocked if one of them makes a blocking system call.

Not suitable for multiprocessors

(b) One-to-one Model:



→ provides more concurrency

→ when ~~a~~ a thread blocks, others can run

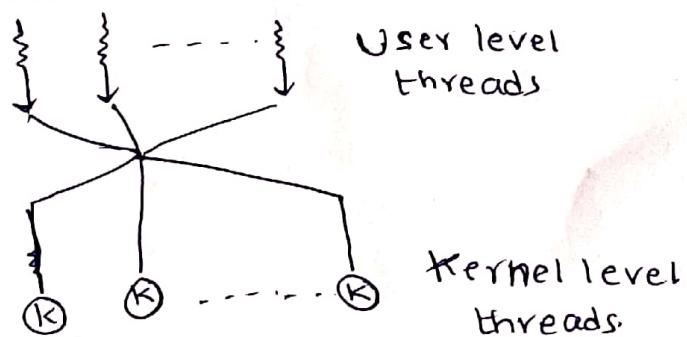
→ Additional overhead w.r.t. kernel.

[.. kernel has to manage many kernel threads]

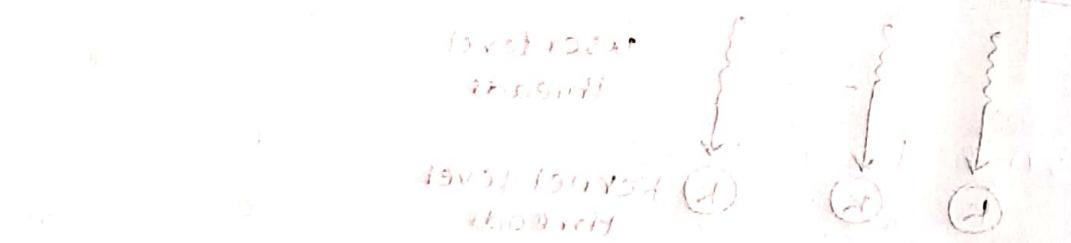
→ Linux / Windows follows this.

→ can be extended to multi processors

(c) Many-to-many Model:



→ Solaris OS used this model



and called *proboscis* snails
because of lobes on dorsal side.

• Conclusions drawn at end of 2014