

→ Web link :-

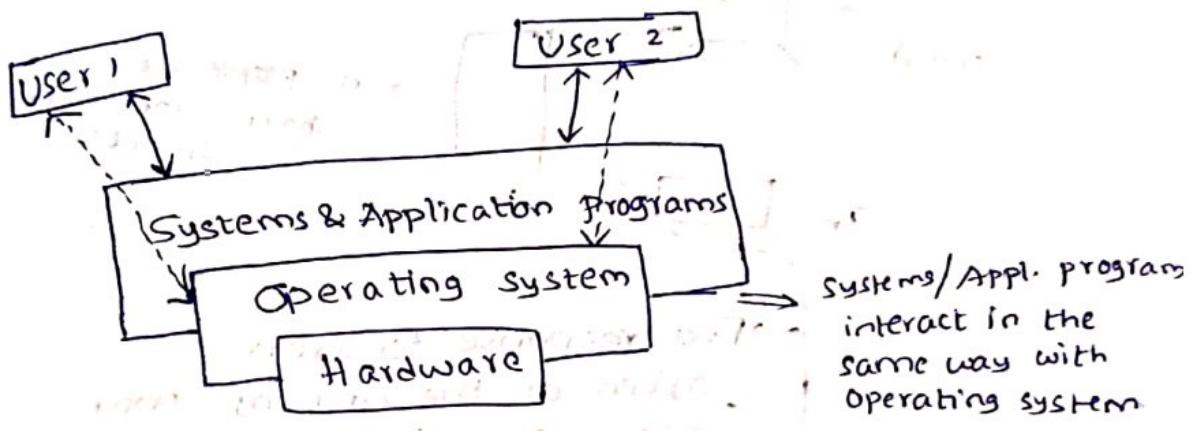
http://www.facweb.iitkgp.ac.in/~isg/os

→ Books :-

→ Operating System Concepts, 9th edition
- Silberschatz, Galvin and

* What is an OS?

→ A program that acts as an interface between user program and computer hardware.



→ Systems programs → programs that assist the user in creating/running programs

Ex: Compiler, loader/linker, editor...

→ Application programs →

Ex: Calculator, word processor...

→ Primary goals of an OS

→ Make a computer system easy to use.

→ Ensures efficient utilization of the

resources

→ CPU, Time, Memory, files, I/O devices

* Evolution :-

→ Mini Computers and Mainframes

Ex: IBM 1130

(1970 - 80)

[Hollerith code → in Punched Cards]

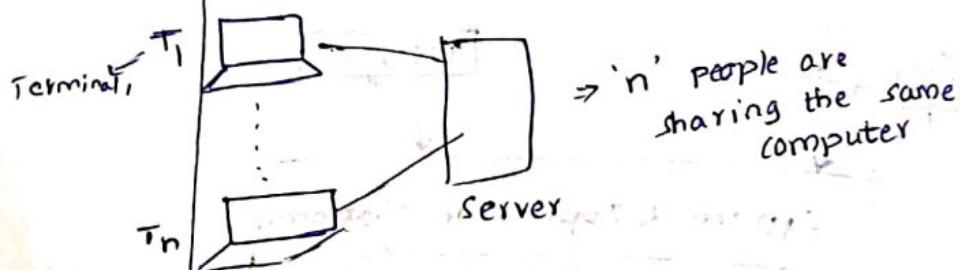
→ Users used to share a computer

→ Best utilization of resources

→ User experience has compromised.

→ Workstations and Terminals

(1980 - Present)



→ Fast response to users
often at the cost of poor
Resource utilization

Ex: P_1
Read
↓
10 min. Compute
↓
. Print

→ while P_1 is in compute,
if P_2 is called, then P_1
has to saved & stopped -
temporarily → more # registers
are used.

→ Desktops^(PCs) and Laptops (1985 - present)

→ Flexible user interface, easy to use

→ Single-user system (fully personal)

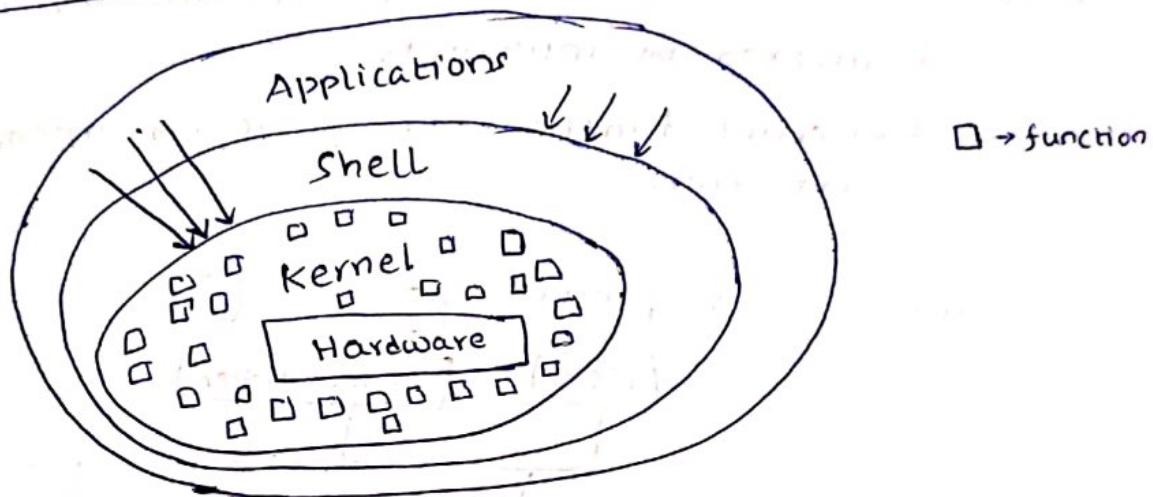
→ Mobile Systems

→ Constrained battery life

→ User experience

terminology

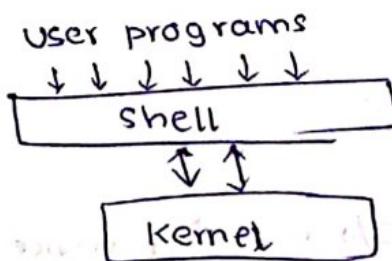
* A Typical structure of OS (w.r.t UNIX)



→ Kernel → A set of functions that provide the essential functionalities of the OS, and the associated data structures.

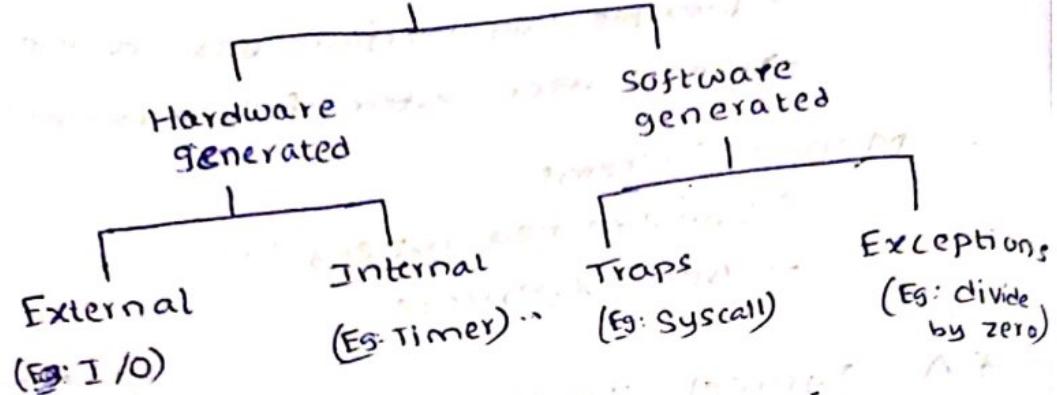
→ kernel runs at all times → INCORRECT Statement

→ shell → A program which allows user programs to access the functions provided by the kernel.



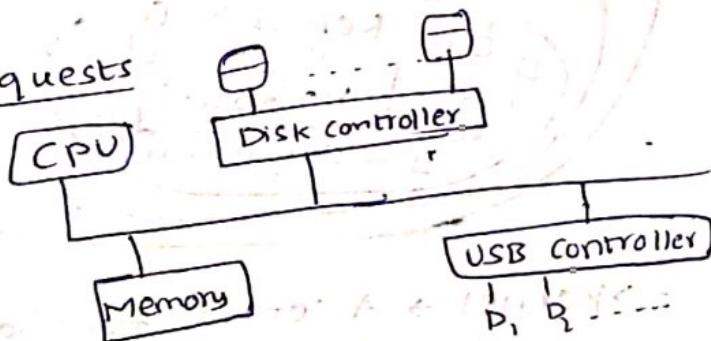


Types of Interrupt



- * How are the OS (i.e kernel) functions invoked by interrupts.
 - Kernel functions are basically interrupt handlers.

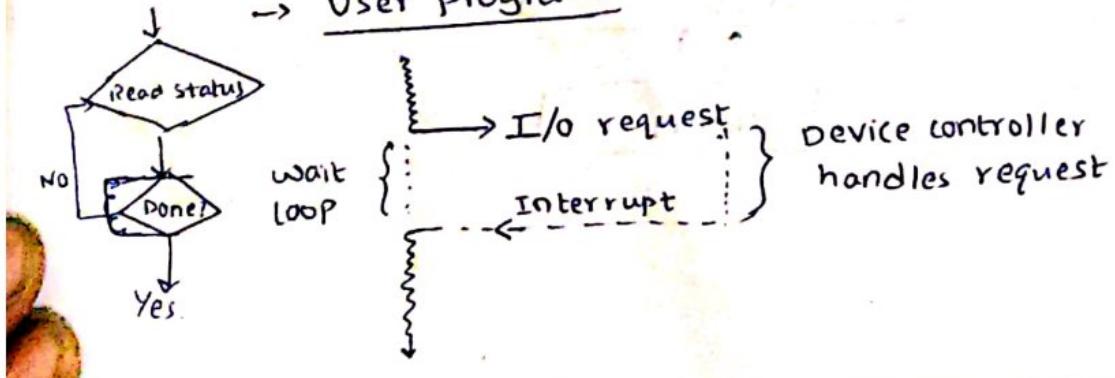
* Handling I/O requests



- If there is no disk controller, CPU has to spend so much time in disk access.
- Disk controller → has a dedicated processor to carryout disk access and store the read data in a buffer

* Alternate :-

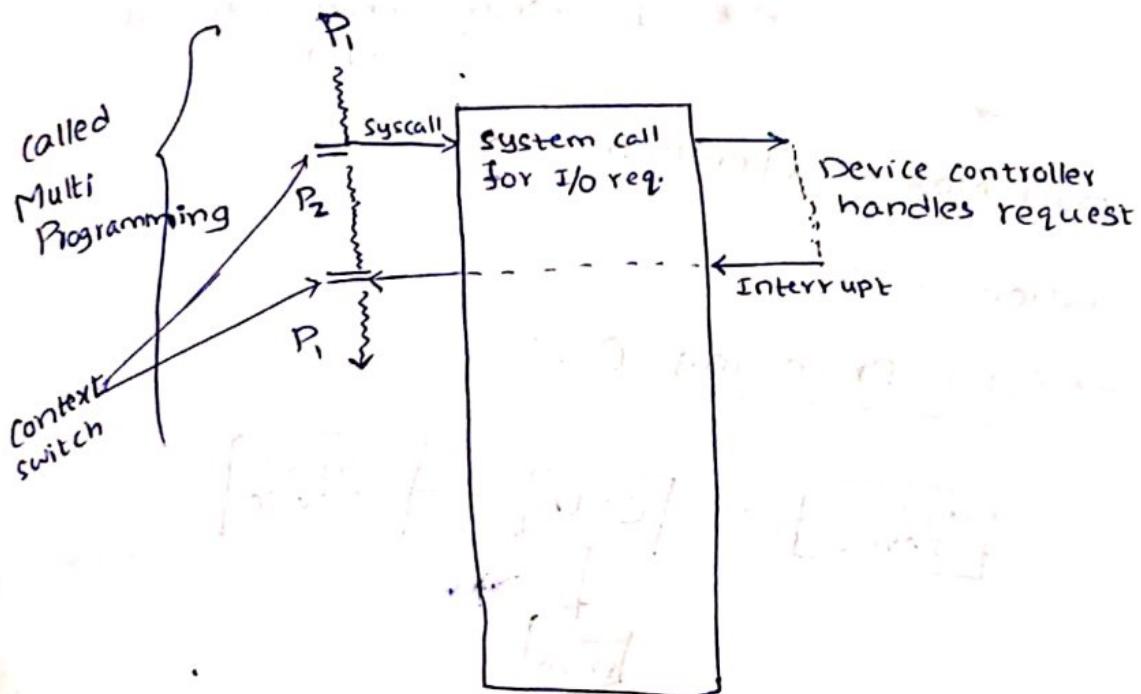
User program



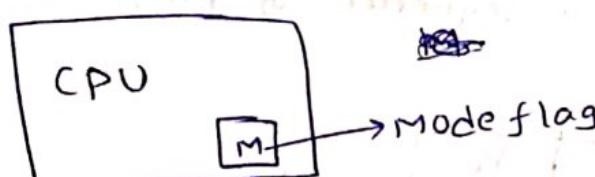
→ The above scheme is called ~~busy waiting~~
"BUSY WAITING".

* Alternative 2 :-

→ Active programs P_1, P_2, P_3 .

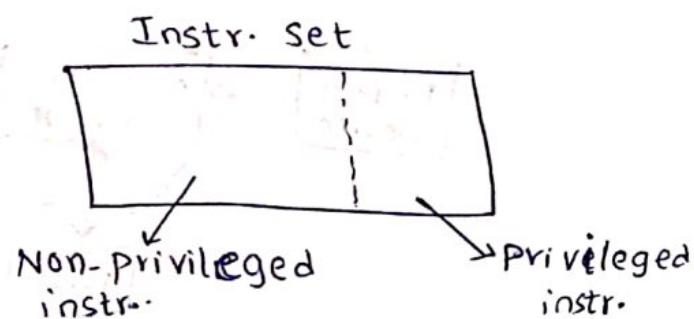


* Syscall

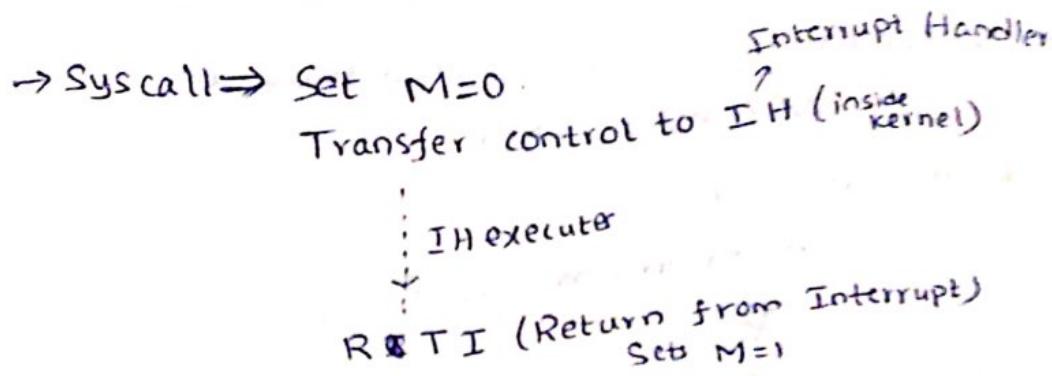


$M=1 \Rightarrow$ User mode of execution

$M=0 \Rightarrow$ Supervisor/Kernel/privilege Mode of Execution.



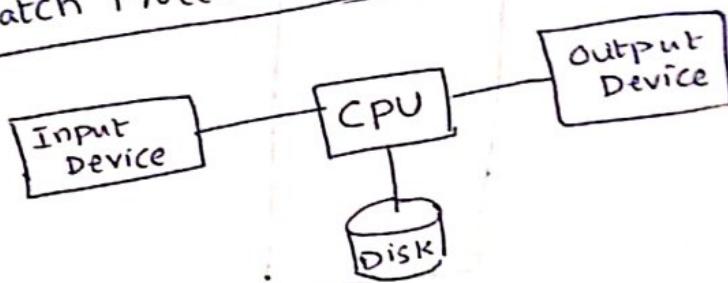
→ CPU can execute a privileged instr. only
when $M=0$
else it will generate an exception.



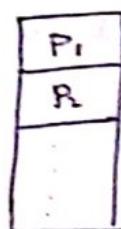
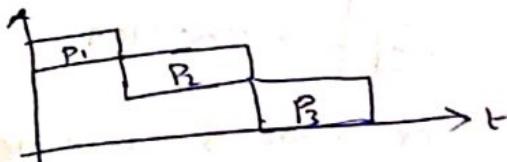
Dual-Mode of Operation.

* Evolution of OS:-

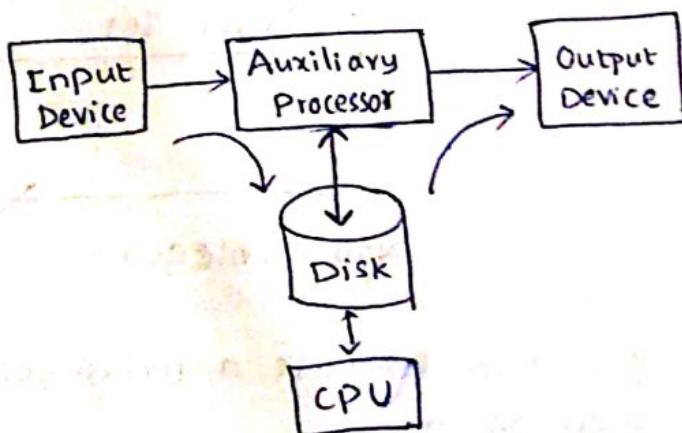
→ Batch Processing OS:-



- programs execute one at a time.
- Batch: a sequence of programs



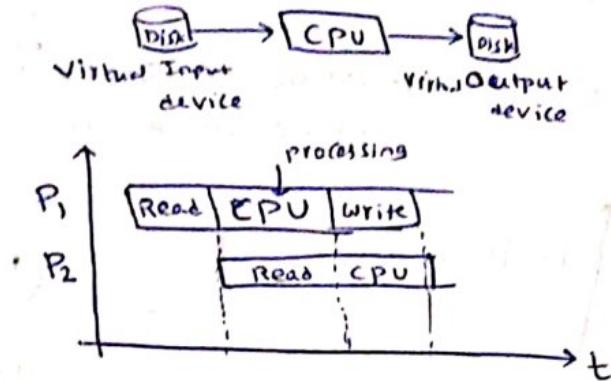
→ Simultaneous Peripheral Operations On Line (SPOOL)



→ Auxiliary Processor

- just reads data from input device & puts it or sends to disk
- (or) reads data from disk and sends to output device

→ Concept of Virtual I/O devices



→ Multiprogrammed OS:-

- Multiple programs loaded in memory at the same time.
- You can switch execution from one program to another.
- Parallelism in CPU & I/O operation.
- Concepts of CPU-bound jobs & I/O-bound jobs
 - programs which do more computation than I/O.

Eg: Matrix multiply ($n \times n$)

Objective
Maximum resource utilization
utilization
by a good mix
of CPU & I/O-bound jobs

I/O: $O(n^2)$
computation: $O(n^3)$

\Rightarrow CPU-bound job

2. Payroll \rightarrow I/O-bound job.

→ Time sharing/Multi-tasking OS:-

- users sit on terminals & interact in real time.

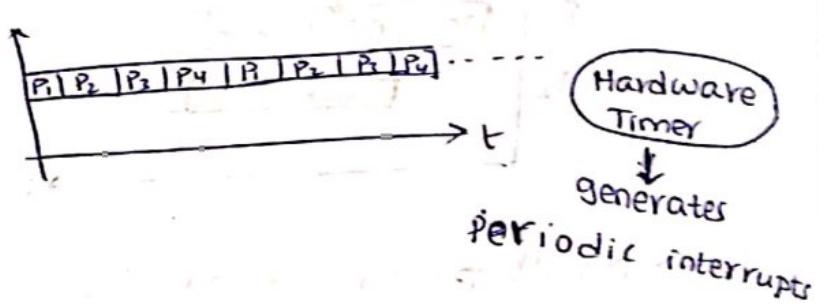
→ Objective \rightarrow Maximum user satisfaction
minimization of response time

→ Very frequent context switching between programs

→ Let's say there are 50 users

if we follow Multiprogramming, one job is being done & remaining 49 users have to wait

→ $P_1, P_2, P_3, P_4 \xrightarrow{\text{say}}$ allocate short bursts of time to each prog.



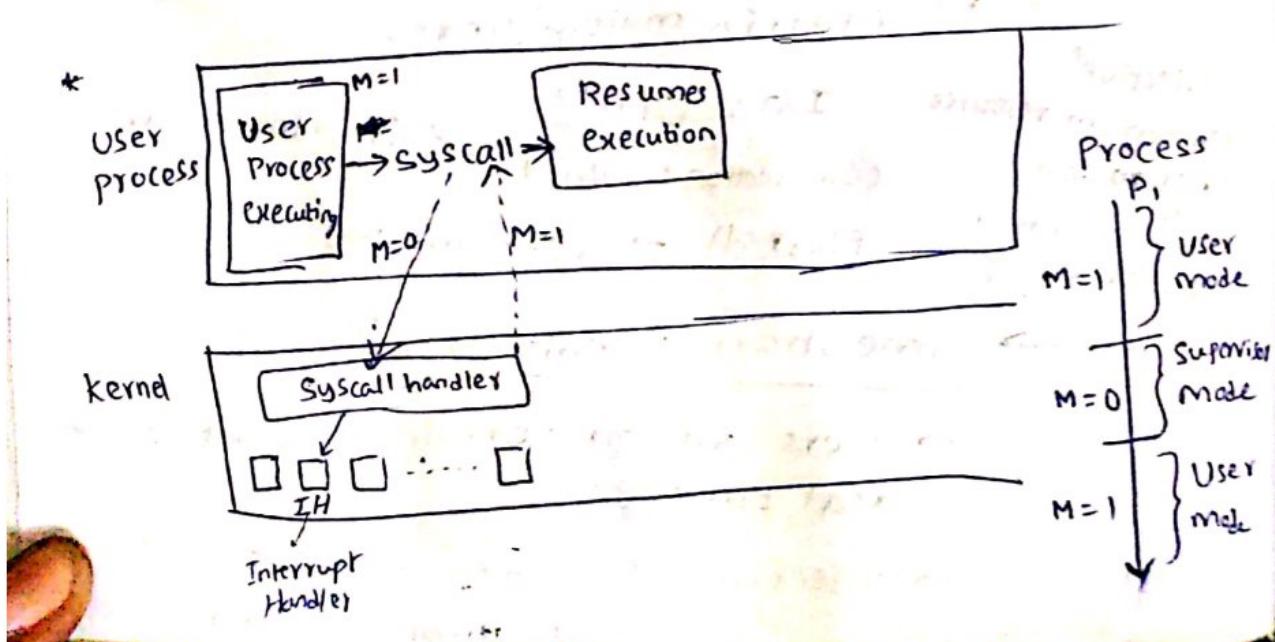
→ Real-time OS:

→ process - user requests based on deadlines
come as interrupts

→ Distributed OS:

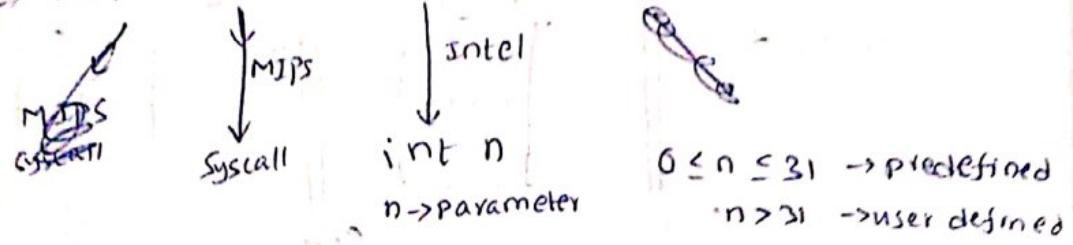
→ part of computation happens in one place and other part in other place.

* process → a program in execution



* System Calls

- It is a machine instr.
- different for different ISA.



→ system call instr. is executed

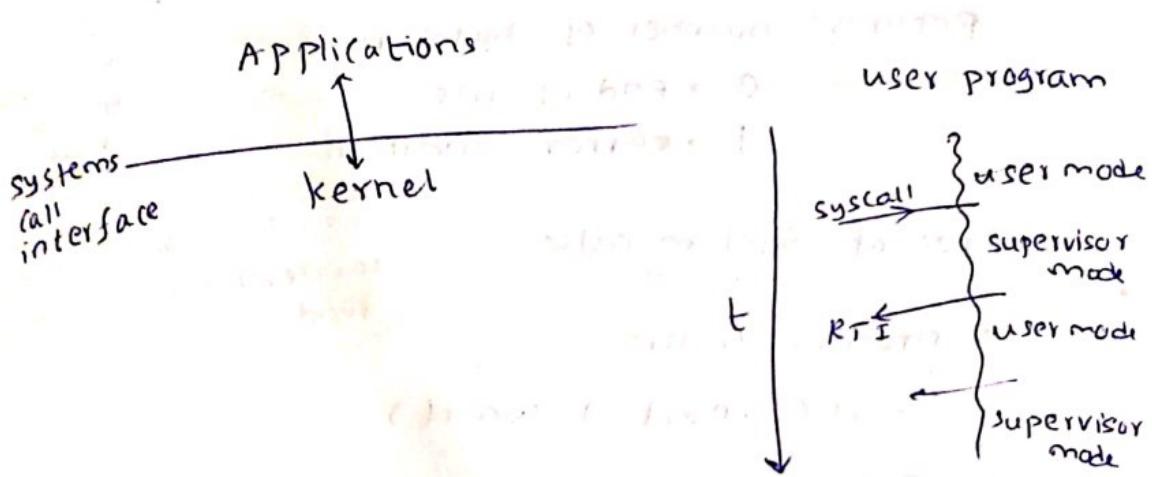
→ software interrupt

→ Mode = supervisor

→ Invoke interrupt handler

→ Return from interrupt

(Mode = user)

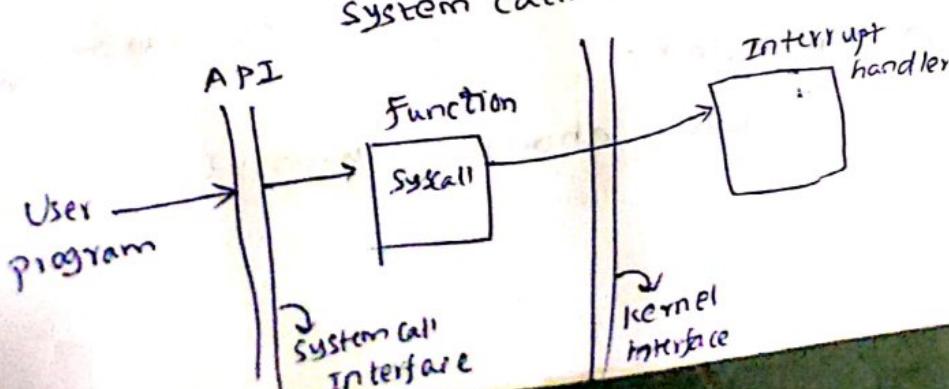


* System call interface :-

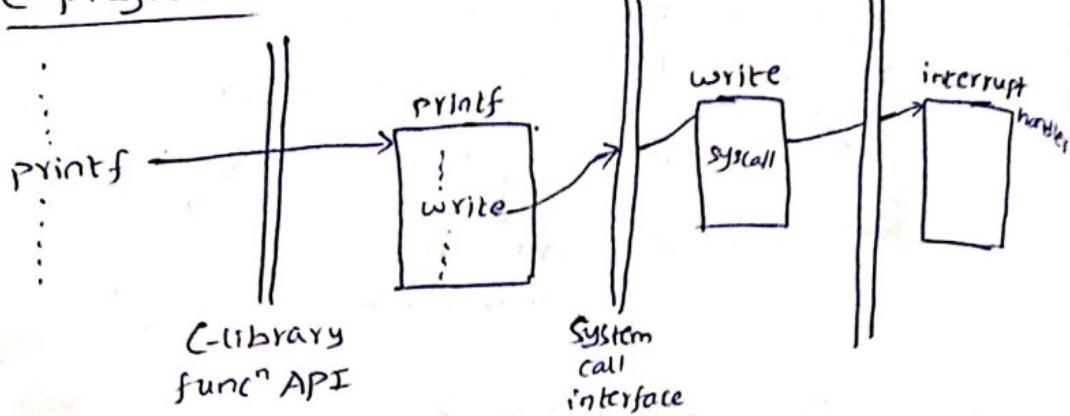
→ provided by the OS

→ A set of low-level func's that can be called from a user program.

→ These low-level func's execute some system call.



* C program



* Example API:

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t count);
```

Returns: number of bytes read

0 → end-of-file

-1 → error condition.

max. no. of bytes to read

→ Types of System calls:

(a) process control:

fork(), exit(), wait().

(b) File manipulation:

open(), close(), read(), write()

(c) Inter-process communication: related to semaphores

pipe(), semget(), semop(), semctl(),

shmget(), shmat(), shmdt(), shmctl(),

kill(), signal().

related to shared memory

(d) protection:

chmod(), chown(), umask()

* Problems

1. Related to SPOOLing

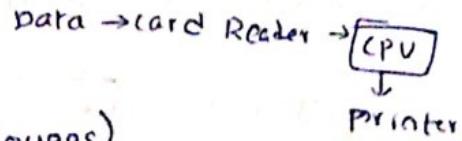
→ Payroll system (1000 employees)

→ Reading a card - 100 ms

→ Printing a line - 100 ms

→ Read/write from disk - 10 ms

→ Processing time - 5 ms/employee



(a) NO SPOOLING

$$\begin{aligned} \text{Elapsed time} &= 1000(100 + 5 + 100) \text{ ms} \\ &= 205 \text{ s.} \end{aligned}$$

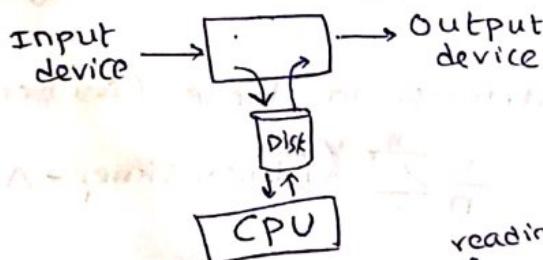
only one line is to be printed

$$\begin{aligned} \text{CPU idle time} &= 205 - 1000 \times 5 \text{ ms} \\ &= 200 \text{ s.} \end{aligned}$$

Card reader reads one card every $(100 + 5 + 100) \text{ ms} \approx 5 \text{ cards/second}$

$$\text{CPU utilization} = \frac{5}{205} \times 100 = 2.5\%$$

(b) SPOOLING



$$\begin{aligned} \text{Elapsed time} &= 1000(10 + 5 + 10) \text{ ms} \\ &= 25 \text{ s} \end{aligned}$$

reading from disk
↓
writing to disk

$$\text{CPU idle time} = 25 - 1000 \times 5 = 20 \text{ s.}$$

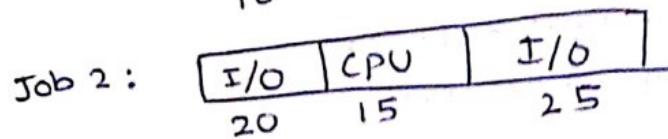
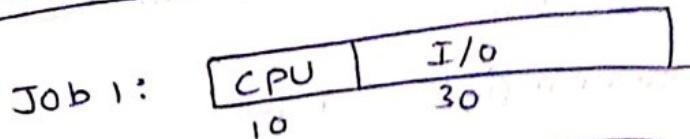
Card reader reads one card every

$$(100 + 10) \text{ ms} \approx 9 \text{ cards/second}$$

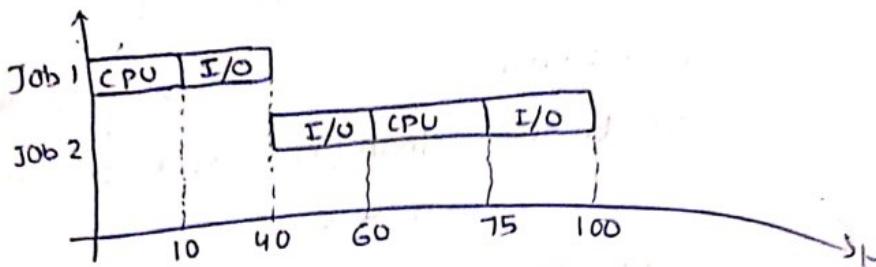
↓
reading from card ↓ writing to disk

$$\text{CPU utilization} = \frac{5}{25} \times 100 = 20\%$$

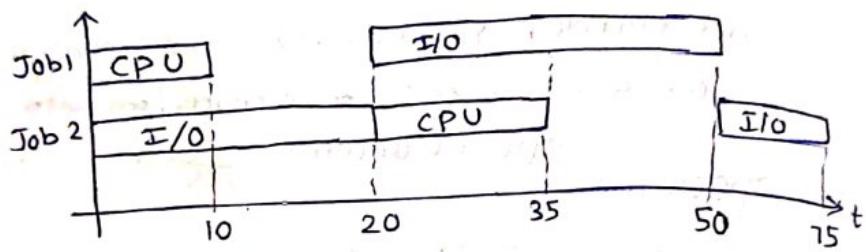
2. Related to multiprogramming :-



(a) Batch execution :-



(b) Multiprogramming :-



→ Average turnaround time (say there are 'n' jobs)

$$= \frac{1}{n} \sum_{i=1}^n (\text{Finish time}_i - \text{Arrival time}_i)$$

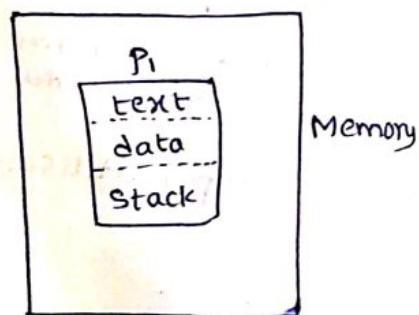
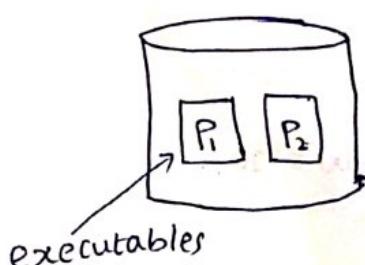
→ Average throughput

= No. of jobs completed per unit time

* Process :-

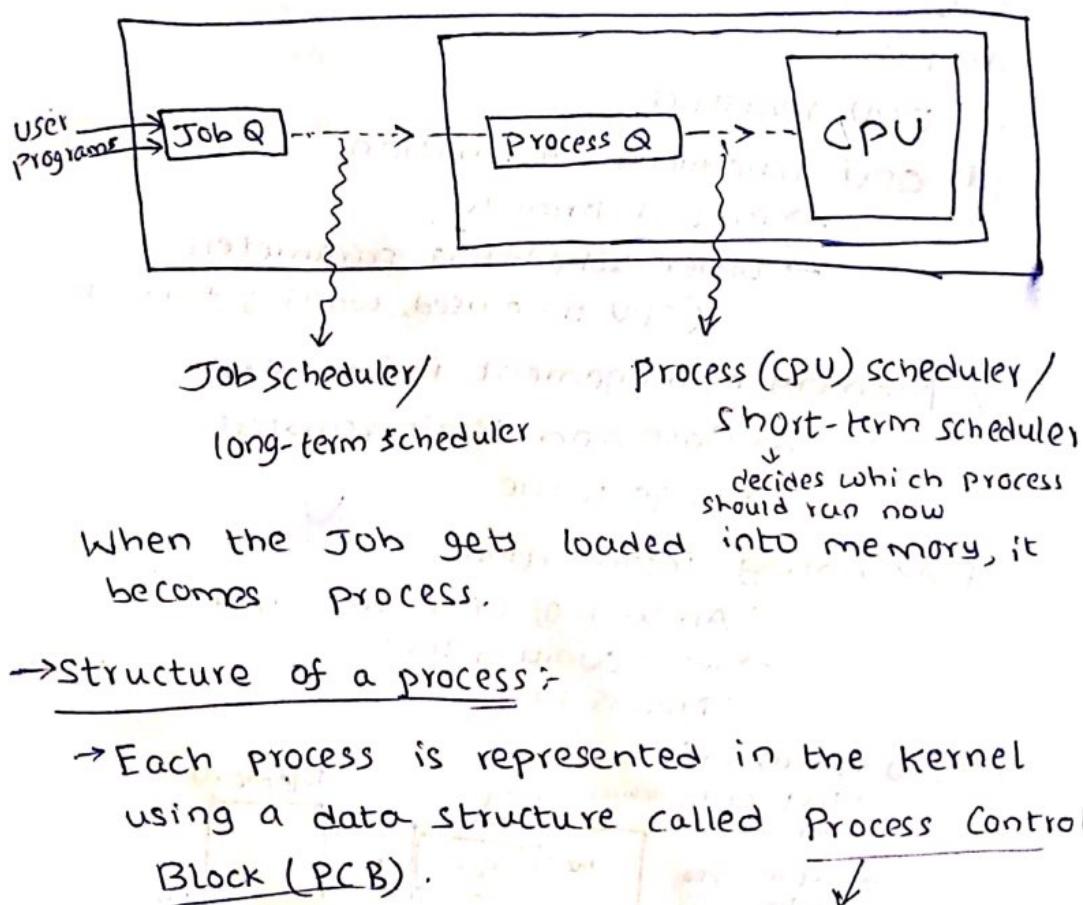
→ A program in execution

(terminologies of UNIX)



- A process consists of several regions
 - text → contains executable code
 - data → contains global data (static data)
 - stack → contains dynamic data (e.g. recursion)
 - heap → dynamic data controlled by the program.
(e.g. malloc).
- shared → Region shared by more than one process.

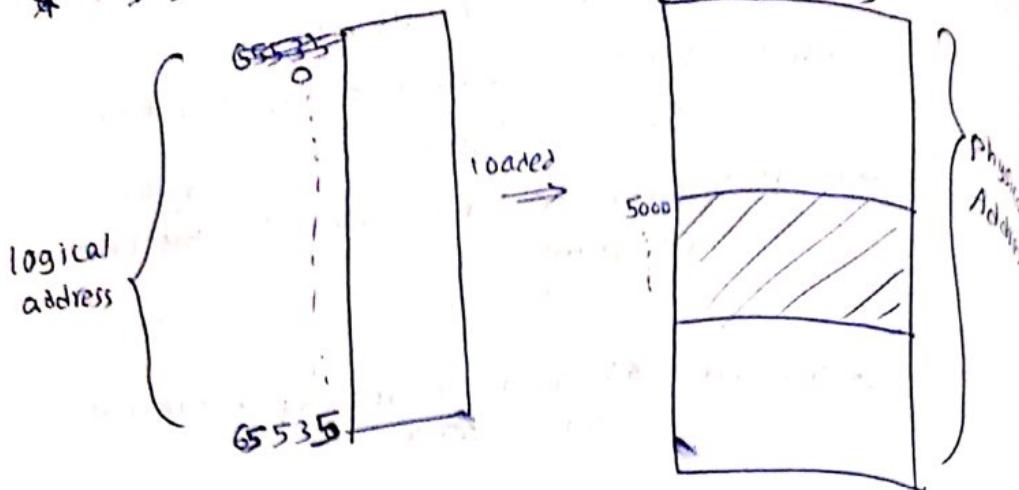
- Batch-processing and multiprogramming
 - ↳ user programs were called jobs



- Each process is represented in the kernel using a data structure called Process Control Block (PCB).



* → say Address - 16 bits



* Typical PCB entries:

(a) process state

(b) PC

(c) CPU registers

(d) CPU scheduling information

→ process priority

→ other scheduling parameters

(CPU time used, waiting time etc.)

(e) Memory management information

→ Base and limit registers

→ Page table.

(f) Accounting information

→ Amount of CPU time used

→ Total waiting time

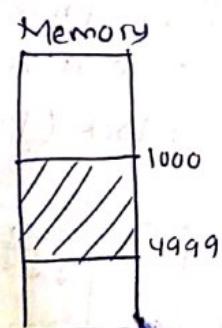
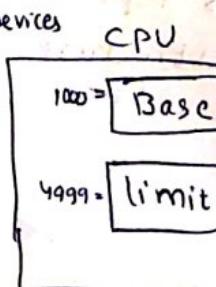
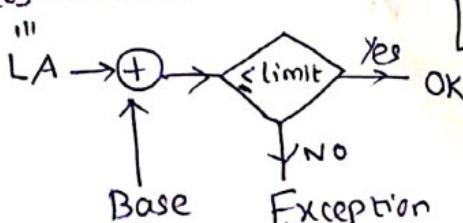
→ process id

(g) I/O status info

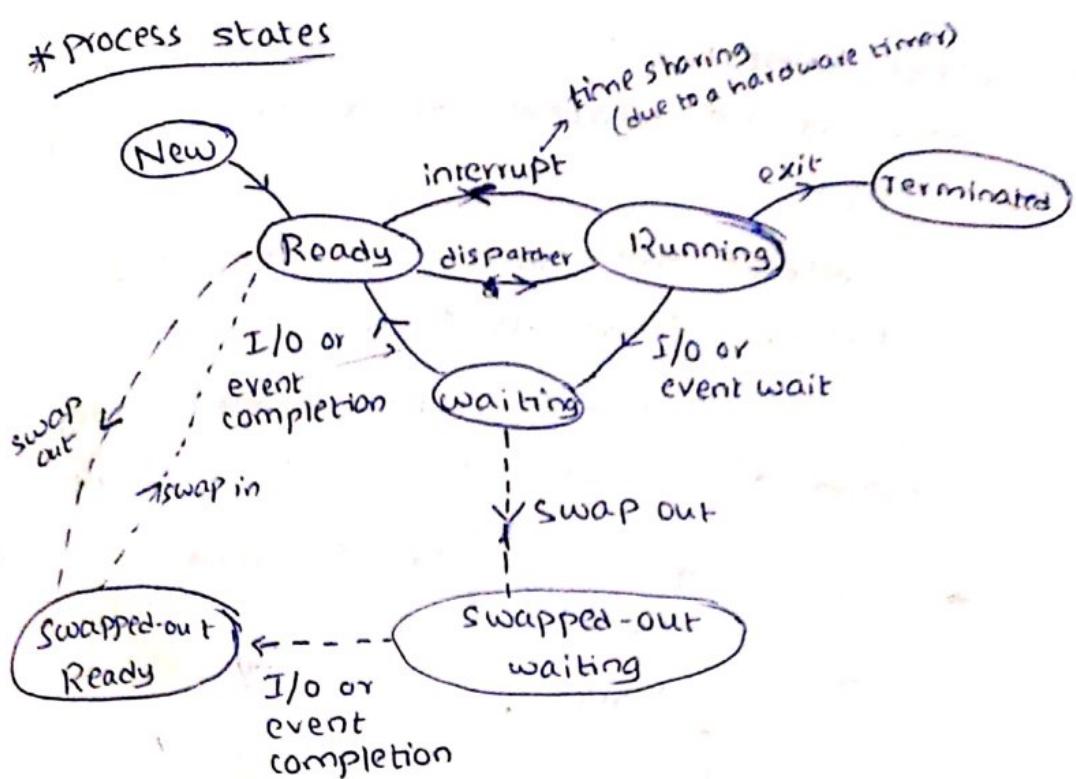
→ list of I/O devices being used

→ list of open files

logical address



* Process States



→ New: A process is being created

→ Ready: The process is ready to execute
(waiting for CPU)

→ Running: The process is running on the CPU.

→ Waiting\Blocked: The process is waiting for some event to occur.
(I/O operation, signal from other process etc.).

→ Terminated: Finished execution and has lost the system.

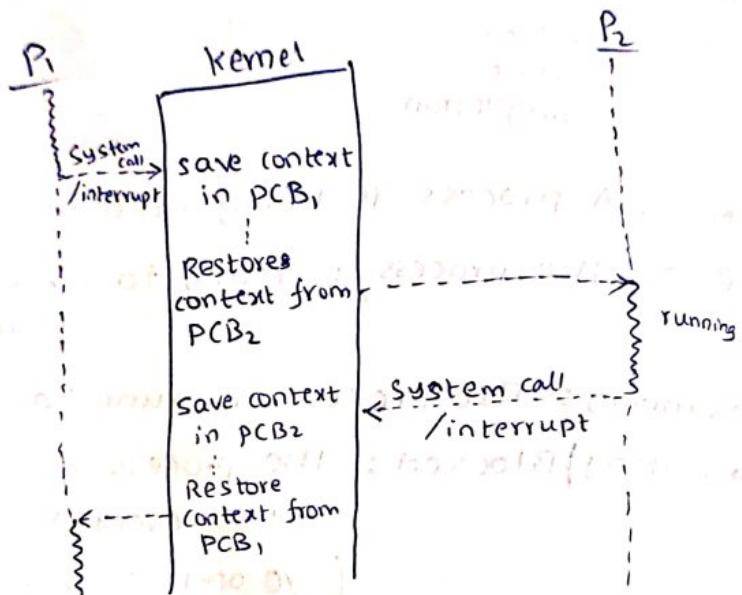
* Memory is limited

⇒ No. of processes that can be accommodated is limited.



- Context switch → overhead
- Saves the context of a process in PCB
- Restores the context of the next process from its PCB.
- 1000's of context switches happen every second → must be as fast as possible

* } = running | = idle (in ready state)



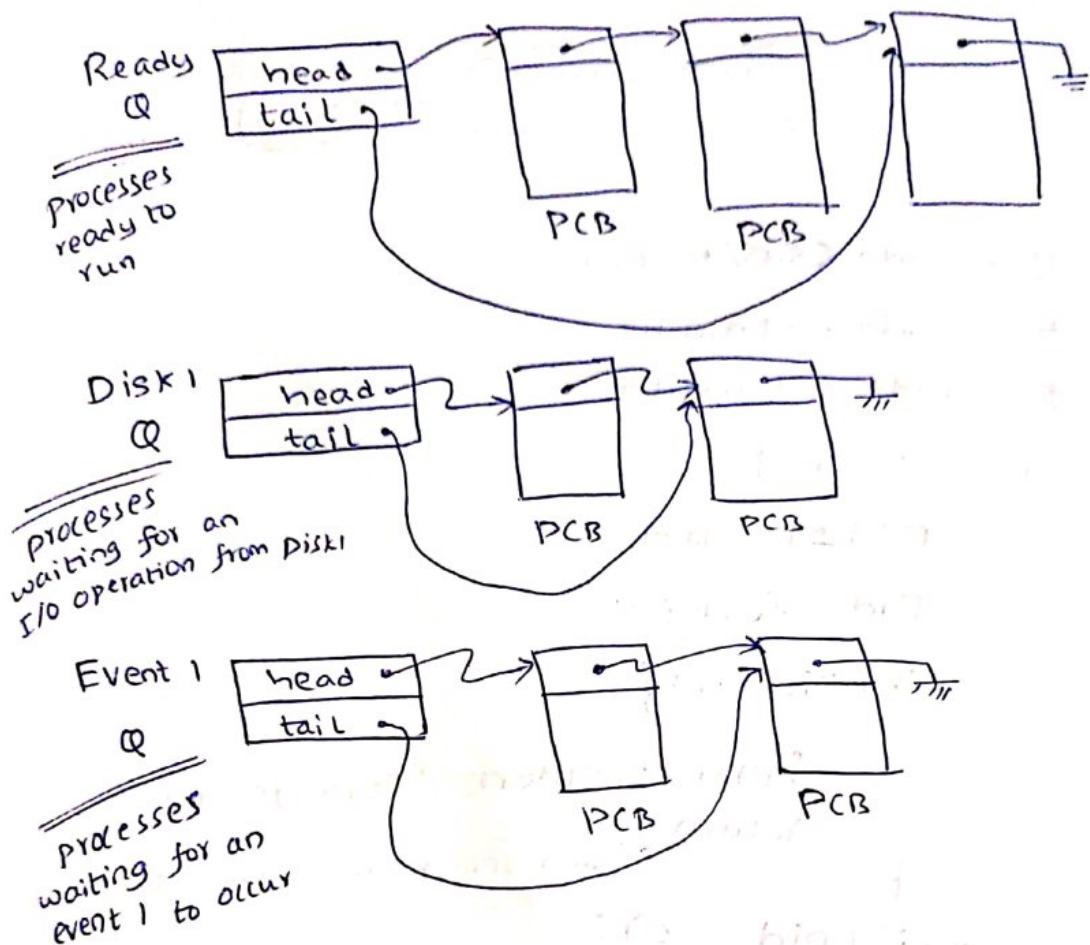
P_1 & P_2 are concurrent processes

* Process scheduling

→ Ready state → queue of processes
 When the running process is interrupted, there must be some algorithm to select a process from the queue to make it running

* Process scheduler or dispatcher selects a process from the ready queue for execution whenever there is a context switch.

* One of the many scheduling queues maintained by kernel.
→ linked list of PCBs



* Operations on Processes:

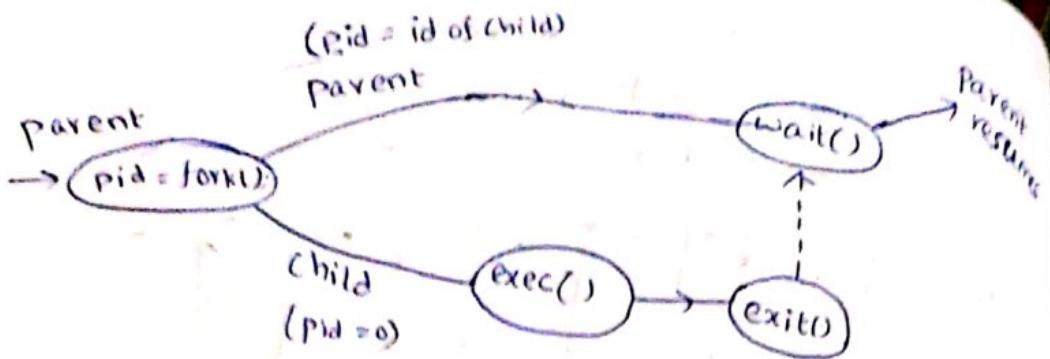
→ Process creation:-

→ 'ps' command lists all the processes active.

→ every process has a unique id (pid).

→ A single process that gets created in the beginning → called "init" with pid=0
i.e. when booted

→ all other processes are created using fork.



```

* #include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main(){
    Pid = fork();
    if (Pid < 0){
        fprintf(stderr, "Fork failed");
        return 1;
    }           ↴ non-zero return value ⇒ some error
    else if (Pid == 0){
        execvp("/bin/ls", {"ls", NULL});
    }
    else {
        wait(NULL);
        printf("Child complete");
    }
    return 0;
}

```

→ Process Termination:

→ The exit() and wait() system calls can be used.

→ child process

```
exit(2);
```

Parent process

```
pid_t pid = 2;
```

```
int status;
```

```
pid = wait(&status);
```

id. of child
who has
terminated.

[∴ Parent may
fork many
processes, it
can find
which process
terminated.]

Value passed
as argument
to exit() in
~~child~~.
i.e., Here it is 2.

→ Normal process termination

→ child executes exit()

→ parent executes wait()

If both these happen, the OS ~~terminates~~
considers that it happened successfully &
all entries are deleted from process tables.

→ A child process terminates but parent
has not called wait().

⇒ The child becomes zombie
process

↳ i.e., some entries will remain
though it's not active

→ A child process terminates & parent also
terminates but without calling wait.

⇒ The child process becomes
orphan process

* All orphans are assigned to "init" as parent (i.e. init is parent) & init calls wait() periodically.

→ ∃ multiple process in various states of execution
 ↳ concurrent process

independent

A process that doesn't affect or be affected by other process running in system (can't share any data with other processes).

co-operative
Process that are dependent.
(Share some data)



* Co-operating Processes:

→ Some mechanism using which the processes can exchange data.

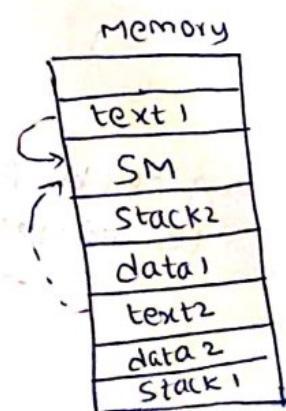
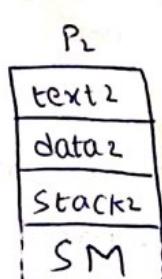
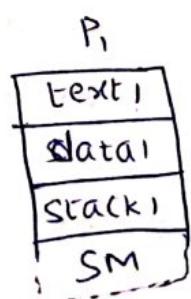
Inter Process Communication (IPC)

Shared Memory (SM)

Message Passing



(a) Shared Memory (SM)



→ System calls:

→ `shmget()` ⇒ creates a SM segment in memory

→ `shmat` ⇒ attaches the SM segment to the address space of a process.

→ `shmdt` ⇒ detaches the SM segment from a process.

→ `shmctl` ⇒ removes the SM segment from memory

* `int pid;`
`char * myseg;`
`key_t key; int shmid;`

`key = 758;`

↑
no. of bytes

`shmid = shmget(key, 250, IPC_CREAT | 0666);`

`pid = fork();`

`if (pid == 0){`

`myseg = shmat(shmid, NULL, 0);`

Create a new SM segment with key if it does not exist

`shmdt(shmid);`

666 = 110 110 110

rwX rwX rwX
for user for group every user has a group id also
other

`}`
`else {`

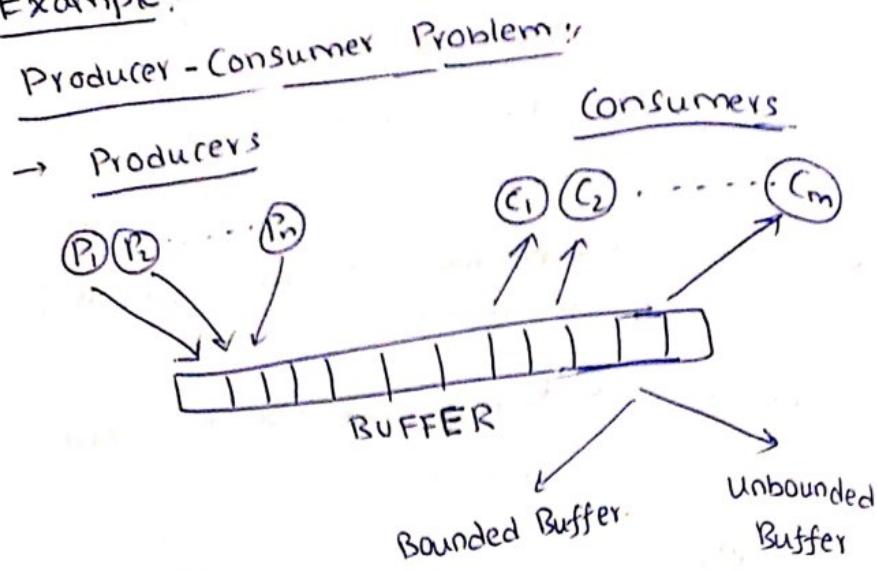
`myseg = shmat(shmid, NULL, 0);`

`}`

`shmdt(shmid);`

`shmctl(shmid, IPC_RMID, NULL);`

* Example:-



If P_i sees that the Buffer is full, then P_i has to wait until next write.

* Pseudo code:-

```
#define BUFF_SIZE 50
```

```
typedef struct {  
    int value;  
} item;
```

```
item buffer[BUFF_SIZE];
```

```
int in=0;  
int out=0;
```

out = pointer to the first data item in the buffer

in = pointer to next free location in Buffer

→ Located in Shared memory

Producer

```
while(true){
```

< Produce an item in next_P >

```
    while((int+1)%BUFF_SIZE == out);
```

```
    buffer[in] = next_P;
```

```
    in = (int+1)%BUFF_SIZE;
```

```
}
```

consumer

```

item next_c;
while(true) {
    while(in == out);
    next_c = buffer[out];
    out = (out + 1) % BUFF_SIZE;
    <consume item in next_c>
}

```

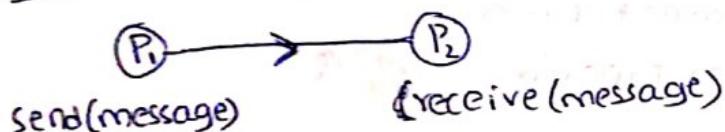


Max data items that can be put in Buffer

= ~~BUFF_SIZE~~

$$= \text{BUFF_SIZE} - 1$$

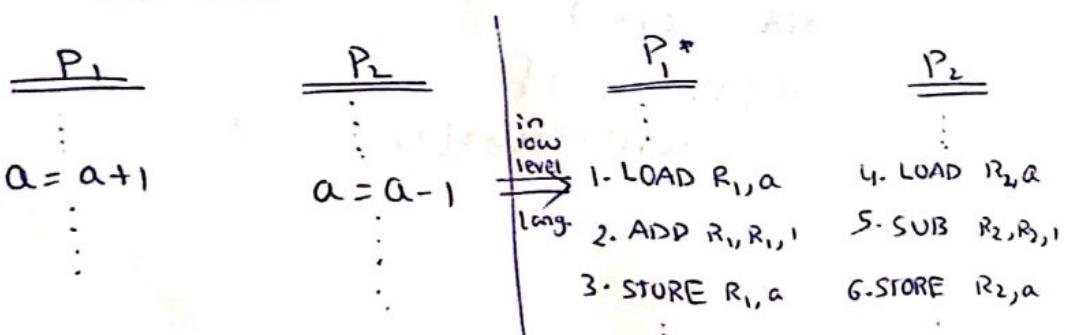
(b) Message-Passing :-



Example .

2 processes trying to update a variable is SM.

SM | $a=100$



let's say St:1 was executing & there's a context switch

The sequence is like say

1 4 5 6 2 3 \Rightarrow Final value of $a = 101$

for 4 5 1 2 3 6 \Rightarrow " " = 99

↓
The above is called RACE condition.

also happens when 1, is given
to SR latch.

Example: Producer - Consumer problem

Producer

```
message next_P;  
while(true){  
    <produce item in next_P>  
    send(next_P);  
    ======  
}
```

Consumer

```
message next_C;  
while(true){  
    receive(next_C);  
    <consume item  
    in next_C>  
    ======  
}
```

→ Two approaches:

(a) unnamed pipes

(b) Named pipes/message Q.

* Unamed pipe:-

→ pipe() system call.

```
int xyz[2];  
Pipe(xyz);  
xyz[0] ← Read  
xyz[1] ← Write
```

Pid = fork();

if (Pid == 0){

write(xyz[1],.....)

}

else {

read(xyz[0],....)

}

→ Parent & child can communicate
since they have access to both
ends of the pipe.

→ unnamed pipe can be used by processes that are generated by a common parent.

* Named Pipe:-

→ They use a common temporary file.

char *myfifo = "/tmp/myfifo"; files in /tmp/ are accessible to every user

mkfifo(myfifo, 0666);

more

parent

```
fd = open(myfifo, O_WRONLY);
;
write(fd, ...);
close(fd);
```

child

```
fd1 = open(myfifo, O_RDONLY);
;
read(fd1, ...);
close(fd1);
```

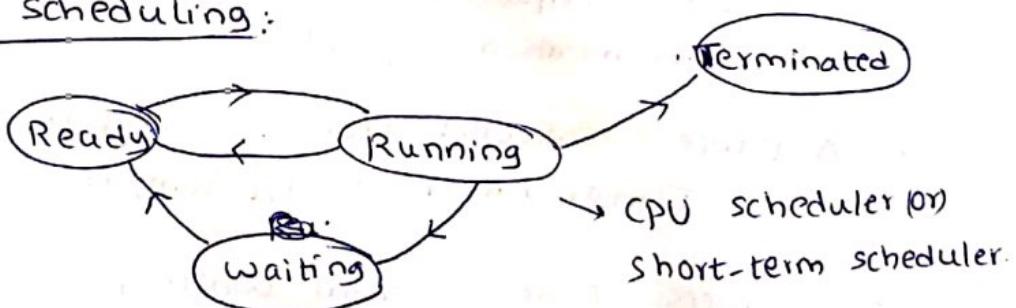
→ If even one of the processes calls mkfifo, then all the other processes can use it as named pipe.

(c) Process Comm' across machines

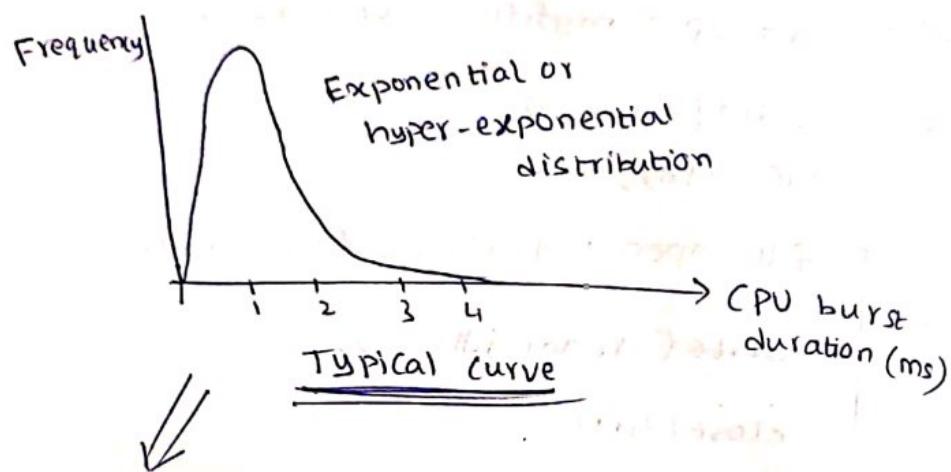
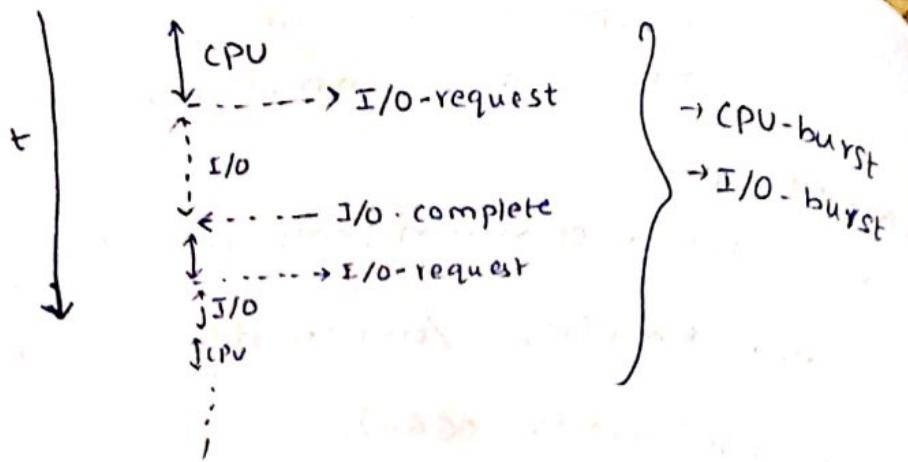
→ Remote procedure call

→ Socket.

* CPU Scheduling:-



* Typical process execution :-



There are large no. of CPU bursts with very small duration.

& Small no. of CPU bursts with large duration.

* Scheduling decisions;

- (a) A process switches from running to waiting
 - I/O request (or) Parent invokes wait() for termination of child.
- (b) A process switches from Running to Ready.
 - Timer interrupt, I/O-completion interrupt
- (c) A process switches from Waiting to Ready.
 - I/O completion interrupt, child has terminated.
 - Then, Parent switches to ready.
- (d) A process terminates
 - exit() system call.

→ (a) and (d)

→ Non-Preemptive scheduling

↳ Allowing a process to continue for as long as it wants.

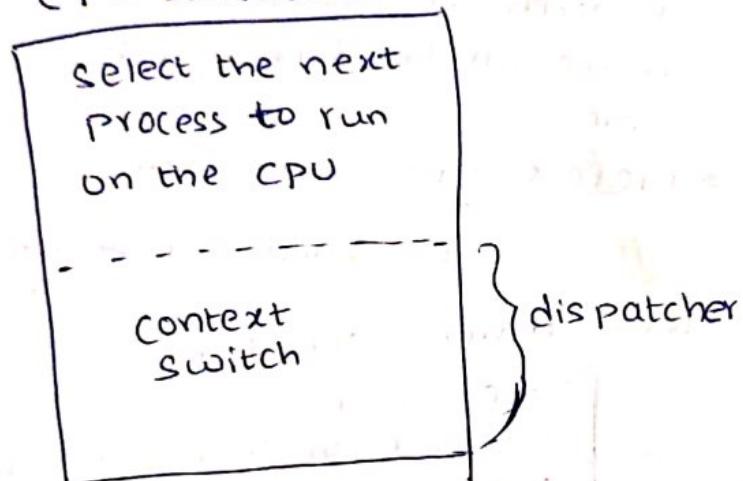
→ Here, the processes are coming out of CPU ~~on~~ on their own.

→ (b) & (c)

→ Preemptive scheduling

↳ forcibly taking the CPU away from a process.

* CPU Scheduler



* Class Test 1

7th February 7-8 PM

Venue:- CSE Dept

Syllabus:- Up to CPU scheduling

* Scheduling criteria

→ (a) CPU utilization

→ (b) Throughput - No. of processes finishing per unit time.

→ (c) Turnaround time - (Finishing time - Arrival time)
when job first arrives,

→ (d) Waiting time - Time taken in the ready and waiting states.

→(e) Response time - after how much time, a response is received

* CPU Scheduling Algorithms

→(a) First-Come First Serve (FCFS)

→ It is a non-preemptive algorithm.

→ Ready list is maintained as a FIFO queue.

* Drawback:

→ Convoy effect

↓
Traffic has to wait when convoy passes

↓
A process with large CPU burst will make all the following processes wait.

→ prefers CPU-bound jobs

↓
Jobs whose CPU-bursts are longer.

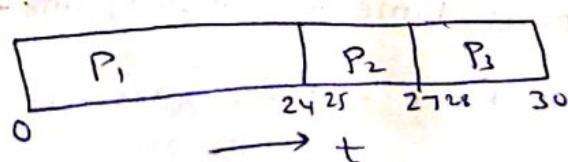
lower device utilization.

Since, if a CPU-bound job is running, I/O-bound jobs behind it in the queue are not even started, then I/O-devices will lie idle

Ex:-

Processes	P ₁	P ₂	P ₃
Arrival Time	0	0	0
CPU Burst	24ms	3ms	3ms

(a) P₁ - P₂ - P₃



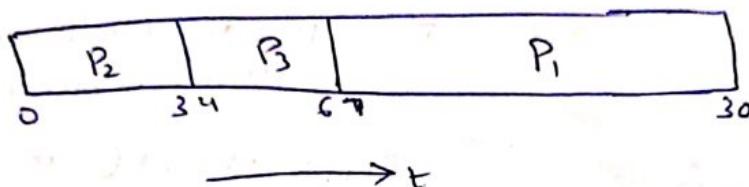
$$\text{Avg. turnaround time} = \frac{(24-0) + (27-0) + (30-0)}{3}$$

$$= 27 \text{ ms}$$

$$\text{Avg. waiting time} = \frac{0 + (27-3) + (30-3)}{3}$$

$$= 17 \text{ ms}$$

(b) $P_2 - P_3 - P_1$



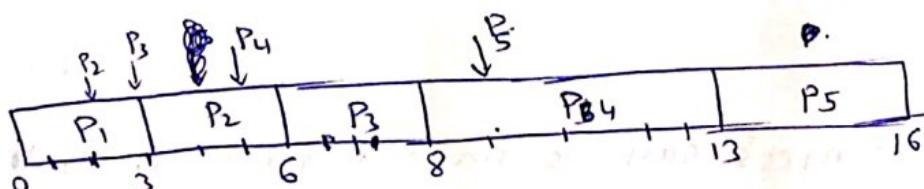
$$\text{Avg. turnaround time} = \frac{(30-0) + (6-0) + (3-0)}{3}$$

$$= 13 \text{ ms}$$

$$\text{AWT} = \frac{(30-24) + (3-3) + (6-3)}{3} = 2 \text{ ms}$$

Ex 2:

Process	P_1	P_2	P_3	P_4	P_5
Arrival time (in ms)	0	2	3	5	9
CPU Burst	3	3	2	5	3



$$\text{AWT} = \frac{0 + 1 + 3 + 3 + 4}{5} = 2.2 \text{ ms}$$

Shortest Job First/Next (SJF/SJN)

→ The process with the smallest next CPU burst is allocated the CPU.
↳ non-preemptive.

→ Predicting the next CPU burst time
→ exponential averaging

t_n - length of n^{th} CPU burst
 T_n - predicted length of n^{th} CPU burst

Then, $T_{n+1} = \alpha t_n + (1-\alpha) T_n$,
where $0 \leq \alpha \leq 1$

$$\Rightarrow T_{n+1} = \alpha t_n + (1-\alpha)t_{n-1} + (1-\alpha)^2 t_{n-2} + \dots + (1-\alpha)^j t_{n-j} + \dots + (1-\alpha)^{n+1} T_0$$

initial value

$$\alpha=0 \Rightarrow T_{n+1} = T_n = T_0$$

→ recent history doesn't have any effect.

$$\alpha=1 \Rightarrow T_{n+1} = t_n$$

→ only recent CPU burst matters

→ The SJF algorithm minimizes the AWT

Avg. Waiting Time

→ Let's say we have "n" processes that all arrive at $t=0$.

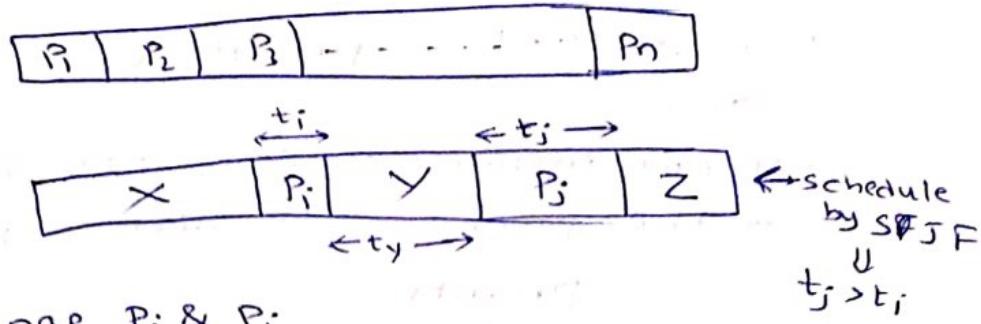
CPU burst times are $t_1 \leq t_2 \leq \dots \leq t_n$.

P_1, P_2, \dots, P_n

Proof that in this case, SJF minimizes AWT.

Proof by contradiction.

→ SJF Schedule



Exchange P_i & P_j:

→ X: no change in waiting time

Z: no change in waiting time

Y: All the processes will have their waiting times increased.

P_j: waiting time decreases by t_y + t_i

P_i: waiting time increases by t_y.

(c) Shortest Remaining Time First (SRT) :-

→ Preemptive version of SJN

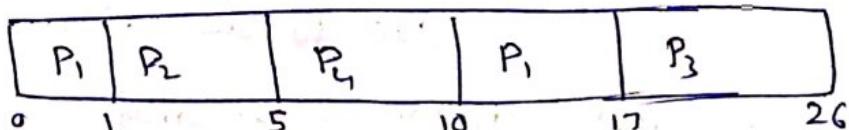
↳ A process with a shorter CPU burst can preempt a running process

→ We decide when

↳ A new process arrives

Ex:-

Process	P ₁	P ₂	P ₃	P ₄
Arrival time	0	1	2	3
CPU Burst	8	4	9	5

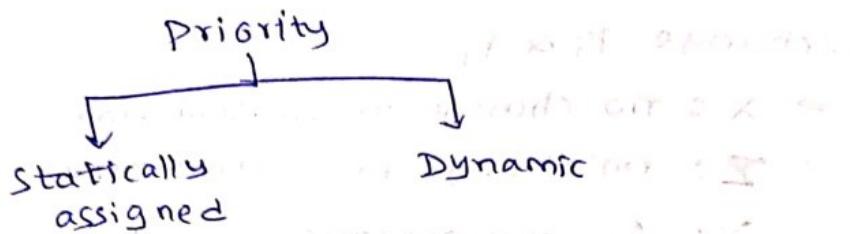


$$AWT = \frac{9+0+15+2}{4} = 6.5$$

(d) Priority Scheduling

→ We associate a priority value with each process

→ At any instance of time, CPU is assigned to the process with highest priority



→ Static priority

→ Many versions of Linux

→ Each process is assigned a priority value

-20 → highest priority

+19 → ~~lowest~~ lowest priority

↳ default priority = 10.

→ Commands

→ nice → decreases the priority of a process

→ renice → increases " "

(only with "root" access)

→ Problem

Starvation or indefinite waiting

A process is waiting indefinitely due to low priority

→ Dynamically changing priority,

→ Try to overcome starvation problem.

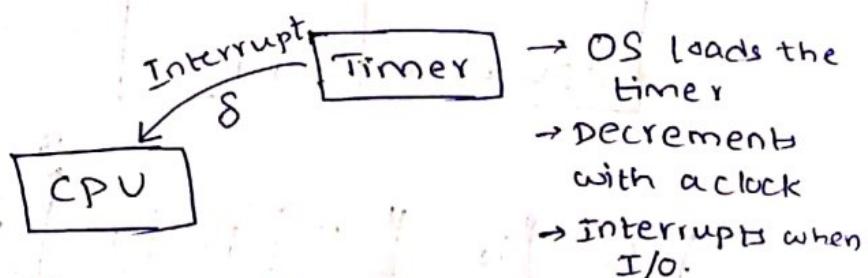
↓
As a process waits, its priority increases

- Response Ratio, $RR = \frac{\text{Time since arrival} + \text{CPU burst time}}{\text{CPU burst time}}$
 jobs arriving first, but are being delayed have high RR.

- RR can indicate priority
- whenever there's a clash of priority, we can choose the one with highest RR

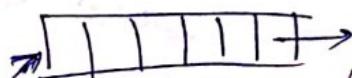
(e) Round-robin scheduling

- Designed primarily for interactive time-stationary systems.
- let's say a small time quantum, δ is predefined ($\delta \approx 1-100\text{ ms}$).

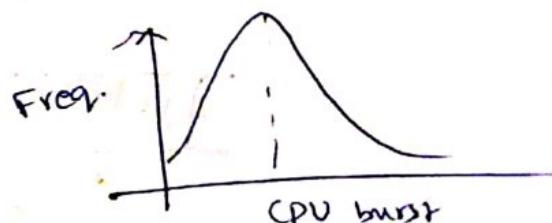


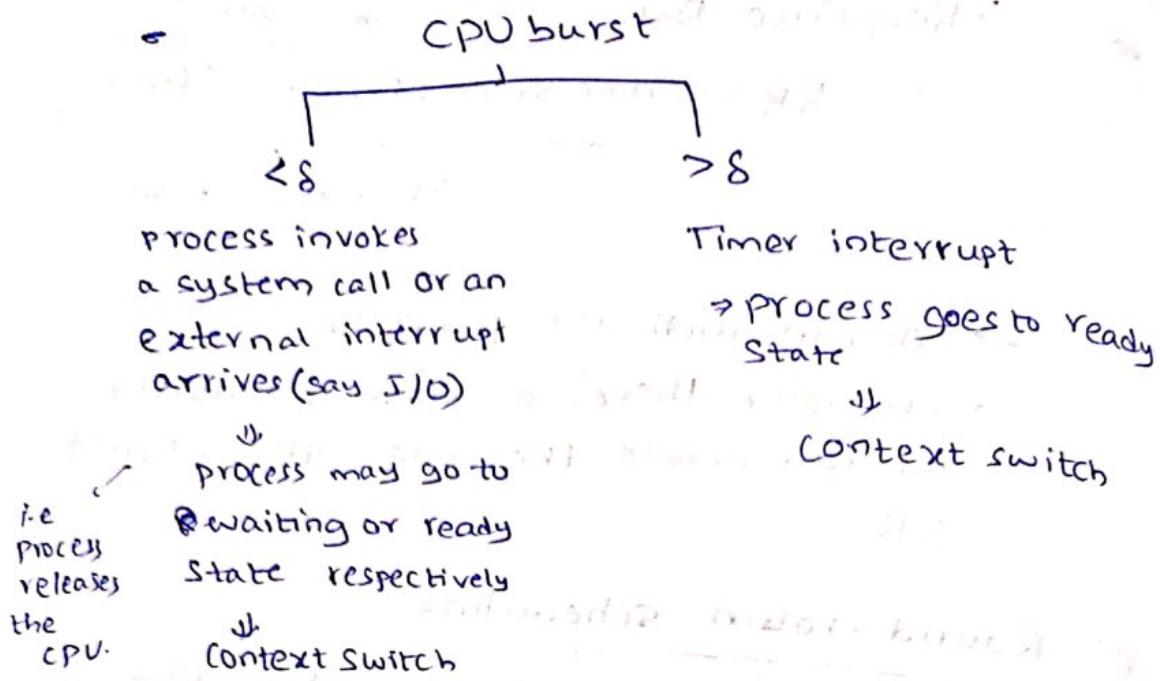
Ready list

→ stored as circular queue.



→ The instruction to set the timer is a privileged instruction.

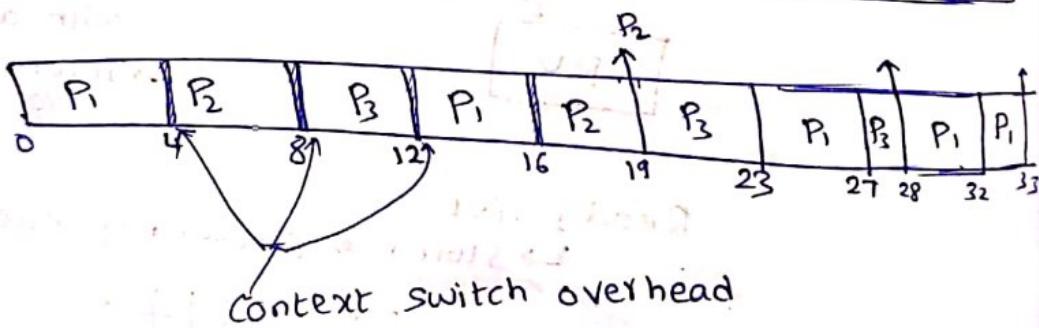




→ Example:

$$\delta = 4 \text{ ms}$$

Process	P ₁	P ₂	P ₃
Arrival Time	0	0	0
CPU burst	17	7	9

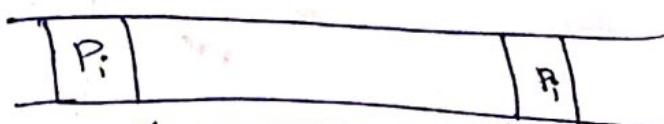


→ A Simple Analysis

→ n Processes in READY queue.

→ δ is the time quantum.

→ σ is the context switch overhead.



$(n-1)(\delta + \sigma)$ ← maximum waiting time of a process before it gets back to CPU

Very high $\delta \Rightarrow$ similar to FCFS.

Very small $\delta \Rightarrow$ large no. of context switches \Rightarrow overhead keeps adding up.

$\delta > \sigma$
typically 10 ms typically 10 μs.

Q.	Process	P ₁	P ₂	P ₃	P ₄
Arrival time	0	0	0	0	0
CPUBurst	6	3	1	7	

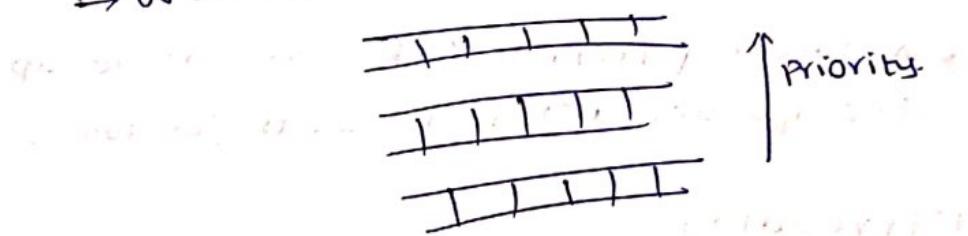
compute AWT for $\delta = 1, 2, 3, 4, 5, 6, 7$.

Ans we get graph like



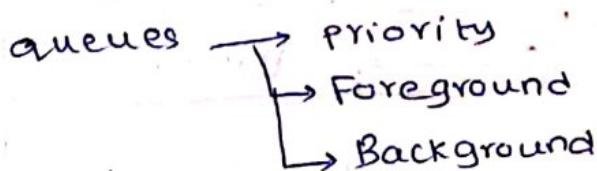
(f) Multi-level queue scheduling

→ We'll have multiple Ready queues



→ Multiple ready queues with different priorities and scheduling ~~processes~~ policies.

→ A process is assigned to one of the queues

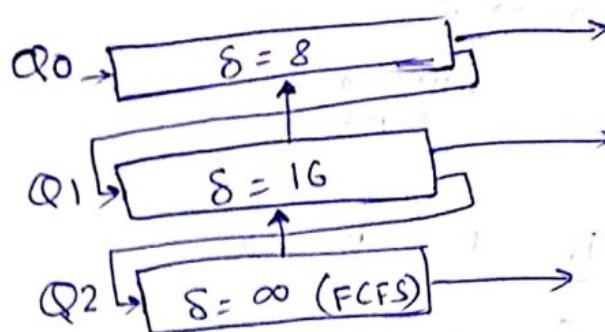


→ processes do not move across queues.

↳ It can lead to starvation.

(G) Multilevel feedback queue scheduling

→ 3 queues



→ Q₀ has the highest priority & Q₂ has the lowest priority.

→ A new process is assigned to Q₀.

→ Current running process is in Q₀

→ & CPU burst < 8 ms

Then, the process remains in Q₀.

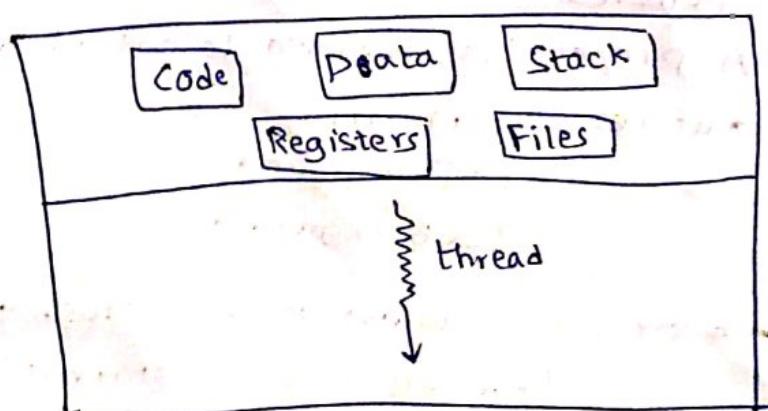
if CPU burst > 8 ms,

Move the process to Q₁:

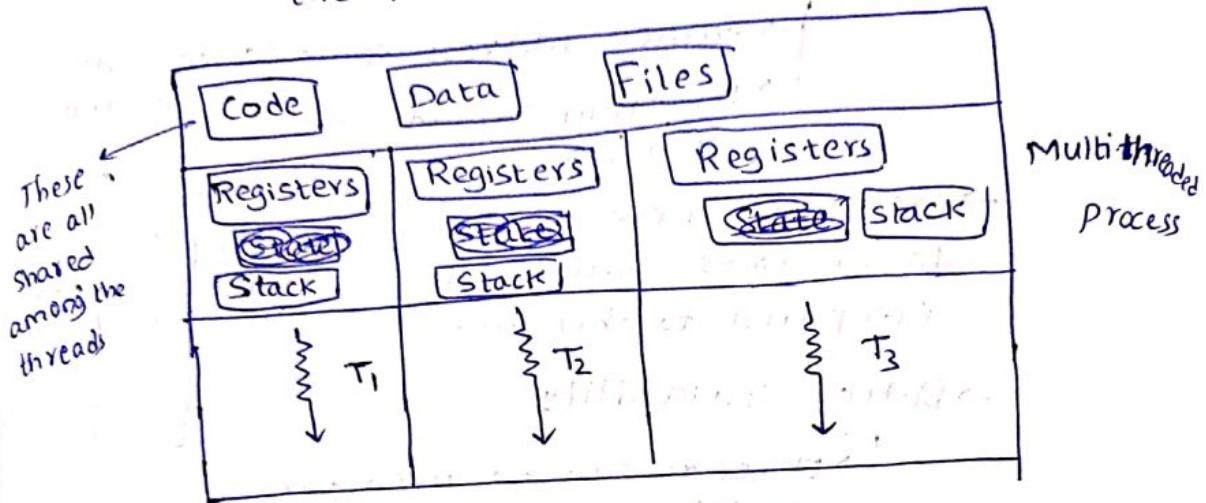
→ A lower priority job can move up the queues also if it waits for too long.

* Multithreading

→ A process is a program in execution with a single thread of control.

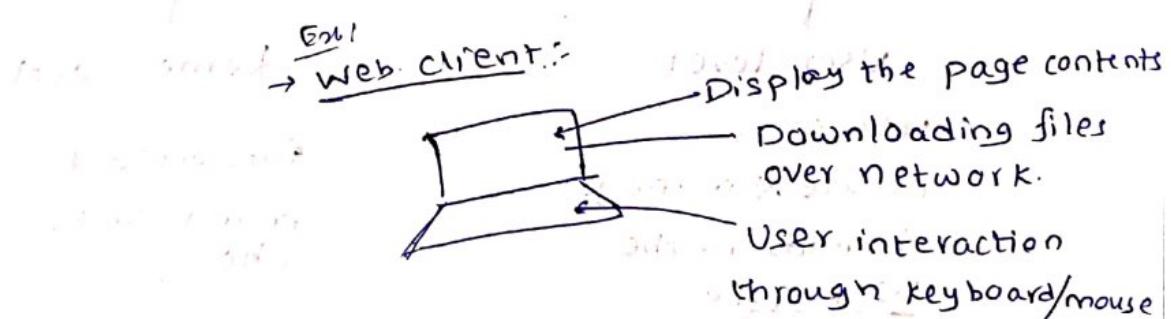


- What is a thread?
 - Also referred to as light-weight process
 - Contains a thread ID, Program Counter (PC), registers and stack.
 - Shares the code, data, files etc. with the peer threads.

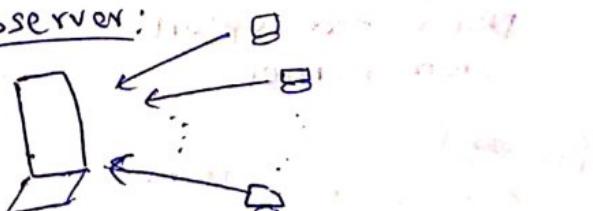


- Why do we require multiple threads?

→ Context switch overhead is significantly reduced. [since only Registers & Stack need to be saved, remaining all are same].



- Webserver:



→ Multithreading benefits

- context switch among threads is faster.
- Better responsiveness in interactive applications.
- Better resource utilization among peer threads
 - ↳ same address space (since code is shared)
 - ↳ different threads of activity
- creating a new process is 10-30 times slower as compared to threads.
- Better scalability
 - ↳ threads can be used for multiprocessor environments.

→ Multithreading models

Types of threads

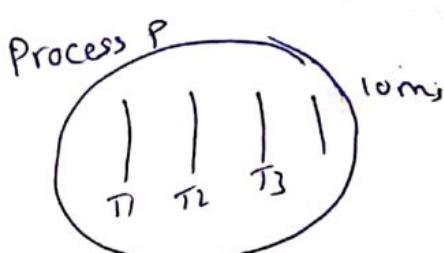
User-level

use APIs to create & manage threads in the user space

Don't need support from kernel.

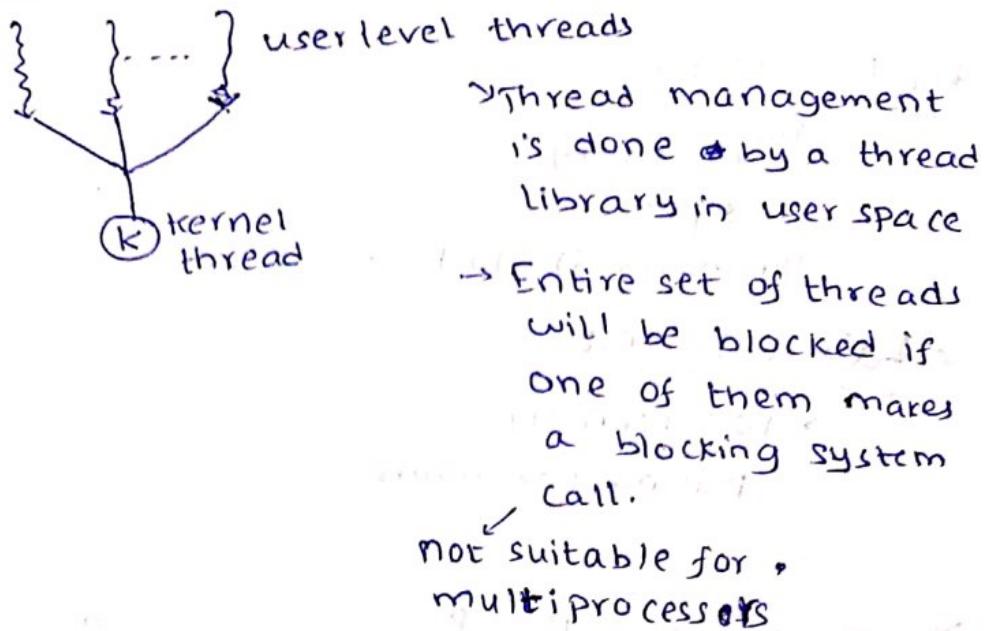
Kernel-level

Supported & managed by the OS.

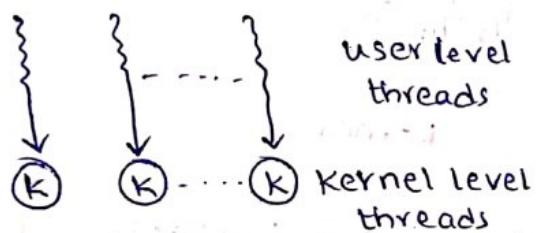


If a thread makes a blocking system call to the kernel, whole process is blocked since the kernel doesn't know about threads. It only sees it as a process

(a) Many-to-one model:-

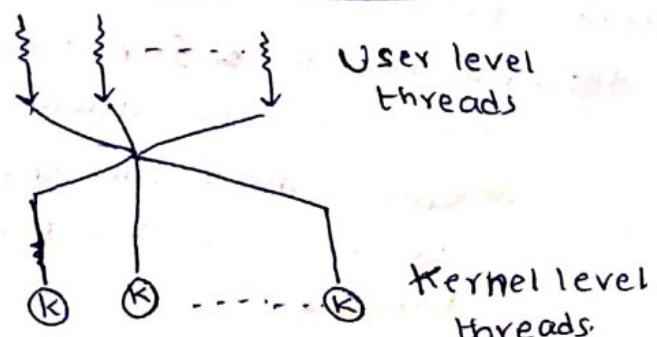


(b) One-to-One Model :-



- provides more concurrency
- When ~~one~~ a thread blocks, others can run.
- Additional overhead w.r.t. kernel.
[∴ kernel has to manage many kernel threads]
- Linux / Windows follows this.
- can be extended to multi processors

(c) Many-to-many Model :-



→ Solaris OS used this model.

* Class Test

→ 7th February

7-8 PM

Rooms: CSE-107 & CSE-108

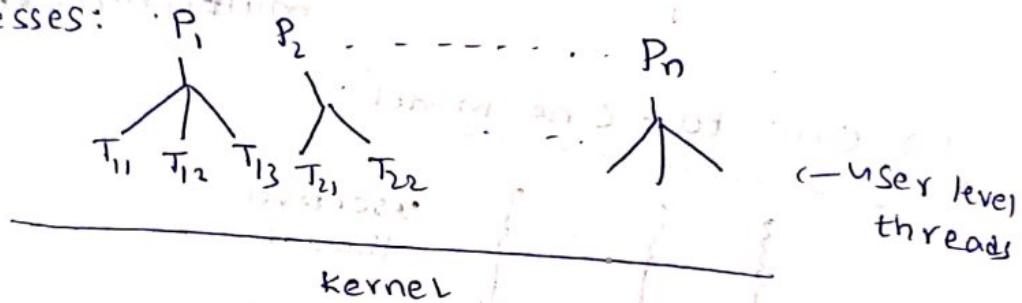
* Lab Quiz Test

→ 10th February

During Lab hours.



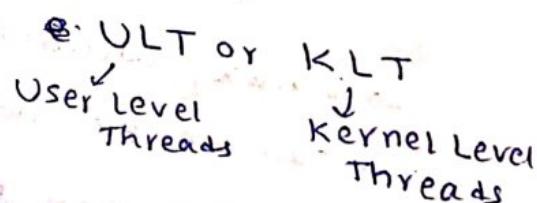
Processes:



- Irrespective of how many threads a process creates, other process's runtime is not hampered.
- Allocating time to threads can be done using thread library.

* POSIX Thread Library (Pthread Library)

→ Define the threads as



- All global data are shared among threads.
- Data defined inside functions → allocated on stack
- Since stacks are separate, separate copies are created for threads.

```

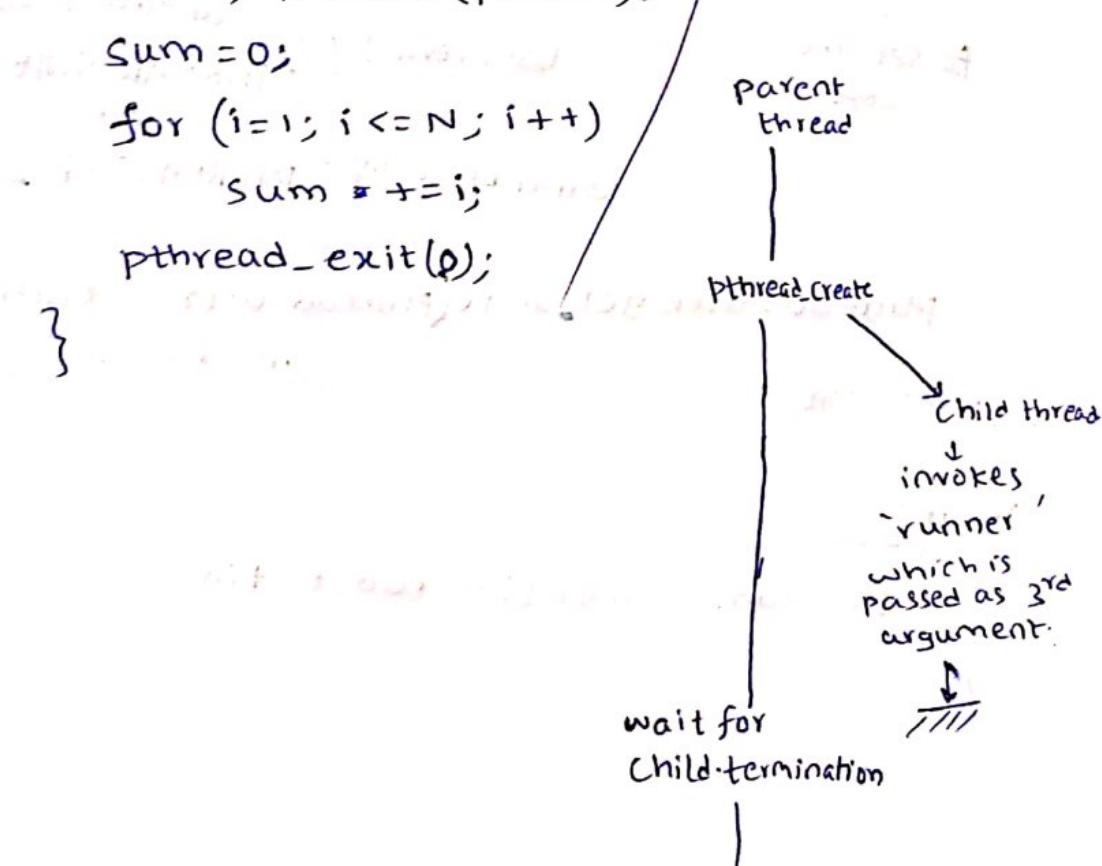
#include<stdio.h>
#include <pthread.h>
int sum;
void *runner(void *param);
int main(int argc, char argv[])
{
    pthread_t tid;
    pthread_attr_t attr;
    pthread_attr_init(&attr);           → retrieving the
    pthread_create(&tid, &attr, runner, argv[1]); → default attributes
    pthread_join(tid, NULL);          → creates
    printf("\nsum is %d", sum);       → a thread
}

```

```

void *runner(void *param);
{
    int i, N = atoi(param);
    sum = 0;
    for (i=1; i <= N; i++)
        sum += i;
    pthread_exit(0);
}

```



* Creating multiple threads

```
#define NTHR 10  
pthread_t mythreads[NTHR]  
{  
    for(int i=0; i<NTHR; ++i)  
        pthread_create(&mythreads[i], &attr,  
                       runner, NULL);  
    for(int i=0; i<NTHR; ++i)  
        pthread_join(mythreads[i], NULL);
```

* Thread scheduling using pthread

→

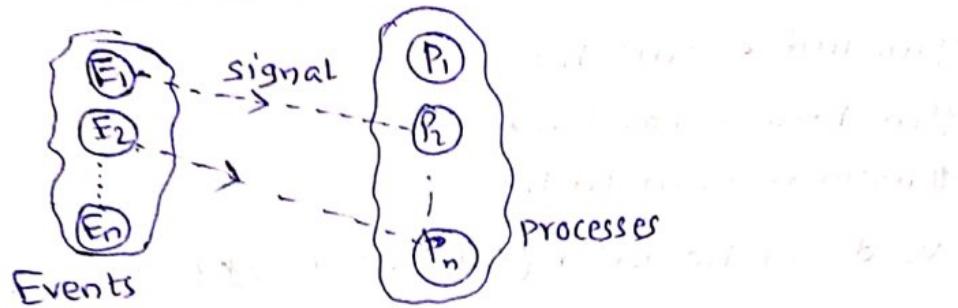
```
pthread_attr_set_scope(pthread_attr_t *attr, int scope);  
to set the scope  
USER LEVEL Thread can take values PTHREAD_SCOPE_PROCESS (only)  
Kernel Level Thread PTHREAD_SCOPE_SYSTEM
```

pthread_attr_get_scope(pthread_attr_t *attr, int *scope);
to get the scope

* Thread Cancelling

```
pthread_cancel(pthread_t tid);
```

*Signal Handling



- A signal notifies a process that some event has occurred.
- An event occurs.
 - A signal corresponding to the event is generated
 - Signal is delivered to a process
 - Once delivered, the signal must be handled appropriately by the process.

- Some of the POSIX Signals

→ SIGABRT → Abort

→ SIGBUS → Bus error

→ SIGILL → Illegal instruction

→ SIGKILL → Kills a process

→ SIGQUIT → terminal quit

→ SIGSEGV → invalid memory reference
(Segmentation fault)

→ SIGUSR1 } → User-defined

→ SIGUSR2 }

→ SIGINT → CTRL + C

* Signal handler in C

```
#include <stdio.h>
#include <signal.h>
#include <unistd.h>
void sig_handler(int sig_no){
    if (sig_no == SIGINT)
        printf("\n Received signal");
}
int main(){
    signal(SIGINT, sig_handler);
    while(1)
        sleep(1);
}
```

↑
install the
signal handler

* Assign - 1:

Batch processing system:-

M : average mounting time

↑ to load the
jobs in
Hard disk

T : average service time per job.

N : number of jobs

S : unit price of service time

W : unit price of waiting time per user.

Find optimum batch size (N) that minimizes total cost.

$$\text{Sol} \quad \text{Total time} = M + NT$$

$$\text{cost per unit time} = S + NW$$

$$\frac{\text{Total cost}}{\text{user}} = \frac{(M+NT)(S+NW)}{N}$$

$$C = \frac{MS + MNW + NTS + N^2WT}{N}$$

$$C = \frac{MS}{N} + \alpha_{WT} N + (\alpha_M + \alpha_S)$$

$$\frac{dC}{dt} = 0 \Rightarrow C = \text{const}$$

✓ Sending signal to a process or thread

→ For process, kill (pid - t, pid, int signal);

→ FOR threads, *join* (used *return*) and *wait* (*join* in signal):

THE VENDEE - HISTOIRE DE LA VENDEE (1793-1802) -

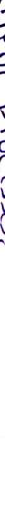
* Process Synchronization ?

→ Multiple process execution parallel execution parallel

Diagram illustrating concurrent processes:

- CPU 1** and **CPU m** are connected by a dashed line labeled **C**.
- The label **concurrent process** is written vertically between the two CPU units.

R_1, R_2, \dots, R_n
Sharing CPU time
(time sharing also)
 R_1, R_2, \dots, R_n
parallel process execution.

* Concurrent processes:

 P₁ P₂ P₃ → context switch (timer interrupt)
 ↓ ↓ ↓ → I/O interrupt (asynchronous)

Ex: Producer - consumer problem

Bounded buffer

```
#define BUFSIZE 50
```

```
typedef struct {  
    void *ptr;  
    size_t size;  
} mem_t;
```

— — — od = force, μ = friction coefficient, α = angle.

Journal of the American Mathematical Society 20 (1987), 371-387.

buffer: [BUFSIZE]:

```

int in=0, out=0
int count=0;
    } → shared
    Memory

```

No. of items in buffer

Producer P

while(1){

 next-prod = ProduceItem();

 while (count == BUF_SIZE);

 buffer[in] = next-prod;

 in = (in + 1) % BUFSIZE;

 Count++;

 1. LOAD R1, 0(count)

 2. ADD R1, R1, 1

 3. STORE R1, 0(count)

Consumer C

while(1){

 while (count == 0);

 next-cons = buffer[out];

 out = (out + 1) % BUFSIZE;

 Count--;

 4. LOAD R5, 0(count)

 5. SUB R5, R5, 1

 6. STORE R5, 0(count)

→ Even if we have 1P & 1C, we don't have problem with in or out but count has the value of count is updated by both of them. ⇒ a problem.

→ suppose initially, count = 10

↳ Both P & C run one cycle each.

Scenarios

1, 2, 3, 4, 5, 6

4, 5, 6, 1, 2, 3

1, 2, 4, 5, 6, 3

4, 5, 1, 2, 3, 6

4, 1, 2, 3, 5, 6

1, 4, 5, 6, 2, 3

Final count

10 (initial)

10 (initial)

11 (initial)

9 (initial)

9 (initial)

11 (initial)

→ RACE condition

* Critical Section Problem :-

→ There are 'n' concurrent processes

P_1, P_2, \dots, P_n are n processes

↳ Each P_i has a segment of code where some shared data is updated.

(e.g. of printing simple numbers, file, etc.)

Writing P_i consists of states 2222222222 2222222222

→ P_i starts in a process section (RS) and then enters a critical section (CS).

After leaving CS, it exits RS and continues in a normal state.

→ Structure of P_i is as follows:

→ When a process P_i is executing in its critical section (CS), no other process P_j should be allowed to enter the CS.

Execution of CS code is mutually excluded

in time by mutual exclusion

Also called Mutual Exclusion Problem.

→ Structure of P_i :

while(1) {

 looping

 RS

 <Entry Section>

 CS

 <Exit Section>

 else part } (written in it's own part)

 } (writing of structure with P)

→ When a process enters at the beginning of CS, it must be prevented from doing so

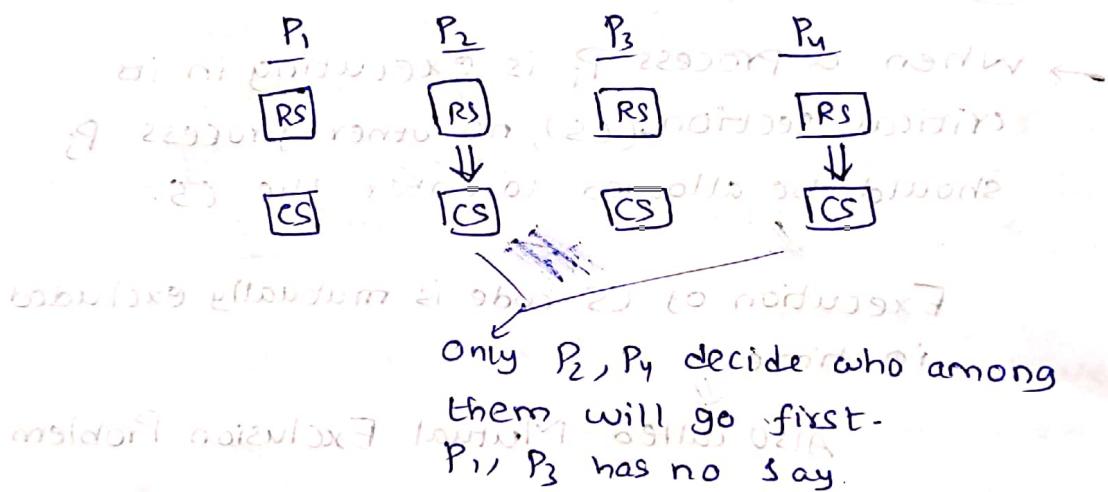
→ If it is allowed to do so, then

* Basic requirements of any correct solution to the critical section problem:

(a) Mutual Exclusion:

No two processes must be inside their CS at the same time.

(b) Progress: If no process is executing in CS, and some processes wish to enter critical section, then only those processes that are not in their RS's, must participate to decide who will enter CS next. This section cannot be defined indefinitely.



(c) Bounded Waiting: There exists a limit on the number of times other processes are allowed to enter their critical sections after a process has made a request to enter its CS & the request is granted.

→ Let's say P₁, P₂, P₃ are visiting CS frequently. P₄ also wants to enter CS.

It should not be the case that P₄ will be always denied in favour of P₁, P₂, P₃.

* Software Solution :- [Peterson's Algorithm]

→ Restrict to two concurrent processes only.

```
int turn=0; //initialization  
boolean flag[2];
```

→ Shared memory.

P₀

```
while(1){  
    Entry section {  
        flag[0]=true;  
        turn=1;  
        while(flag[1]&&turn==1);  
    }  
    CS {  
        RS {  
            flag[0]=false;  
        }  
    }  
}
```

P₁

```
while(1){
```

```
    flag[1]=true;  
    turn=0;  
    while(flag[0]&&turn==0);
```

CS

```
    flag[1]=false;  
    RS {  
        }  
    }
```

→ flag[0] → willingness of P₀ to enter CS

→ flag[1] → willingness of P₁ to enter CS

→ turn → specifies who will enter next

* Modern computer Architectures

→ Typical execution sequence

P_i: turn = 0 → flag[1] = false

P_o: turn = 1 → flag[0] = true

P_o: flag[0] = true → flag[0]

P_o: flag[1] && turn == 1 → flag[1]

⇒ P_o enters CS → turn = [0] (Ques??)

P_i: flag[1] = true → turn = [1] (Ques??)

P_o: flag[0] && turn == 0 → WRONG

⇒ P_o enters CS → turn = [0] (Ques??)

This error occurs for those architectures which can trigger out-of-order execution.

* Hardware Solution for Synchronization:

lock = false shared

false → unlocked, true → locked

P_i

while(1) {

 acquire(&lock);

 CS

 release(&lock);

check if the lock is available, if not, it waits, else it gets lock = true;

Must be atomic..

They must be either executed completely or not executed at all.

(a) Disabling & Enabling Interrupts:

Pi

```
while(1) {
```

DI

CS

EI

RS

we are disabling interrupts before the CS. So in this sense, we are making CS execute like it's atomic.

↓ But not so good.

}

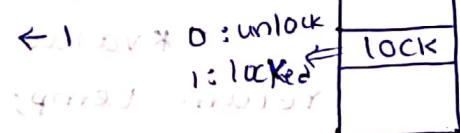
↑ If it fails

(b) Test-and-set. Instructions

TAS R1, 0(1000)

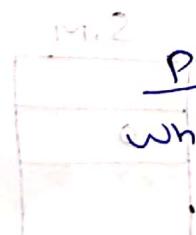
III: $R_1 \leftarrow M[1000]$; unlock

$M[1000] \leftarrow R_1$; lock



As a pseudocode,

```
boolean test_and_set(boolean *target)
{
    boolean rv = *target;
    if (*target == TRUE)
        return rv;
    else
        *target = TRUE;
    return rv;
}
```



while(1) {

. while(test_and_set(&lock));

CS

lock = FALSE;

RS

can be used for any number of processes.

also V7 commands ~ 1960

(C) Compare and swap instruction

Ex:

CAS R₁, R₂, 0(1000)

if (R₁ == M[1000])
M[1000] = R₂

→ Pseudo code → int CAS
→ Process P_i

while(1){

→ Pseudo code

int CAS(int *value, int expected, int new)

{

int temp = *value;

if (*value == expected)

*value = new;

return temp;

}

→ Process P_i

while(1){ Initially, lock=0.

if(CAS(&lock, 0, 1) != 0);

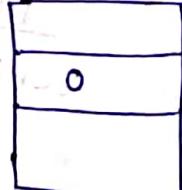
Called
spin lock

CS

LOCK=0;

RS

S.M.

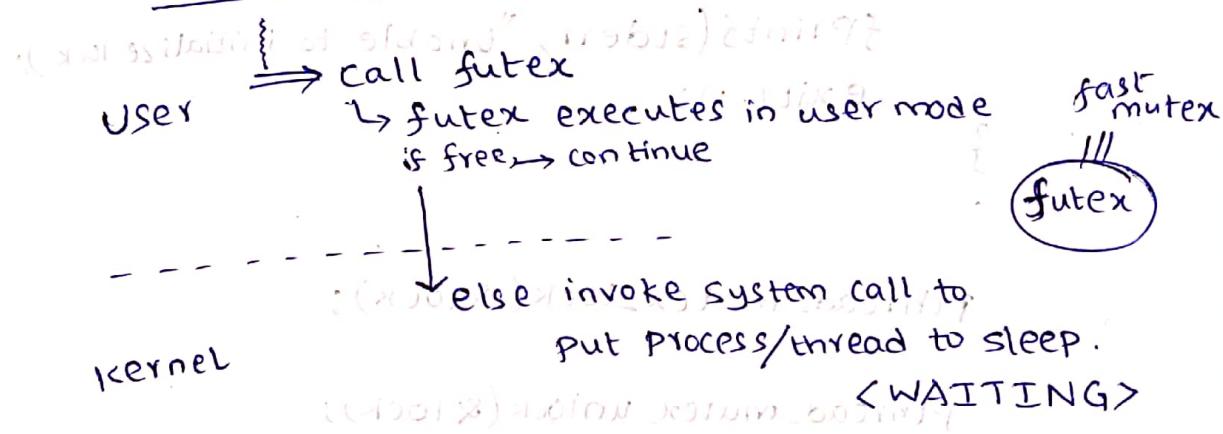


* Spin lock

The process that is waiting in the entry section is actually running in a loop → consuming CPU cycles.

- * Mutex locks :- (about 8-10) principles :-
→ simple mechanism for locking & unlocking.
- You require one of the special hardware instructions (multiple general purpose)
- You may or may not wait on spin lock (if the CS is occupied).

- Typical mutex implementation that avoids spinlock (about 10-12)



- Spinlocks
 → Drawback :- It wastes CPU time.

- Advantages :-
 (a) It avoids context switch during waiting. → To put the process to waiting state.
 (b) May be good if the CS code is fairly small.

- * acquire() and release() :-

```

    P_i
    while(1) {
        acquire(&lock);      Spinlock or context switch
                            may occur.
        CS
        release(&lock);
        RS
    }
  
```

```

→ acquire(int *lock)           release(int *lock)
    {
        while(*lock != 0);      *lock = 0;
    }
}

no spin lock → no waiting
no race condition
when using spin lock.

```

~~no waiting no race for when no spin lock~~

→ Mutex lock for pthread library:

```

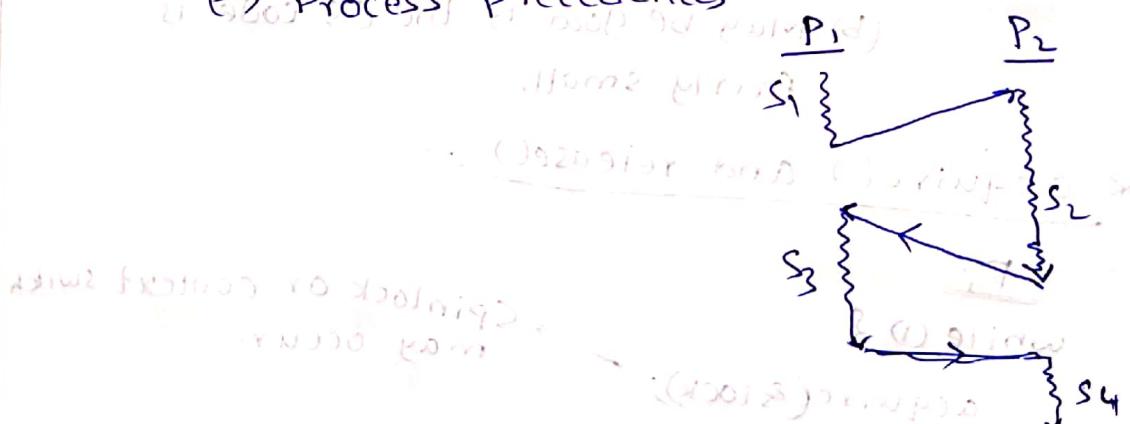
pthread_mutex_t lock;
if(pthread_mutex_init(&lock, NULL) != 0){
    fprintf(stderr, "Unable to initialize lock");
    exit(-1);
}
pthread_mutex_lock(&lock);
pthread_mutex_unlock(&lock);
if(pthread_mutex_destroy(&lock) != 0)
{
    fprintf(stderr, "Error destroying lock");
    exit(-1);
}

```

* Process synchronization

(a) Mutual exclusion

(b) Process precedences



(c) Resource counting

count++ → atomic instruction

