

Question 1

1. How many lines does the RDD contain?

Answer:- 9669739

Method:- We create an RDD that is a collection of the lines in the file. Then we count the number of items in it.

2. Count the number of “WARN” messages.

Answer:- 132158

Method:- We create an RDD that is a collection of *case classes* that represent a line in the log file. Then filter those messages whose *debug_level* is WARN.

3. How many repositories were processed in total when the retrieval_stage is “api_client” ?

Answer:- 71981

Method:- We create an RDD that is a collection of *case classes* that represent a line in the log file. Then filter those messages whose *retrieval_stage* is api_client.rb. Then we process the repository name for each entry from the *rest* field. After that, we use MapReduce with key as the repository name to find the number of unique repositories processed.

4. Using retrieval_stage as “api_client”, find which clients did the most HTTP requests and FAILED HTTP requests from the download_id field.

Answer:- ghtorrent-13(85528), ghtorrent-13(79623)

Method:- We create an RDD that is a collection of *case classes* that represent a line in the log file. Then filter those messages whose *retrieval_stage* is api_client.rb and check whether the request is failed or not from the *rest* field. Then we count the number of items in the resultant RDD.

5. Find the most active hour of the day and most active repository.

Answer:- 10(2662487), greatfakeman/Tabchi(79524)

Method:- We create an RDD that is a collection of *case classes* that represent a line in the log file. After that, we use MapReduce, with the key as the hour (from *timestamp*) or repository name (from *rest*) and count the number of entries for each hour/repo. After that we find the key with the maximum value.

6. Which access key is failing most often?

Answer:- ac6168f8776 (79623)

Method:- We create an RDD that is a collection of *case classes* that represent a line in the log file. After that, we use MapReduce, with the key as the access key (from *rest*) and count the number of entries for each key. After that we find the key with the maximum value.