Distributed Mutual Exclusion

Mutual Exclusion

- Well-understood in shared memory systems
- Requirements:
  - at most one process in critical section (safety)
  - if more than one requesting process, someone enters (liveness)
  - a requesting process enters within a finite time (no starvation)
  - requests are granted in some order (fairness)

Some Complexity Measures

- No. of messages per critical section entry

- Synchronization delay

- Response time

- Throughput

Classification of Distributed Mutual Exclusion Algorithms

- **Permission based**
  - Node takes permission from all/subset of other nodes before entering critical section
  - Permission from all: costly, good for small systems
  - Permission from subset: scalable, widely used
    - Main problem: How to choose the subsets for each node?
- **Token based**
  - Single token in the system
  - Node enters critical section if it has the token
  - Algorithms differ in how the token is circulated among requesting nodes

Adaptive vs. Non-Adaptive

- Performance under low load (less requests for CS) should be better than performance under high load (lots of requests)

- **Adaptive** mutual exclusion algorithms: performance is dependent on load

Permission based Algorithms

Lamport's Algorithm

- Permission from all
- Given as an example of use of logical clocks for ordering requests
- Model: asynchronous system, completely connected topology, reliable FIFO channels
- Better algorithms exist even in the class of "permission from all", but good to study first as a base to understand those

- Every node i has a request queue q_i, keeps requests sorted by logical timestamps (total ordering enforced by including process id in the timestamps)
- To request critical section:
  - send timestamped REQUEST (ts_i, i) to all other nodes
  - put (ts_i, i) in its own queue
- On receiving a request (ts_i, i):
  - send timestamped REPLY to the requesting node i
  - put request (ts_i, i) in the queue

- To enter critical section:
  - i enters critical section if (ts_i, i) is at the top if its own queue, and i has received a message (any message) with timestamp larger than (ts_i, i) from ALL other nodes.

- To release critical section:
  - i removes its request from its own queue and sends a timestamped RELEASE message to all other nodes
  - On receiving a RELEASE message from i, i's request is removed from the local request queue

Some Observations

- Timestamp of REPLY message > timestamp of REQUEST message it is replying to

- Purpose of REPLY messages from node i to j is to ensure that j knows of all requests of i prior to sending the REPLY (and therefore, possibly any request of i with timestamp lower than j's request)

- Requests are granted in order of increasing timestamps

- Requires FIFO channels.

- $3(n-1)$ messages per critical section invocation

- Synchronization delay $=$ max. message transmission time

Ricart-Agarwala Algorithm

- Improvement over Lamport's
- Main Idea:
  - node j need not send a REPLY to node i if j has a request with timestamp lower than the request of i (since i cannot enter before j anyway in this case)
    - REPLY sent after j satisfies its own request
    - No separate RELEASE message needed
- Requests still granted in order of increasing timestamps
- Does not require FIFO
- Still requires knowledge of all other nodes in the network
- $2(n-1)$ messages per critical section invocation
- Synchronization delay = max. message transmission time

- To request critical section:
  - send timestamped REQUEST message (ts_i, i)
- On receiving request (ts_i, i) at j:
  - send REPLY to i if j is neither requesting nor executing critical section or if j is requesting and i's request timestamp is smaller than j's request timestamp. Otherwise, defer the request.
- To enter critical section:
  - i enters critical section on receiving REPLY from all nodes
- To release critical section:
  - send REPLY to all deferred requests

Maekawa's Algorithm

- Permission obtained from only a subset of other processes, called the Request Set (or Quorum)

- Separate Request Set R_i for each process i

- Requirements:
  - for all i, j: $R_i \cap R_j \neq \Phi$
  - for all i: $i \in R_i$
  - for all i: $|R_i| = K$, for some K
  - any node i is contained in exactly D Request Sets, for some D

- K = D (easy to see)

- For minimum K, K ≈ sqrt(N) (why?)

A simple version

- To request critical section:
  - i sends REQUEST message to all processes in R_i
- On receiving a REQUEST message:
  - send a REPLY message if no REPLY message has been sent since the last RELEASE message is received. Update status to indicate that a REPLY has been sent. Otherwise, queue up the REQUEST
- To enter critical section:
  - i enters critical section after receiving REPLY from all nodes in R_i

- To release critical section:
  - send RELEASE message to all nodes in R_i
  - On receiving a RELEASE message, send REPLY to next node in queue and delete the node from the queue. If queue is empty, update status to indicate no REPLY message has been sent since last RELEASE is received.

- Message Complexity: 3*sqrt(N)
- Synchronization delay =

    2 *(max message transmission time)

- Major problem: Deadlock possible
- Can you update the protocol with additional messages to solve this problem?
  - Good practice ☺
  - Maekawa's protocol already does that, we just looked at a part of it
- Building the request sets?

Some Observations

- Permission based algorithms with permission from a subset are widely used
  - Voting/Quorum based protocols
    - In Maekawa's algorithm,
      - each process has one vote
      - A process needs a certain number of votes to proceed
- Questions/Issues
  - How to choose the quorums?
  - Should the quorum be the same for read and write?
  - Should each process have one vote only? Same number of votes for all?
  - Dynamic quorums/votes

Token based Algorithms

Token based Algorithms

- Single token circulates, enter CS when token is present
- Mutual exclusion obvious
- Algorithms differ in how to find and get the token
  - Token circulates, nodes use it when it passes through them
  - Token stays at node of last use, other nodes request for it when needed
    - Need to differentiate between old and current requests

Suzuki Kasami Algorithm

- Broadcast a request for the token
- Process with the token sends it to the requestor if it does not need it

Issues:

- Current vs. outdated requests
- Determining sites with pending requests
- Deciding which site to give the token to

- The token:
  - Queue (FIFO) Q of requesting processes
  - LN[1..n] : sequence number of request that j executed most recently

- The request message:
  - REQUEST(i, k): request message from node i for its k^{th} critical section execution

- Other data structures
  - $RN_i[1..n]$ for each node i, where $RN_i[j]$ is the largest sequence number received so far by i in a REQUEST message from j.

- To request critical section:
  - If i does not have token, increment $RN_i[i]$ and send REQUEST(i, $RN_i[i]$) to all nodes
  - if i has token already, enter critical section if the token is idle (no pending requests in token Q), else follow rule to release critical section

- On receiving REQUEST(i, sn) at j:
  - set $RN_j[i] = \max(RN_j[i], sn)$
  - if j has the token and the token is idle, send it to i if $RN_j[i] = LN[i] + 1$. If token is not idle, follow rule to release critical section

- To enter critical section:
  - enter CS if token received after sending REQUEST

- To release critical section:
  - set $LN[i] = RN_i[i]$
  - For every node j which is not in Q (in token), add node j to Q if $RN_i[\,j\,] = LN[\,j\,] + 1$
  - If Q is non empty after the above, delete first node from Q and send the token to that node

Points to note:

- No. of messages: 0 if node holds the token already and token is idle, n otherwise
- Synchronization delay: 0 (node has the token) or max. message delay (token is elsewhere)
- No starvation