

# Accelerating large graph algorithms on GPU using CUDA

Archisman Pathak  
Koushik Raj  
Koustav Chowdhury  
Rounak Patra

Satyam Porwal  
Siddhant Agarwal  
Sriyash Poddar



# Table of Contents



Problem & Motivation



Breadth First Search



Single Source  
Shortest Path




All Pair  
Shortest Path






**01**

# **Problem and Motivation**



# Understanding the Problem

...



## Graph Algorithms

Graph algorithms are used to develop intelligent solutions and enhance various machine learning models.

...




## Impractical Sequential Algorithms

Fast implementations of sequential graph algorithms are fast, but the hardware used in them is very expensive.

...

# Understanding the Motivation


...



## CUDA

Nvidia CUDA provides a development environment for creating high performance GPU-accelerated applications.

...




## Previous Works in Parallel Graph Algorithms

Previous works in parallel graph algorithms achieved practical times on basic graph operations but at high hardware cost.

...

# CUDA Graph Representation



- Each value of posV contains an index of the posE array
- The pointer contains the number of neighbours of the node, say  $n$ .
- The next  $n$  elements are the ids of neighbouring vertices.

# Graphs used for Experiments

## APSP Graphs

(Due to time and memory constraints restricted the size)

\*Graphs are numbered in increasing number of vertices

Graph Number	# Nodes	# Edges	Average Degree
1	2,700	1,808,853	1,340
2	4,039	88,234	44
3	4412	108,818	49
4	5,881	21,492	7
5	7,500	837,083	233

# Graphs used for Experiments

BFS and SSSP Graphs  
(No constraints on the size)


\*Graphs are numbered in increasing number of vertices

Graph Number	# Nodes	# Edges	Average Degree
1	2,700	1,808,853	1,340
2	21,853	12,323,648	1,128
3	82,168	504,230	12
4	1,694,616	11,094,209	13
5	16,777,214	132,557,200	16





# BREADTH FIRST SEARCH





# Our Implementations

**Parallel BFS**

**Queue BFS**

**Scan BFS**

# First Approach: Parallel BFS

## Host

- Level-synchronous approach
- Runs  $O(V^2 + E)$  operations
- *Vertex frontiers* : Nodes that are currently being visited
- *Edge frontiers* : Nodes that will be visited in the next iteration
- *flag* checks for termination

...

---

**Algorithm 1** parallelBFS\_Host

---

```
1: Input:  $V_a, E_a, S$  ▷ The graph  $G(V, E)$  and source  $S$   
2: Create distance array  $Dist_a$ , and parent array  $P_a$  of size  $|V|$   
3: Initialise all elements of  $Dist_a, P_a$  to  $\infty$   
4:  $D_a[S] = 0$   
5:  $level = 0$   
6:  $flag = True$   
7: while  $flag$  do  
8:    $flag = False$   
9:   Invoke parallelBFS_kernel( $level, V_a, E_a, Dist_a, flag$ ).  
10:   $level = level + 1$ 
```

---

...

...

# First Approach: Parallel BFS

## Kernel

- Threads are launched for each vertex
- Each vertex checks if it is *frontier vertex*
- If yes, it updates the distances of neighbours and populates the *edge frontiers*
- Terminates when no *frontier vertex* updates its neighbour

...

---

### Algorithm 2 parallelBFS\_kernel

---

```
1: Input:  $level, V_a, E_a, Dist_a, flag$ 
2:  $tid = \text{getThreadID}$ 
3:  $f = \text{False}$ 
4: if  $tid < V_{a\_size}$  and  $Dist_a[tid] = level$  then
5:    $u = tid$ 
6:   for all  $v = \text{neighbours of } u$  do
7:     if  $level + 1 < Dist_a[v]$  then
8:        $Dist_a[v] = level + 1$ 
9:        $f = \text{True}$ 
10:  if  $f = \text{True}$  then
11:     $flag = \text{True}$ 
```

---

...

...

# Second Approach: Queue BFS

## Host

- Level-synchronous
- Runs  $O(V + E)$  operations
- *Vertex frontier* and *edge frontier* is maintained in form of a queue
- Intuition is similar to sequential BFS.
- Terminates when there are no *vertex frontiers*.

---

### Algorithm 3 queueBFS\_Host

---

```
1: Input:  $V_a, E_a, S$  ▷ The graph  $G(V, E)$  and source  $S$   
2: Create cost array  $Dist_a$  and parent array  $P_a$  of size  $|V|$  and  
   initialise all values to  $\infty$   
3: Create two array  $cQ$  and  $nQ$ , and initialise it to  $S$  and null  
   respectively.  
4:  $Dist_a[S] = 0$   
5:  $P_a[S] = -1$   
6:  $l = 0$  ▷ Start with the source vertex  
7: while  $cQ_{size} > 0$  do  
8:   Invoke queueBFS( $l, V_a, E_a, Dist_a, P_a, cQ, nQ$ )  
9:    $swap(cQ, nQ)$   
10:  Set  $nQ$  to null  
11:   $l = l + 1$ 
```

---

# Second Approach: Queue BFS

## Kernel

- Threads are launched for each node in *vertex frontier queue*
- For all the neighbours of the node, update the distance if the node can be reached in fewer steps
- Add that node in the *edge frontier queue*
- Involves atomic operations

---

### Algorithm 4 queueBFS\_kernel

---

```
1: Input:  $l, V_a, E_a, Dist_a, P_a, cQ, nQ$    $\triangleright$  The graph  $G(V, E)$  and  
   source  $S$   
2:  $tid = \text{getThreadID}$   
3: if  $tid < cQ_{size}$  then  
4:    $u = cQ[tid]$   
5:   for all  $v = \text{neighbours of } u$  do  
6:     if  $Dist_a[v] = \infty$  and  $\text{atomicMin}(Dist_a[v], l + 1) = \infty$   
       then  
7:        $P_a[v] = u$   
8:        $pos = \text{atomicAdd}(nQ_{size}, 1)$   
9:        $nQ[pos] = v$ 
```

---



# Third Approach: Scan BFS

...

## Host

- Level-synchronous approach
- Needs 4 global synchronization
- Perform  $O(V + E)$  operations
- Improves on *parallelBFS* by populating vertex and edge frontier queue in linear operations
- Terminates when there are no *vertex frontiers*.

...

### Algorithm 5 ScanBFS\_Host

```
1: Input:  $V_a, E_a, S$  ▷ The graph  $G(V, E)$  and source  $S$ 
2: Create updating cost array  $Deg_a, PreDeg_a$  of size  $|V|$  and initialise all values to 0
3: Create cost array  $Dist_a$  of size  $|V|$  and initialise all values to  $\infty$ 
4: Create mask array  $P_a$  of size  $|V|$  and initialise all values to  $-1$ 
5: Create two array  $cQ$  and  $nQ$ , and initialise it to  $S$  and null respectively.
6:  $Dist_a[S] = 0$ 
7:  $P_a[S] = -1$ 
8:  $l = 0$  ▷ Start with the source vertex
9: while  $cQ_{size} > 0$  do
10:   Invoke nextLayer( $l, V_a, E_a, P_a, Dist_a, cQ$ )
11:   Invoke countDegrees( $V_a, E_a, P_a, cQ, Deg_a$ )
12:   Invoke scanDegrees( $cQ_{size}, Deg_a, PreDeg_a$ )
13:   Perform Prefix Sum on  $Deg_a$ , and store the results in  $PreDeg_a$ 
14:    $nQ = PreDeg_a[cQ_{size}/NUM\_THREADS]$ 
15:   Invoke populateNextQueue( $V_a, E_a, P_a, cQ, nQ, Deg_a, PreDeg_a$ )
16:    $cQ = nQ$ 
17:    $l = l + 1$ 
```





# Third Approach: Scan BFS

...

## Kernels

- Threads are launched for each node in the vertex frontier
- Uses Blelloch's prefix sum on the number of edge frontiers contributed by a vertex.
- Fills up the the edge frontier using the computed prefix sum.

...



---

### Algorithm 6 nextLayer

---

```
1: Input:  $l, V_a, E_a, P_a, Dist_a, cQ$ 
2:  $tid = \text{getThreadId}()$  ▷ Get the Id of the thread
3: if  $tid < cQ_{size}$  then
4:    $u = cQ[tid]$ 
5:   for all  $v = \text{neighbours of } u$  do
6:     if  $Dist_a[v] > l + 1$  then
7:        $Dist_a[v] = l + 1$ 
8:        $P_a[v] = u$ 
```

---

---

### Algorithm 7 countDegrees

---

```
1: Input:  $V_a, E_a, P_a, cQ, Deg_a$ 
2:  $tid = \text{getThreadId}()$  ▷ Get the Id of the thread
3: if  $tid < cQ_{size}$  then
4:    $u = cQ[tid]$ 
5:    $d = 0$ 
6:   for all  $v = \text{neighbours of } u$  do
7:     if  $P_a[v] = E_a.index(v) \text{ and } v \neq u$  then
8:        $d = d + 1$ 
9:    $Deg_a[tid] = d$ 
```

---





# Third Approach: Scan BFS

...

## Kernels


### Algorithm 9 populateNextQueue

```
1: Input:  $V_a, E_a, P_a, cQ, nQ, Deg_a, PreDeg_a$ 
2:  $tid = \text{getThreadId}()$  ▷ Get the Id of the thread
3: if  $tid < cQ_{size}$  then
4:   Initialise a shared variable  $i$ 
5:   if  $threadId.x = 0$  then
6:      $i = PreDeg_a[ NUM\_THREADS ]$ 
7:    $sync\_threads$ 
8:    $s = 0$ 
9:   if  $threadId.x \neq 0$  then
10:     $s = Deg_a[tid - 1]$ 
11:    $u = cQ[tid]$ 
12:    $c = 0$ 
13:   for all  $v = \text{neighbours of } u$  do
14:     if  $P_a[v] = E_a.index(v)$  and  $v \neq u$  then
15:        $nQ[i + s + c] = v$ 
16:        $c = c + 1$ 
```



### Algorithm 8 scanDegrees

```
1: Input:  $cQ_{size}, Deg_a, PreDeg_a$ 
2:  $tid = \text{getThreadId}()$  ▷ Get the Id of the thread
3: if  $tid < cQ_{size}$  then
4:   Create a shared array  $preSum$  of size  $NUM\_THREADS$ 
5:    $m = threadId.x$ 
6:    $preSum[m] = Deg_a[tid]$ 
7:    $sync\_threads$ 
8:    $n = 2$ 
9:   while  $n \leq NUM\_THREADS$  do
10:    if  $\text{bitwiseAnd}(m, n - 1) = 0$  and  $tid + (2 * n) < cQ_{size}$ 
11:    then
12:       $preSum[m] += preSum[tid + (2 * n)]$ 
13:       $sync\_threads$ 
14:       $n = 2 * n$ 
15:    if  $m = 0$  then
16:       $PreDeg_a[tid / NUM\_THREADS + 1] = preSum[m]$ 
17:       $n = NUM\_THREADS$ 
18:      while  $n > 1$  do
19:        if  $\text{bitwiseAnd}(m, n - 1) = 0$  and  $tid + (n/2) < cQ_{size}$ 
20:        then
21:           $temp = preSum[m]$ 
22:           $preSum[m] += preSum[tid + (n/2)]$ 
23:           $preSum[tid + (n/2)] = temp$ 
24:           $sync\_threads$ 
25:           $n = n/2$ 
26:       $Deg_a[tid] = preSum[m]$ 
```






# Results and Analysis



## Various Observations:

- *parallelBFS* gives the best results for all possible graphs
- *queueBFS* performs well on **dense** graphs
- *scanBFS* performs well on **sparse** graphs


# Single Source Shortest Path



# Our Approach

Bugged Parallel  
Dijkstra

Coarsening



Corrected  
Version

# First Approach: Bugged Parallel Dijkstra

- Algorithm proposed by Harish et. al.
- Uses an updating cost array  $\mathbf{U}_a$  as intermediate to update the actual cost array  $\mathbf{C}_a$ . Prevents RAW and WAR data hazards.
- Operates in two sequential phases over multiple iterations.
- The boolean array  $\mathbf{M}_a$  and variable **flag** determines the termination of the algorithm.

...

---

**Algorithm 10** SSSP\_Host

---

- 1: **Input:**  $V_a, E_a, W_a, S$      $\triangleright$  The graph  $G(V, E, W)$  and source  $S$
  - 2: Create updating cost array  $U_a$  of size  $|V|$  and initialise all values to  $\infty$
  - 3: Create cost array  $C_a$  of size  $|V|$  and initialise all values to  $\infty$
  - 4: Create mask array  $M_a$  of size  $|V|$  and initialise all values to *false*
  - 5:  $U_a[S] = 0$
  - 6:  $C_a[S] = 0$
  - 7:  $M_a[S] = \text{flag} = \text{true}$      $\triangleright$  Start with the source vertex
  - 8: **while** *flag* **do**
  - 9:    *flag* = *false*
  - 10:    **for all**  $v \in V$  in parallel **do**
  - 11:        Invoke SSSP\_Phase1( $V_a, E_a, W_a, C_a, U_a, M_a$ )
  - 12:        Invoke SSSP\_Phase2( $C_a, U_a, M_a, \text{flag}$ )
- 

...

...

# First Approach: Bugged Parallel Dijkstra

- In Phase 1, the vertices in  $\mathbf{M}_a$  are treated as potential intermediaries for a shortest path.
- The distance to neighbours of such vertices are updated (Line 5 - 7).
- In Phase 2,  $\mathbf{C}_a$  is updated using  $\mathbf{U}_a$ , and corresponding bit is set in  $\mathbf{M}_a$ .
- If no such update, the algorithm is terminated through the **flag** variable.

...

---

## Algorithm 11 SSSP\_Phase1

---

```
1: Input:  $V_a, E_a, W_a, C_a, U_a, M_a$ 
2:  $tid = \text{getThreadId}()$  ▷ Get the Id of the thread
3: if  $M_a[tid] = \text{true}$  then
4:    $M_a[tid] = \text{false}$ 
5:   for all neighbours  $nid$  of  $tid$  do ▷ Line 6, 7 must be atomic
6:     if  $U_a[nid] > C_a[tid] + W_a[nid]$  then
7:        $U_a[nid] = C_a[tid] + W_a[nid]$ 
```

---

---

## Algorithm 12 SSSP\_Phase2

---

```
1: Input:  $C_a, U_a, M_a, \text{flag}$ 
2:  $tid = \text{getThreadId}()$  ▷ Get the Id of the thread
3: if  $C_a[tid] > U_a[tid]$  then
4:    $M_a[tid] = \text{flag} = \text{true}$ 
5:    $C_a[tid] = U_a[tid]$ 
```

---

...

...

# Second Approach: Corrected Parallel Dijkstra

...

- Above presented algorithm is bugged, pointed out by Martin et. al.
- Using  $\mathbf{U}_a$  prevents RAW and WAR in  $\mathbf{C}_a$ , but does nothing for **WAW** dependencies in  $\mathbf{U}_a$  (Line 6 and 7).
- During simultaneous update of  $\mathbf{U}_a[nid]$  in Line 7, we need the smallest value to be retained.
- Solve it indirectly, by executing Line 6 and 7 **atomically**.

...

---

## Algorithm 11 SSSP\_Phase1


---

```
1: Input:  $V_a, E_a, W_a, C_a, U_a, M_a$ 
2:  $tid = \text{getThreadId}()$  ▷ Get the Id of the thread
3: if  $M_a[tid] = \text{true}$  then
4:    $M_a[tid] = \text{false}$ 
5:   for all neighbours  $nid$  of  $tid$  do ▷ Line 6, 7 must be atomic
6:     if  $U_a[nid] > C_a[tid] + W_a[nid]$  then
7:        $U_a[nid] = C_a[tid] + W_a[nid]$ 
```


---

...






# Results and Analysis



## Observations


- Heavily dependent on number of **edges** (Graph 2 vs Graph 3 & 4)
- Speedup of GPU compared to CPU is dependent on the **density of graph**.
- Lower the density, higher the speedup offered.
- **Density** [G1(0.5) >> G5( $10^{-6}$ )],  
**SpeedUp** [G1(4x) << G5(95x)]








# Third Approach: Thread Coarsening

- Smaller number of more Coarse-grained threads are being executed
- Instructions executed by a number of different threads are merged into a single thread.
- For finding the optimal thread coarsening factor, we used the manual approach of merging.





# Results and Analysis



## Various Observations:

- A consistent fall of performance with increase in the thread coarsening factor ( **c.f** ).
- Reason can be reduction in parallelism with increasing **c.f**
- Another reason can be the subsequent increase in pressure on the kernel.

# All Pair Shortest Path



# First Approach: Using SSSP

What if we call SSSP on all the vertices?

- Works well on sparse graphs.
- Serial time complexity:  $O(V^2 \log V + EV)$

...

---

**Algorithm 13** APSP\_Using\_SSSP

---

```
1: Input:  $V_a, E_a, W_a$  ▷ The graph  $G(V, E, W)$ 
2: Create updating cost array  $U_a$  of size  $|V|$ 
3: Create cost array  $C_a$  of size  $|V|$ 
4: Create mask array  $M_a$  of size  $|V|$  and initialise all values to false
5: Create a 2d output array  $O_a$  of size  $|V| \times |V|$ 
   for all  $S \in V$  do
6:   Assign  $\infty$  to all values of  $U_a$  and  $C_a$ 
7:    $U_a[S] = C_a[S] = 0$ 
8:    $M_a[S] = \text{flag} = \text{true}$  ▷ Start with the source vertex
9:   while flag do
10:    flag = false
11:    for all  $v \in V$  in parallel do
12:      Invoke SSSP_Phase1( $V_a, E_a, W_a, C_a, U_a, M_a$ )
13:      Invoke SSSP_Phase2( $C_a, U_a, M_a, \text{flag}$ )
14:    Copy the distances in  $C_a$  to  $O_a[S]$ 
```

---

...


...

# Second Approach: Floyd Warshall

...

## Floyd Warshall in Brief

all intermediate vertices in  $\{1, 2, \dots, k-1\}$     all intermediate vertices in  $\{1, 2, \dots, k-1\}$



...

$$d_{ij}^{(k)} = \begin{cases} w_{ij} & \text{if } k = 0, \\ \min(d_{ij}^{(k-1)}, d_{ik}^{(k-1)} + d_{kj}^{(k-1)}) & \text{if } k \geq 1. \end{cases}$$

...

...

# Second Approach: Floyd Warshall

...

## Implementing FW in GPU

---

**Algorithm 14** APSP\_Naive\_FW

---

```
1: Input:  $E$      $\triangleright G(V, E, W)$  as adjacency matrix of size  $|V| \times |V|$   
2: Initialise the non-edges in  $E$  as  $\infty$   
3: for all  $k \in V$  do  
4:    for all  $(u, v) \in V \times V$  in parallel do  
5:      $E[u, v] = \min(E[u, v], E[u, k] + E[k, v])$ 
```

---

$|V|$  iterations

One thread per  
element

- Uses the adjacency matrix representation rather than the adjacency list representation discussed earlier.
- $O(|V|^2)$  threads and  $O(|V|)$  iterations.

...

...

# Third Approach: Blocked Floyd Warshall

...

A deeper look into FW

		1		
1	2		4	5
		3		
		4		
		5		

Solution: Use Tiling

...

...

# Third Approach: Using Blocked Floyd Warshall

First Phase


**Diagonal Blocks**

Second Phase

**Blocks row and column  
aligned to diagonal**

Third Phase


**All other blocks**





# Third Approach: Using Blocked Floyd Warshall


Phases when block (1,1) is the self-dependent block



Phase 1




Phase 2




Phase 3

Phases when block (t,t) is the self-dependent block




Phase 1



Phase 2




Phase 3





# Third Approach: Using Blocked Floyd Warshall

Independent Phase




$K = 1 \text{ to } 4$





# Third Approach: Using Blocked Floyd Warshall

Partially dependent Phase




$K = 1 \text{ to } 4$






# Third Approach: Using Blocked Floyd Warshall

Double dependent Phase





$K = 1 \text{ to } 4$





# Results and Analysis



## Various Observations:

- Naive GPU FW performs better than using SSSP for dense graphs.
- **Blocked FW has a massive improvement in performance.**
- **FW does not depend on sparsity** of the graph.

A decorative graphic consisting of blue circles of varying sizes connected by thin black lines, forming a network-like structure around the central text. Some circles have concentric rings. Ellipses (...) are placed at several points along the lines.

**Thanks!**