

Compiler Support for Branch Handling

Using compilers for hard to predict branches

Prediction:

In modern processors, penalty for mispredicted branch is ~~large~~ huge.

A pipeline may be filled with instructions executed in parallel — say a 12 stage pipe with 4 IPC.
~ 48 instructions in the pipe

→ will go waste if prediction is wrong.

Workaround → Do the work both ways
[taken & not taken]

In a branch you always ~~waste~~ waste exactly 50% of the ~~prediction~~ executed instruction if you use predication.

Where useful??


Consider a loop → in every iteration you execute loop instr & also outside loop.



→ lot of work wasted ~~as~~ if loop is deep
Better to predict

Fn call / return: Unless the call is conditional it is always taken, ~~but~~ hence no point of predication.

A large if-then-else




~ amount of waste may be too much w.r.t. B.P.

∴ for large if-then-else we should be predicting.

Small if-then-else → less wastage

∴ makes sense to do predication.



→ wastage 5 instr

Consider 100 instances

→ we waste 500 instr

using BP with 90% accuracy & penalty of 50 instr

→ we waste $\frac{10 \times 50}{100} = 500$

↓
wrong instances

∴ for smaller if-then-else, we want to predicate.

if conversion

```
if (cond) {
    x = arr[i];
    x = x + 1;
}
```

```
else {
    x = arr[j];
    x = x - 1;
}
```

⇒

```
x1 = arr[i];
```

```
x2 = arr[j];
```

```
x1 = x1 + 1;
```

```
x2 = x2 - 1;
```

```
x = cond ? x1 : x2
```

```
x = cond ? x1 : x2
```

↑

still need a branch instr

alternate idea
have hw support
for a conditional 'mov' instr

Conditional Mov

Simplest predication \rightarrow mov x x, flag

MIPS \rightarrow MOVZ R_d R_n R_t

\Rightarrow if (R_t == 0)
R_d = R_s ;

\Rightarrow there is not branching to anywhere.

MOVN

if (R_t != 0)

X = cond ? X1 : X2 ;

\Downarrow

R₃ = cond ;
 \swarrow
result

R₁ = X₁ expression result

R₂ = -- X₂

MOVN X, R₁, R₃

MOVZ X, R₂, R₃

X86 ! \rightarrow CMov instructions

Example in MOVN/MOVZ

80%
40 instr penalty

\rightarrow BEQZ R₁, Else

ADDI R₂, R₂, 1

100% \rightarrow B END

Else:

ADDI R₃, R₃, 1

END:

if
conv

ADDI R₄, R₂, 1
ADDI R₅, R₃, 1
MOVN R₂, R₄, R₁
MOVZ R₃, R₅, R₁

\downarrow
4 instr

Performance

3 instr on one path

2 instr on other path

consider 50% taken & 50% not taken

2.5 instr + 0.2 x 40

= 10.5 instr on average

(4)
Need compiler support for if conversion.
Such conversion require more registers.

A branch hard to predict as ~~define~~
can be found by profiling \rightarrow we can
perform if ~~conversion~~ conversion.

In general case \rightarrow if conv requires
more instructions.

full predication \rightarrow make all instr in
conditional,

this alleviates the problem of
requiring more instructions or extra
registers.

Full predication HW support

EX: \rightarrow ITANIUM ISA

\rightarrow every instr has conditional bits.
 \rightarrow 6 bits specify condition.

\rightarrow tells which of 64 registers (1 bit)
to use to determine execution.

Earlier example

(i) MPEOZ P1, P2 R1
predicate at instr, P1 & P2 are always not
opposite
if R1 = 0, P1 true, P2 false
& vice versa.

(ii) (P2) ADDI R2, R2, 1

(iii) (P1) ADDI R3, R3, 1

5

Compiler Support for ILP

HW has a limited window of view of the program \rightarrow cannot track dependency beyond the window limited by ROB size.

Co-Compiler has an overall view.

True hight Reduction technique: \rightarrow

$$R8 = R2 + R3 + R4 + R5$$

ADD R8, R2, R3
ADD R8, R8, R4
ADD R8, R8, R5

3 cycles

exploit associativity of '+'.
True hight Reduction opt

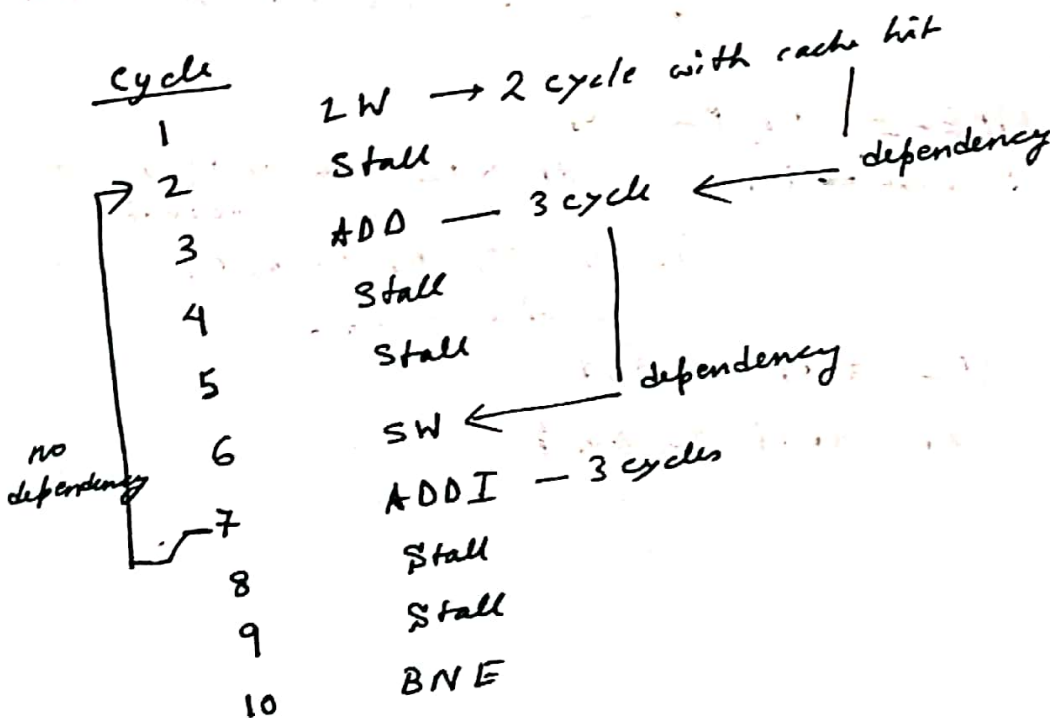
ADD R8, R2, R3
ADD R8, R8, R4
ADD R8, R8, R5
2 cycles.

Intra scheduling by compiler
 ϕ denotes address in R1.

Loop:

LW R2, ϕ (R1)
ADD R2, R2, R0
SW R2, ϕ (R1)
ADDI R1, R1, 4
BNE R1, R3, Loop

Consider a simple Processor,
w/o R/s & o-o-o support.



A transformed but equivalent code


loop:

```

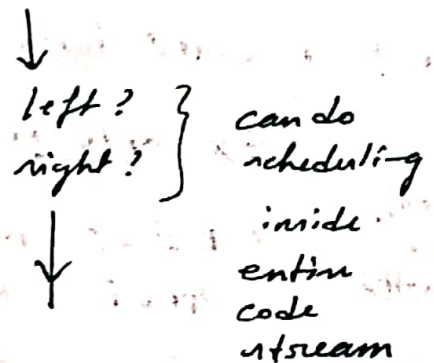
LW R2, 0(R1)
ADDI R1, R1, 4
ADD R2, R2, R0
SW R2, -4(R1)
BNE R1, R3, loop
    
```

<u>Cycle</u>	
1.	LW
2.	ADDI
3.	ADD
4.	Stall
5.	Stall
6.	SW
7.	BNE

Scheduling & If - conversion



\Rightarrow



if - convert a loop

Loop:

```

LW R2, 0(R1) ← 1 stall cycle
ADD R2, R2, R3 ← 1 stall
SW R2, -4(R1) ← no stall, addI don't depend on SW
ADDI R1, R1, 4 ← 1 stall
BNE R1, R2, loop
    
```

assume
all
instr
has 2
cycle
exec
time

after scheduling →

```

Loop: LW R2, 0(R1)
      ADDI R1, R1, 4
      ADD R2, R2, R3
      SW R2, -4(R1) ← 1, stall
      BNE R1, R5, Loop
  
```

A potential if-conversion can eliminate the branch & then some code after BNE can fill in the stall cycle, but as we have seen, if conversion of loops insert 1 predicate per iteration which is very costly.

Consider a loop which sometimes executes 10^4 times but sometimes only once, in predication it is always ^{exponentially} ~~that~~ many predicates, considering all the possibilities.

Alternative is to unroll the loop

for (i = 1000; i != 0; i--)

a[i] = a[i] + 5;

Loop: LW R2, 0(R1)
 ADD R2, R2, R3
 SW R2, -4(R1)
 ADDI R1, R1, -4
 BNE R1, R5, Loop

No. of instr = $5 \times 1000 = 5000$

for (; ; i = i - 2)
 { a[i] = a[i] + 5;
 a[i-1] = a[i-1] + 5;
 }

Loop: LW R2, 0(R1)
 ADD R2, R2, R3
 SW R2, -4(R1)
 LW R2, -4(R1)
 ADD R2, R2, R3
 SW R2, -4(R1)
 ADDI R1, R1, -8
 BNE R1, R5, loop

exectime
 = #instr × CPI × Cycle time

reduced
 definitely

same
 or may
 change

No. of instr
 = 8×500


= 4000

Reducing loop overhead

Loop unrolling benefit \rightarrow CPI \downarrow

Consider 4 issue in-order, perfect BP


Loop: LW R2, 0(R1)
ADD R2, R2, R3
SW R2, 0(R1)
ADDI R1, R1, -4
BNE R1, R5, loop




$$CPI = 3/5 = 0.6$$

perform previous re-ordering optimization

Loop: LW R2, 0(R1) \rightarrow ADDI R1, R1, -4
ADD R2, R2, R3
SW R2, 0(R1)
ADDI R1, R1, -4
BNE R1, R5, loop




Loop: LW R2, 0(R1)
ADDI R1, R1, -4
ADD R2, R2, R3
SW R2, 0(R1)
BNE R1, R5, loop



$$\therefore CPI = \frac{2}{5} = 0.4$$

Unroll w/o scheduling



$$CPI = \frac{5}{8} = 0.62$$

similar to original.

Unroll once + schedule

LW R2, 0(R1)

LW R10, -4(R1) → parallel load in ~~the~~ new register

loop:

LW R2, 0(R1)

LW R10, -4(R1)

ADD R2, R2, R3

ADD R10, R10, R3

ADDI R1, R1, -8

SW R2, 8(R1)

SW R10, 4(R1)

BNE R1, R5, loop

CPI

$$= \frac{3}{8} = 0.38$$

#instr reduced ↓ ← unrolling

CPI reduced ↓ ← more scheduling scope due to unrolling.

Downside to unrolling

Code bloating → increased code size

What if no. of iterations is unknown, we exit after executing extra instr.

Also, what if no. of iterations is not a multiple of the unroll factor,

Fn call inlining

Fn call

Code
parameter set up

fn call

code

⋮

Fn...

preparing fn params is not required with inlining.

→ enables instr scheduling for larger code sequence

(10)

Also, less instr \rightarrow CPI \downarrow .

Good for small fns which have less scope of scheduling inside fn body.

Fn call inlining downsize \rightarrow code bloat
If too many calls exist \rightarrow inlining increase
instr count very much.

Other IPC enhancing techniques.

Software pipelining

loop {
 LW a[i]
 ADD S
 SW a[i]
 ADDI i, 1
 BNE

\rightarrow think of loop as a pipeline


added
in prev
iteration

NO Dependency

SW a[i]
ADD S, a[i+1]
LW a[i+2]

Value added has been loaded
in previous iteration of the loop

Trace Scheduling



Say, $B_1 \rightarrow B_2 \rightarrow B_4$ is common path

