Multi Processing

Flynn's Taxonomy of Parallel machines.

		Instruction Stream	Data Stream
uniprocessor	SISD	1	1
Vector	SIMD	1	> 1
SSE/MMX			
Stream Processors	MISD	> 1	1 (not really used much)
	MIMD	> 1	> 1

MIMD
Multi Processor

Multi Processors need parallel programs

original single threaded code → difficult → parallel code development difficult
debugging also difficult

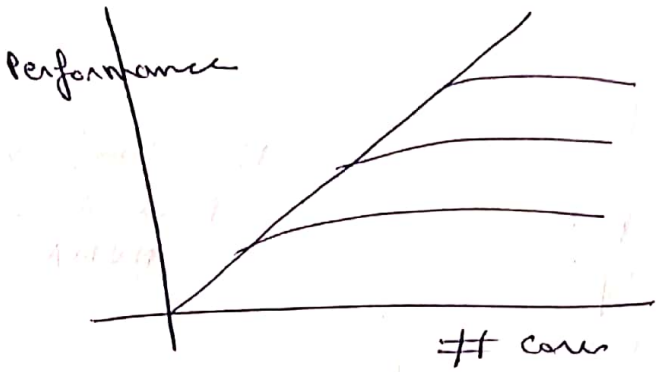

difficult to write scalable parallel code.

Centralized Shared Mem

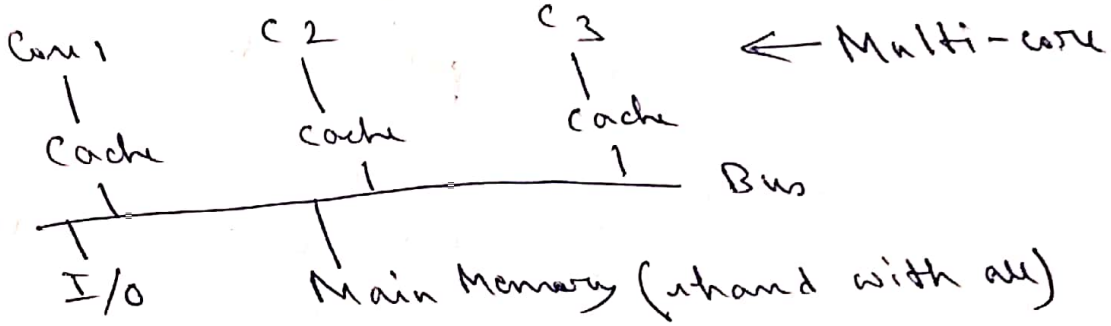

UMA : uniform memory access (time)

SMP : Symmetric Multi Processor.

Problem with centralized memory

A large memory is slow (access time more)

Memory B/W — all cores try to access one big slow mem

— B/W contention

— accesses get serialized. (queued)

∴ Centralized main memory is not good solution for machines with large no. of cores.

ok for 2 ~ 16 cores.

↓

Distributed ~~Shared~~ Memory

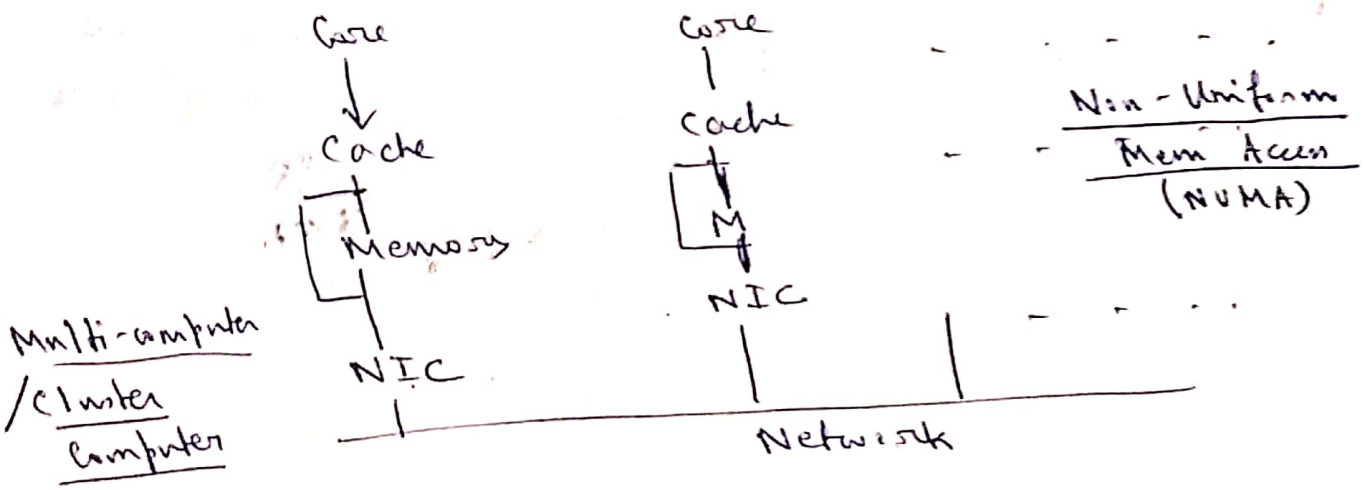

Core → req data → cache → Memory

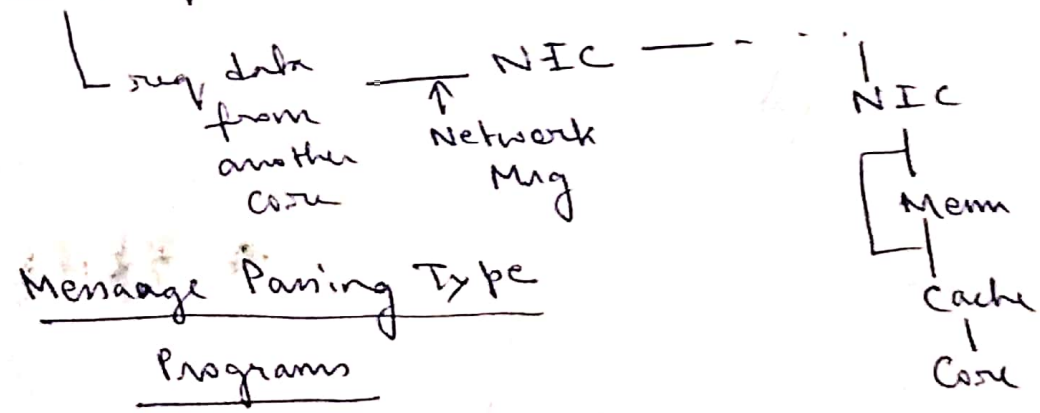

Message Passing Type Programs

Scalable Solution ⟶ support large # processors.

↓

programmer manages communication explicitly.

In shared mem, this is oblivious to the programmer.

Ex of MP program (message passing)

```
#define ASIZE 1024
#define Numproc 4
double myArray (ASIZE/Numproc);   // assume each proc already has 1/4 of array
double mySum = 0;                 // local sum variable
for (int i=0; i < ASIZE/Numproc; i++)
    mySum += myArray[i];          // local sum
if (myPID = 0) {
    for (int p=1; p < Numproc; p++) {
        int pSum; recv (p, pSum);
        mySum += pSum;
    }
    printf ("Sum: %lf \n", mySum);
}
else
    send (0, mySum);
```

one proc, compute overall sum

Shared Mem Program

```
#define ASIZE 1024
#define Numproc 4
shared double array[ASIZE];   <- array in shared mem
shared double allSum = 0;
shared mutex sumLock;
double myPID = 0;
for(int i = myPID * ASIZE / Numproc;
    i < (myPID+1) * ASIZE / Numproc; i++)
    mySum += array[i];
lock(sumLock);
allSum += mySum;
unlock(sumLock);
if(myPID = 0) printf("Sum: %lf \n", allSum);
```

job / proc →

i) no send / receive

ii) no array distribution

} critical section

insert a barrier here

Message Passing	Shared Mem
Communication: Programmer	Auto
Data distribution: Manual	Auto
HW support: Simple (need network HW)	Extensive
Program Correctness: Difficult to guarantee	Less difficult
Program Performance: Difficult	Very diff

↓

with correctness, performance also comes.

Shared Mem HW

Hyperthreading / Simultaneous Multi-threading (SMT)

using core-level HW to switch among threads
across clock ~~cly~~ cycles.

ii) execute multiple thread context in parallel. (saving & restoring registers for thread contexts)

HPCA notes

19 April 2020 13:43

Performance of Multithreaded code

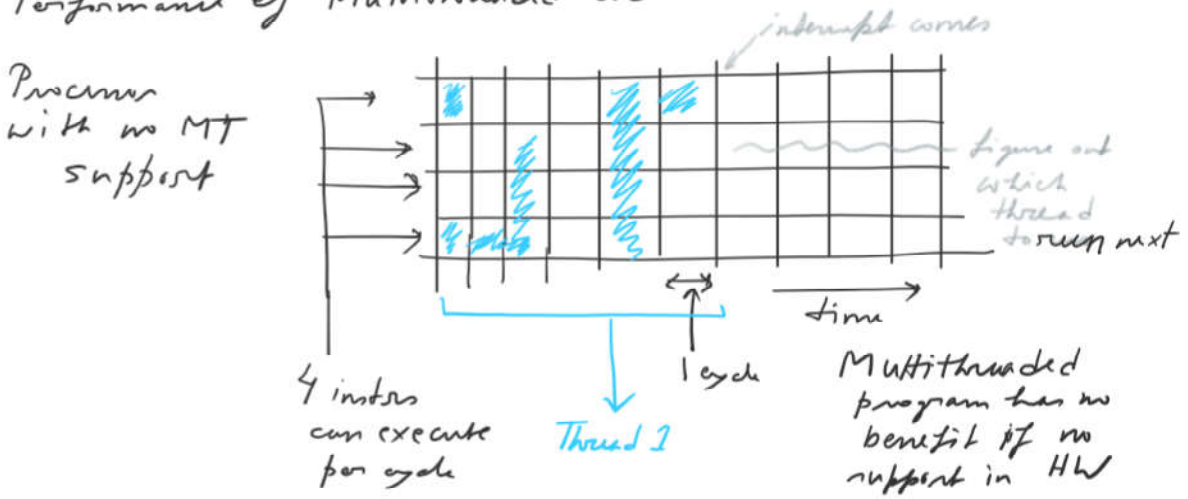

Chip Multiprocessor (CMP)

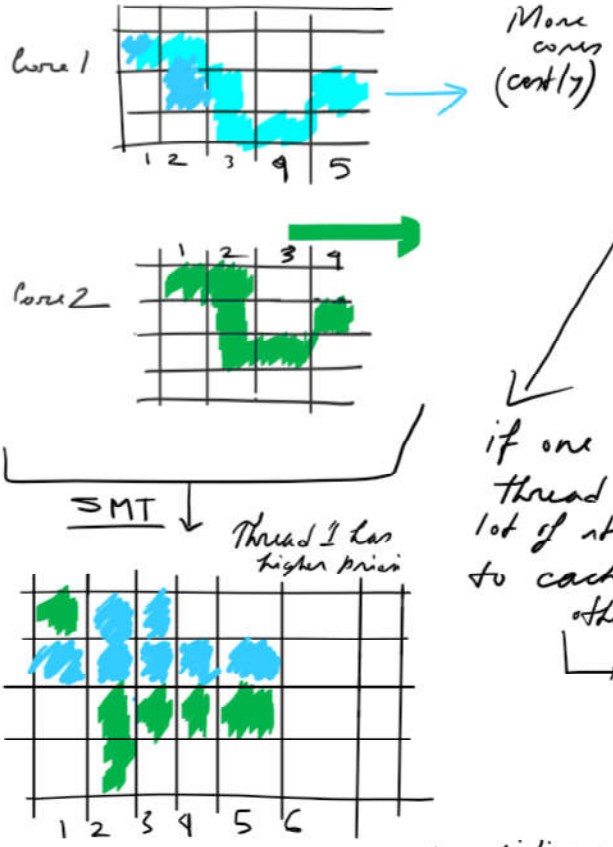

Fine grained MT (in one core)

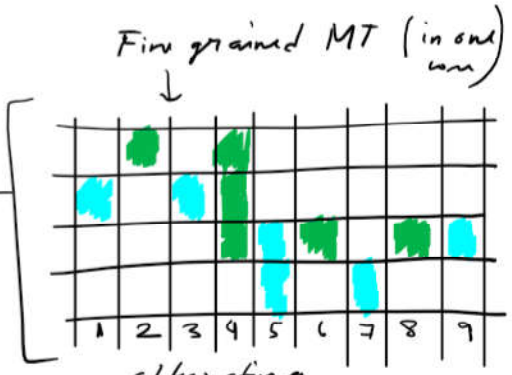

alternating (in every cycle) among threads using HW support

Coarse grain — alternate among threads after multiple cycles, not after every cycle

if one thread has lot of stalls (due to cache miss & other cores)
→ other threads with 'ready' instructions can fill in the next cycles.

using the underutilized issue slots of the multi-issue 4 way O-O-O superscalar processor.

HW cost → fetch stage need to be able to fetch from two PCs.
Register files → 2 set of Architectural registers.
Performance almost like CMP, little cost overhead of dependence/priority check during issue slot fill-up.

Normal O-O-O multi issue

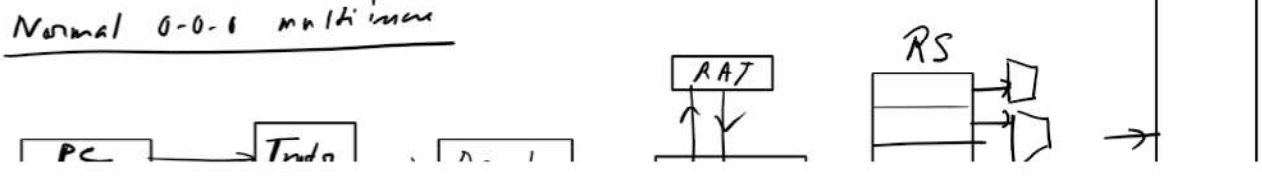

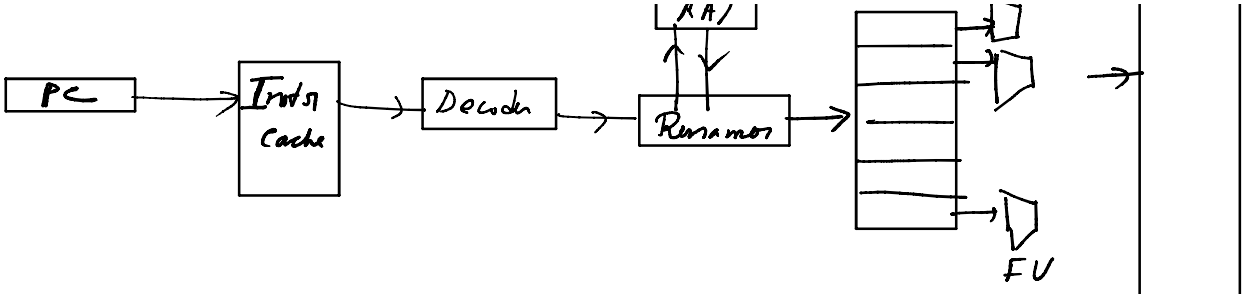

Modifications